

**FIWARE
Global
Summit**

Building real-time inference services based on ML with FIWARE

José Andrés Muñoz Arcentales, Javier Conde Díaz, Joaquín
Salvachúa

Vienna, Austria

12-13 June, 2023

#FIWARESummit

**From Data
to Value**

OPEN SOURCE
OPEN STANDARDS
OPEN COMMUNITY



**FIWARE
Global
Summit**

ML-OPS

Joaquín Salvachúa

Vienna, Austria

12-13 June, 2023

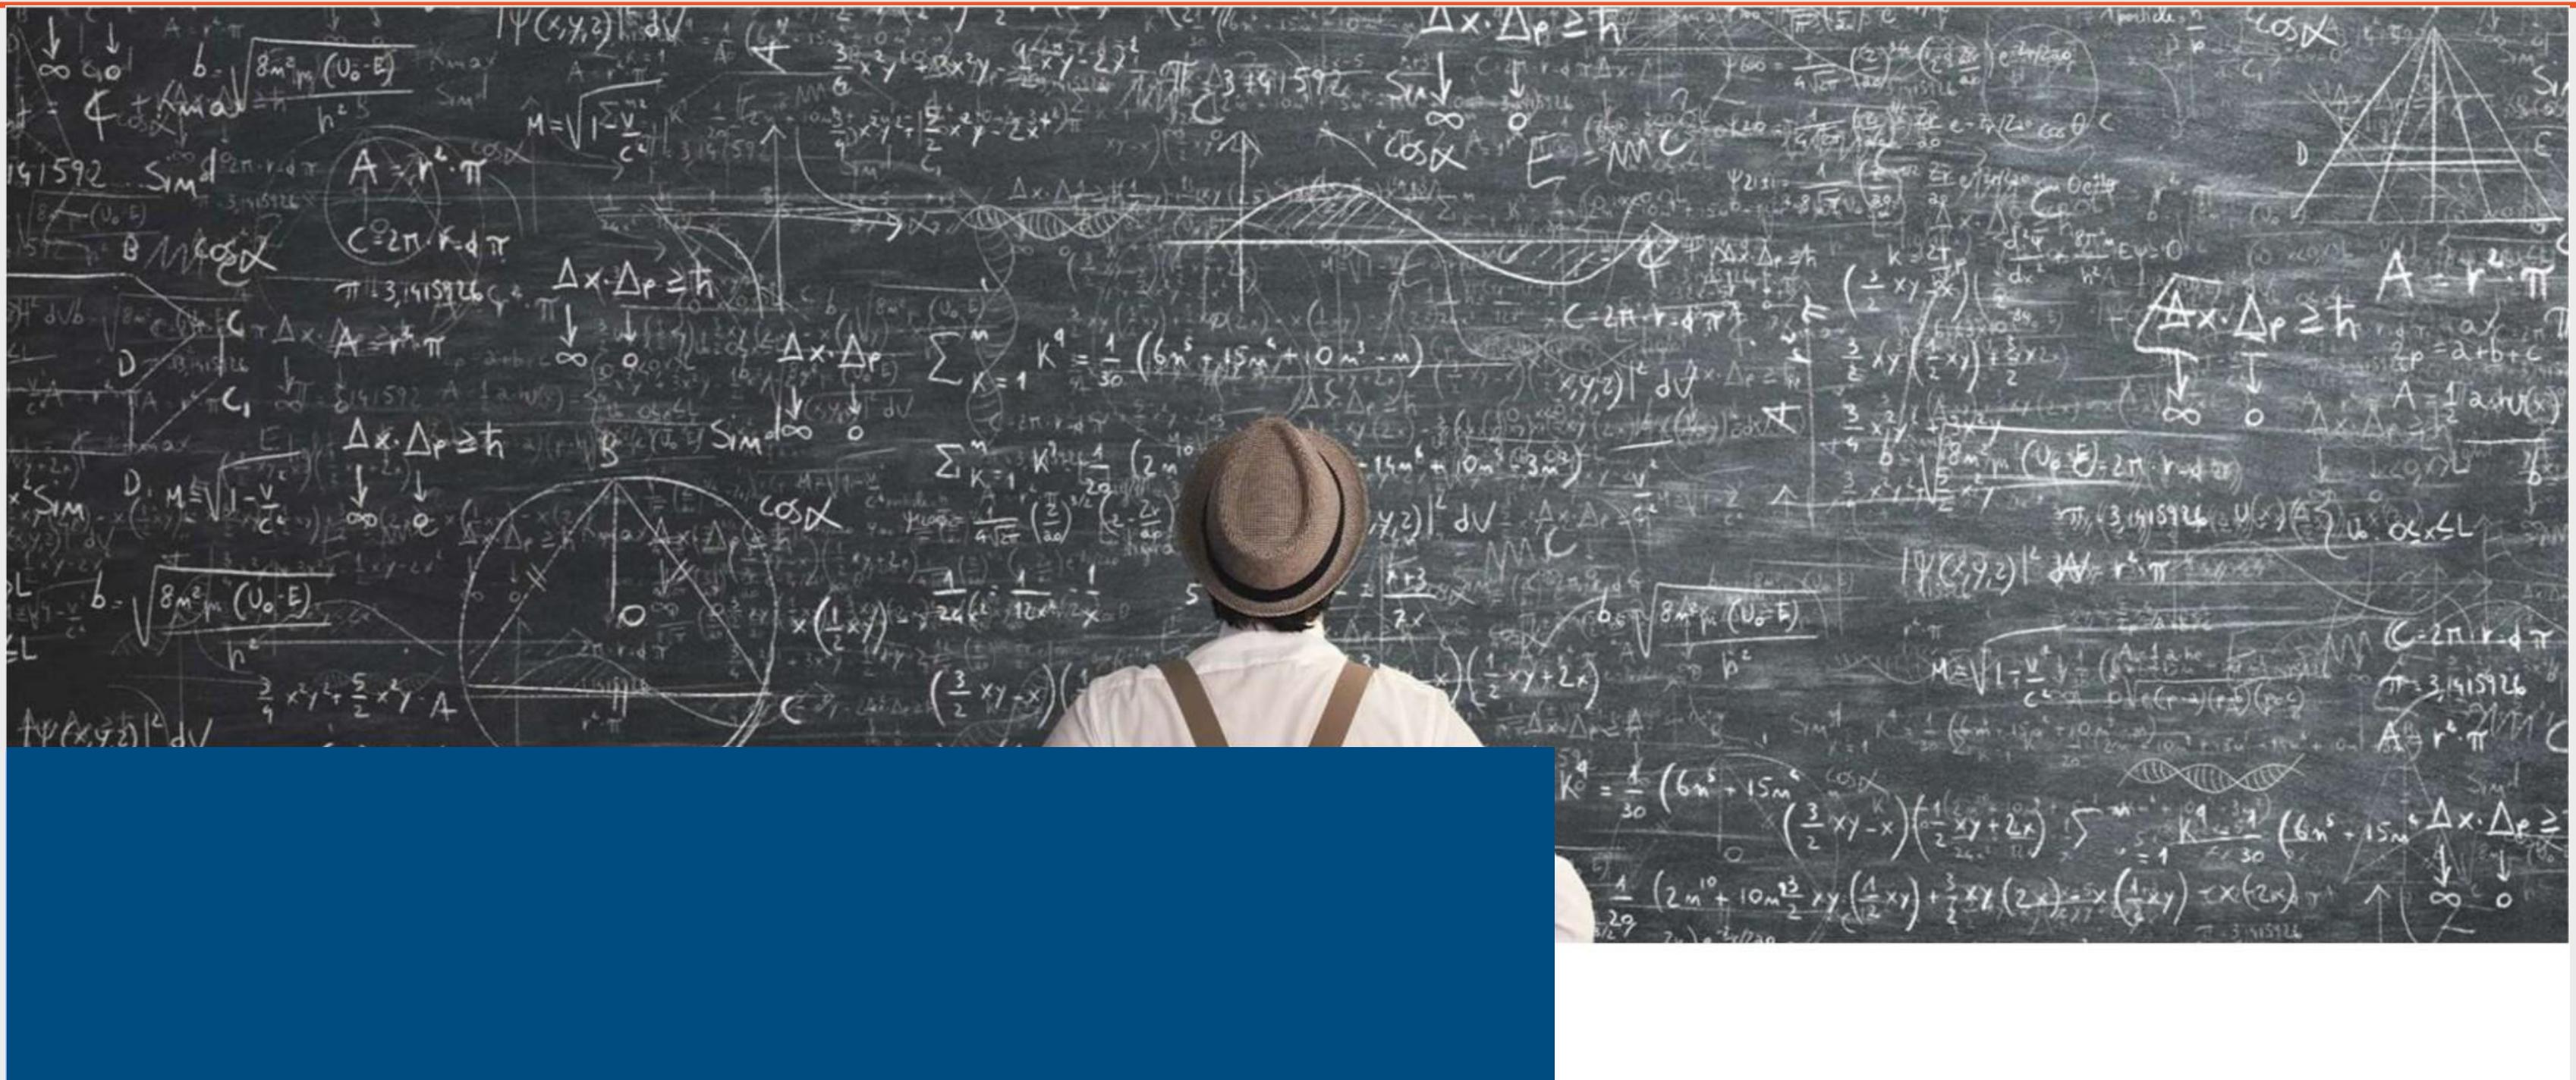
#FIWARESummit

**From Data
to Value**

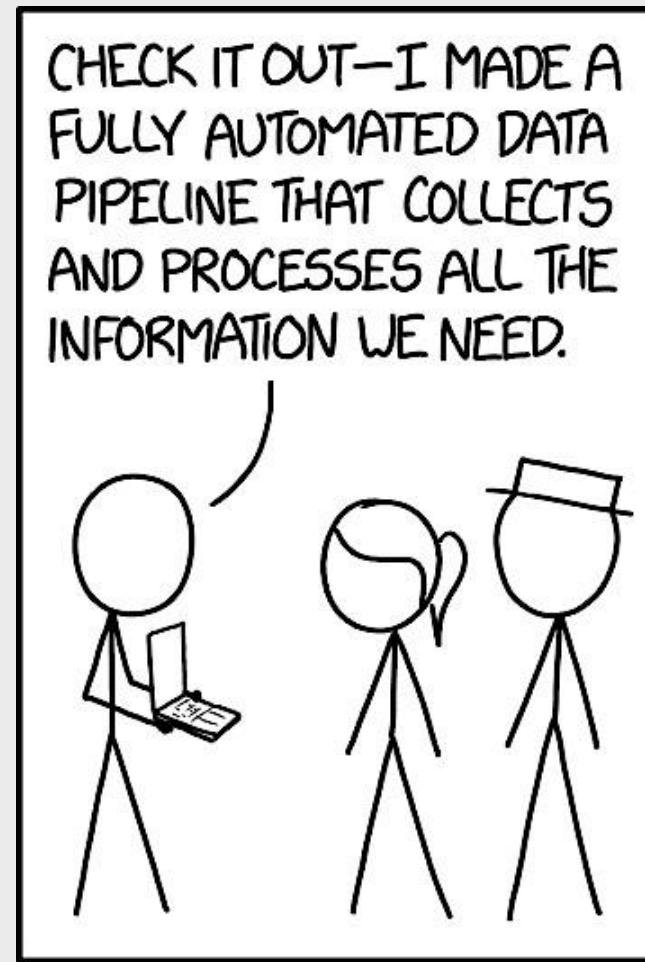
OPEN SOURCE
OPEN STANDARDS
OPEN COMMUNITY



ML is hard!

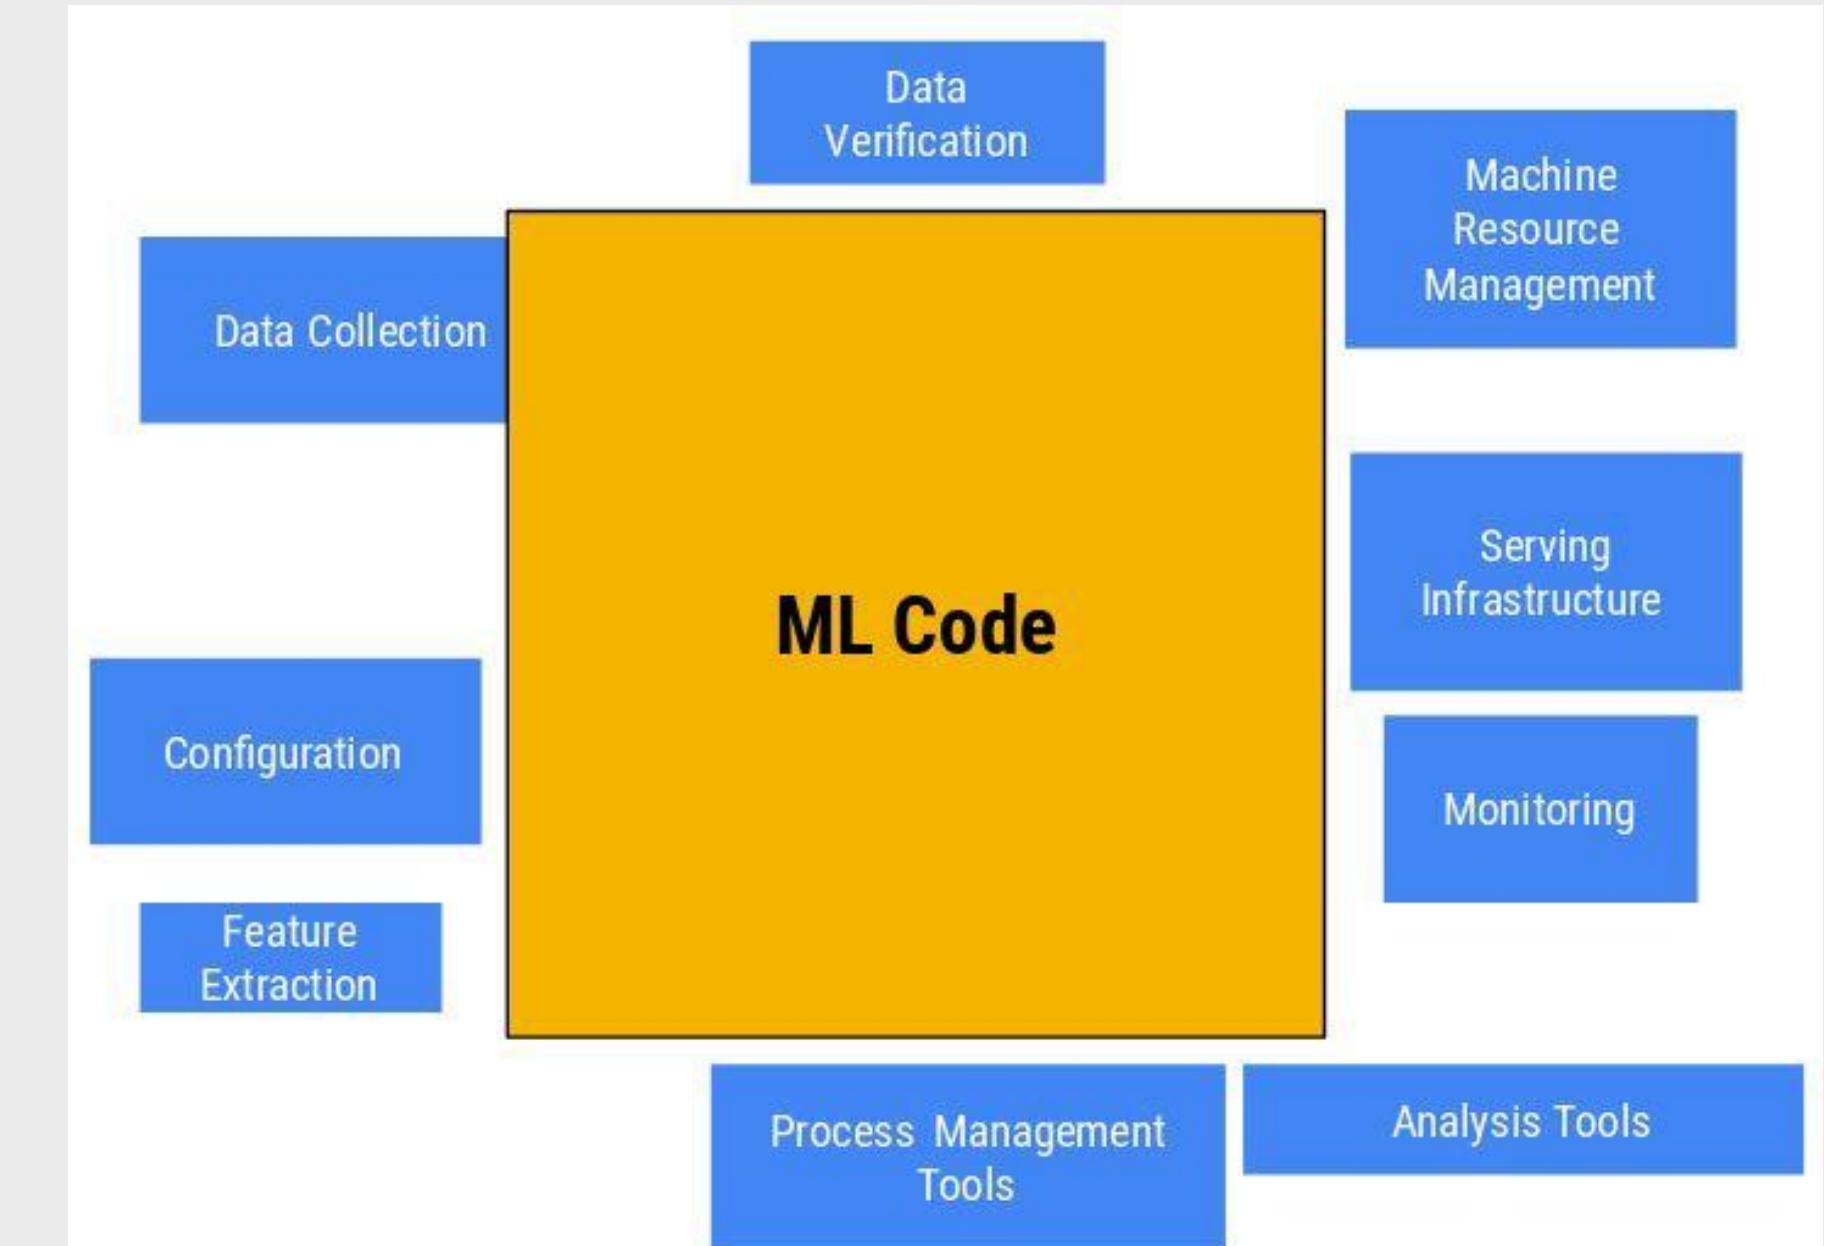


MAD · NOV 23-24 · 2018



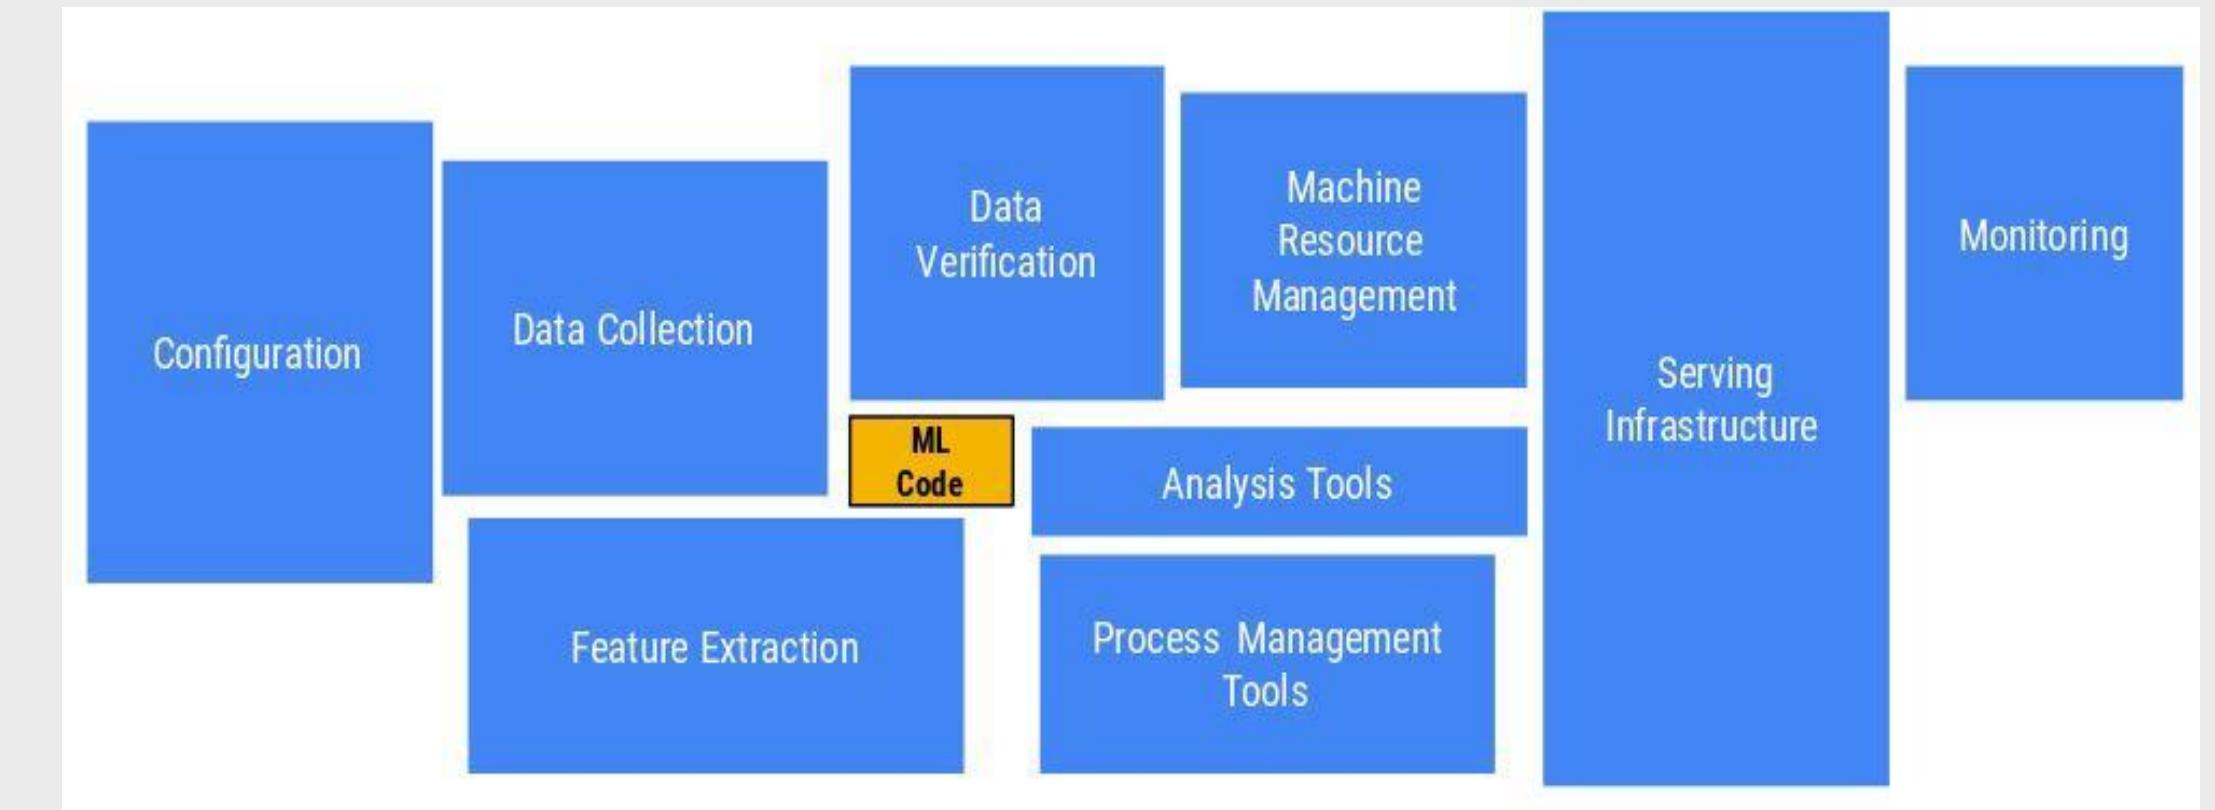
Current ML workflow

What you think



Current ML workflow

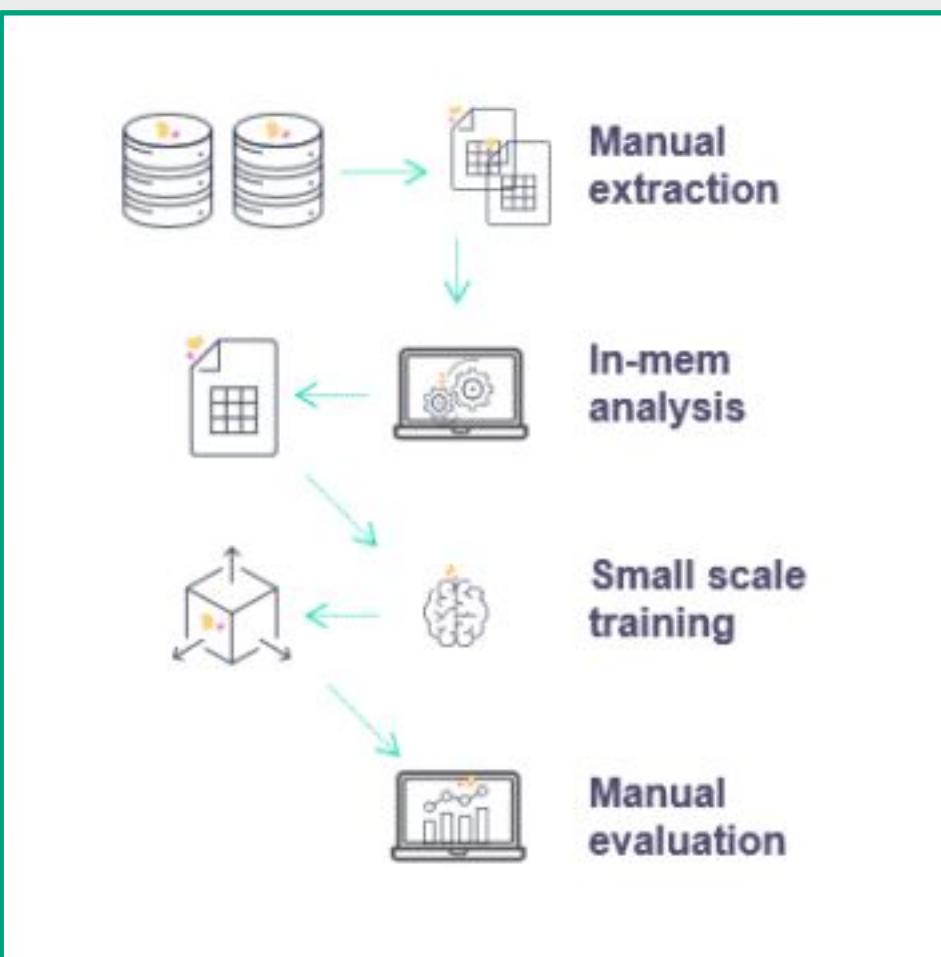
The reality



Source: <https://papers.nips.cc/paper/5656-hidden-technical-debt-in-machine-learning-systems.pdf>

Introduction

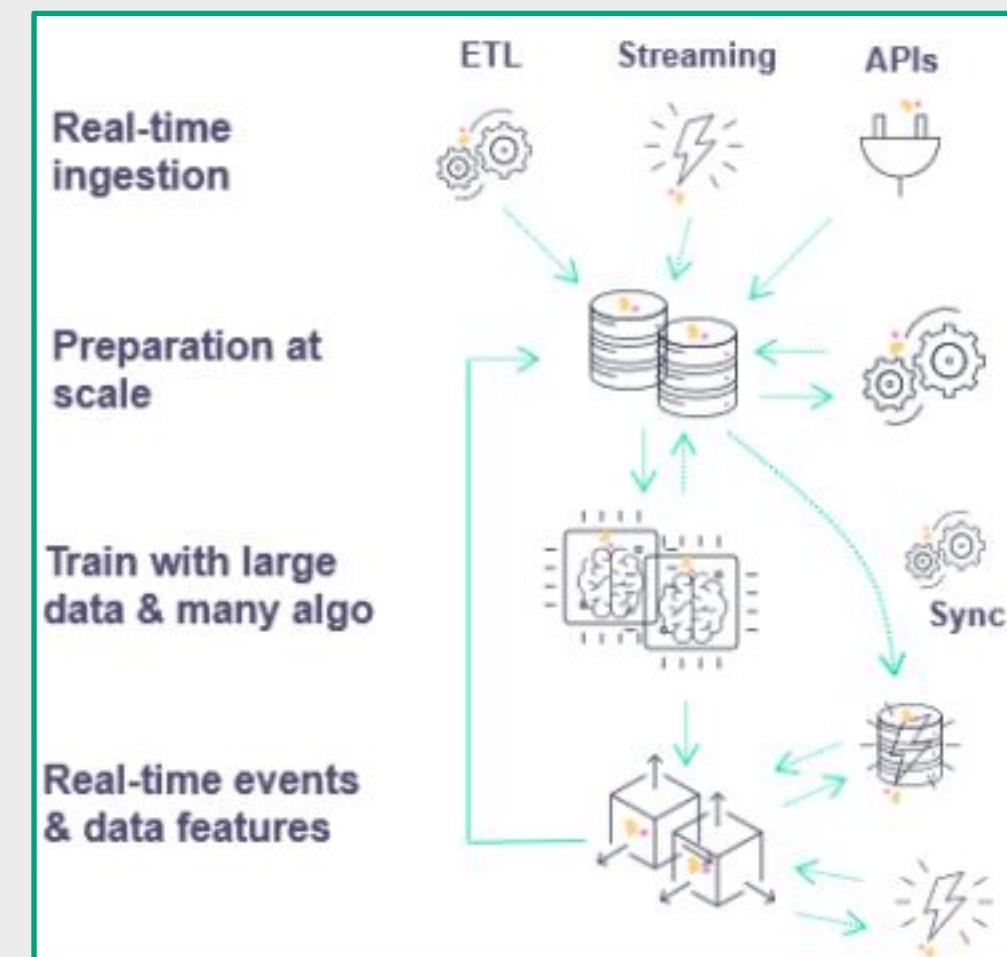
Research Environment



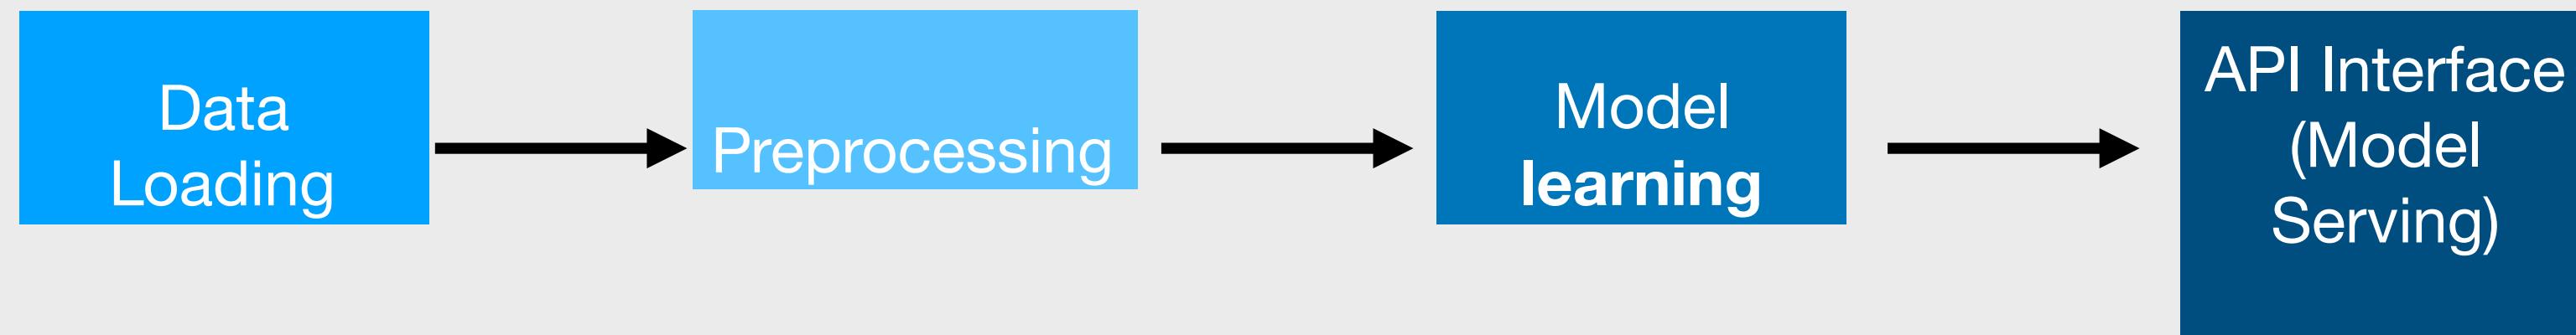
Build from Scratch

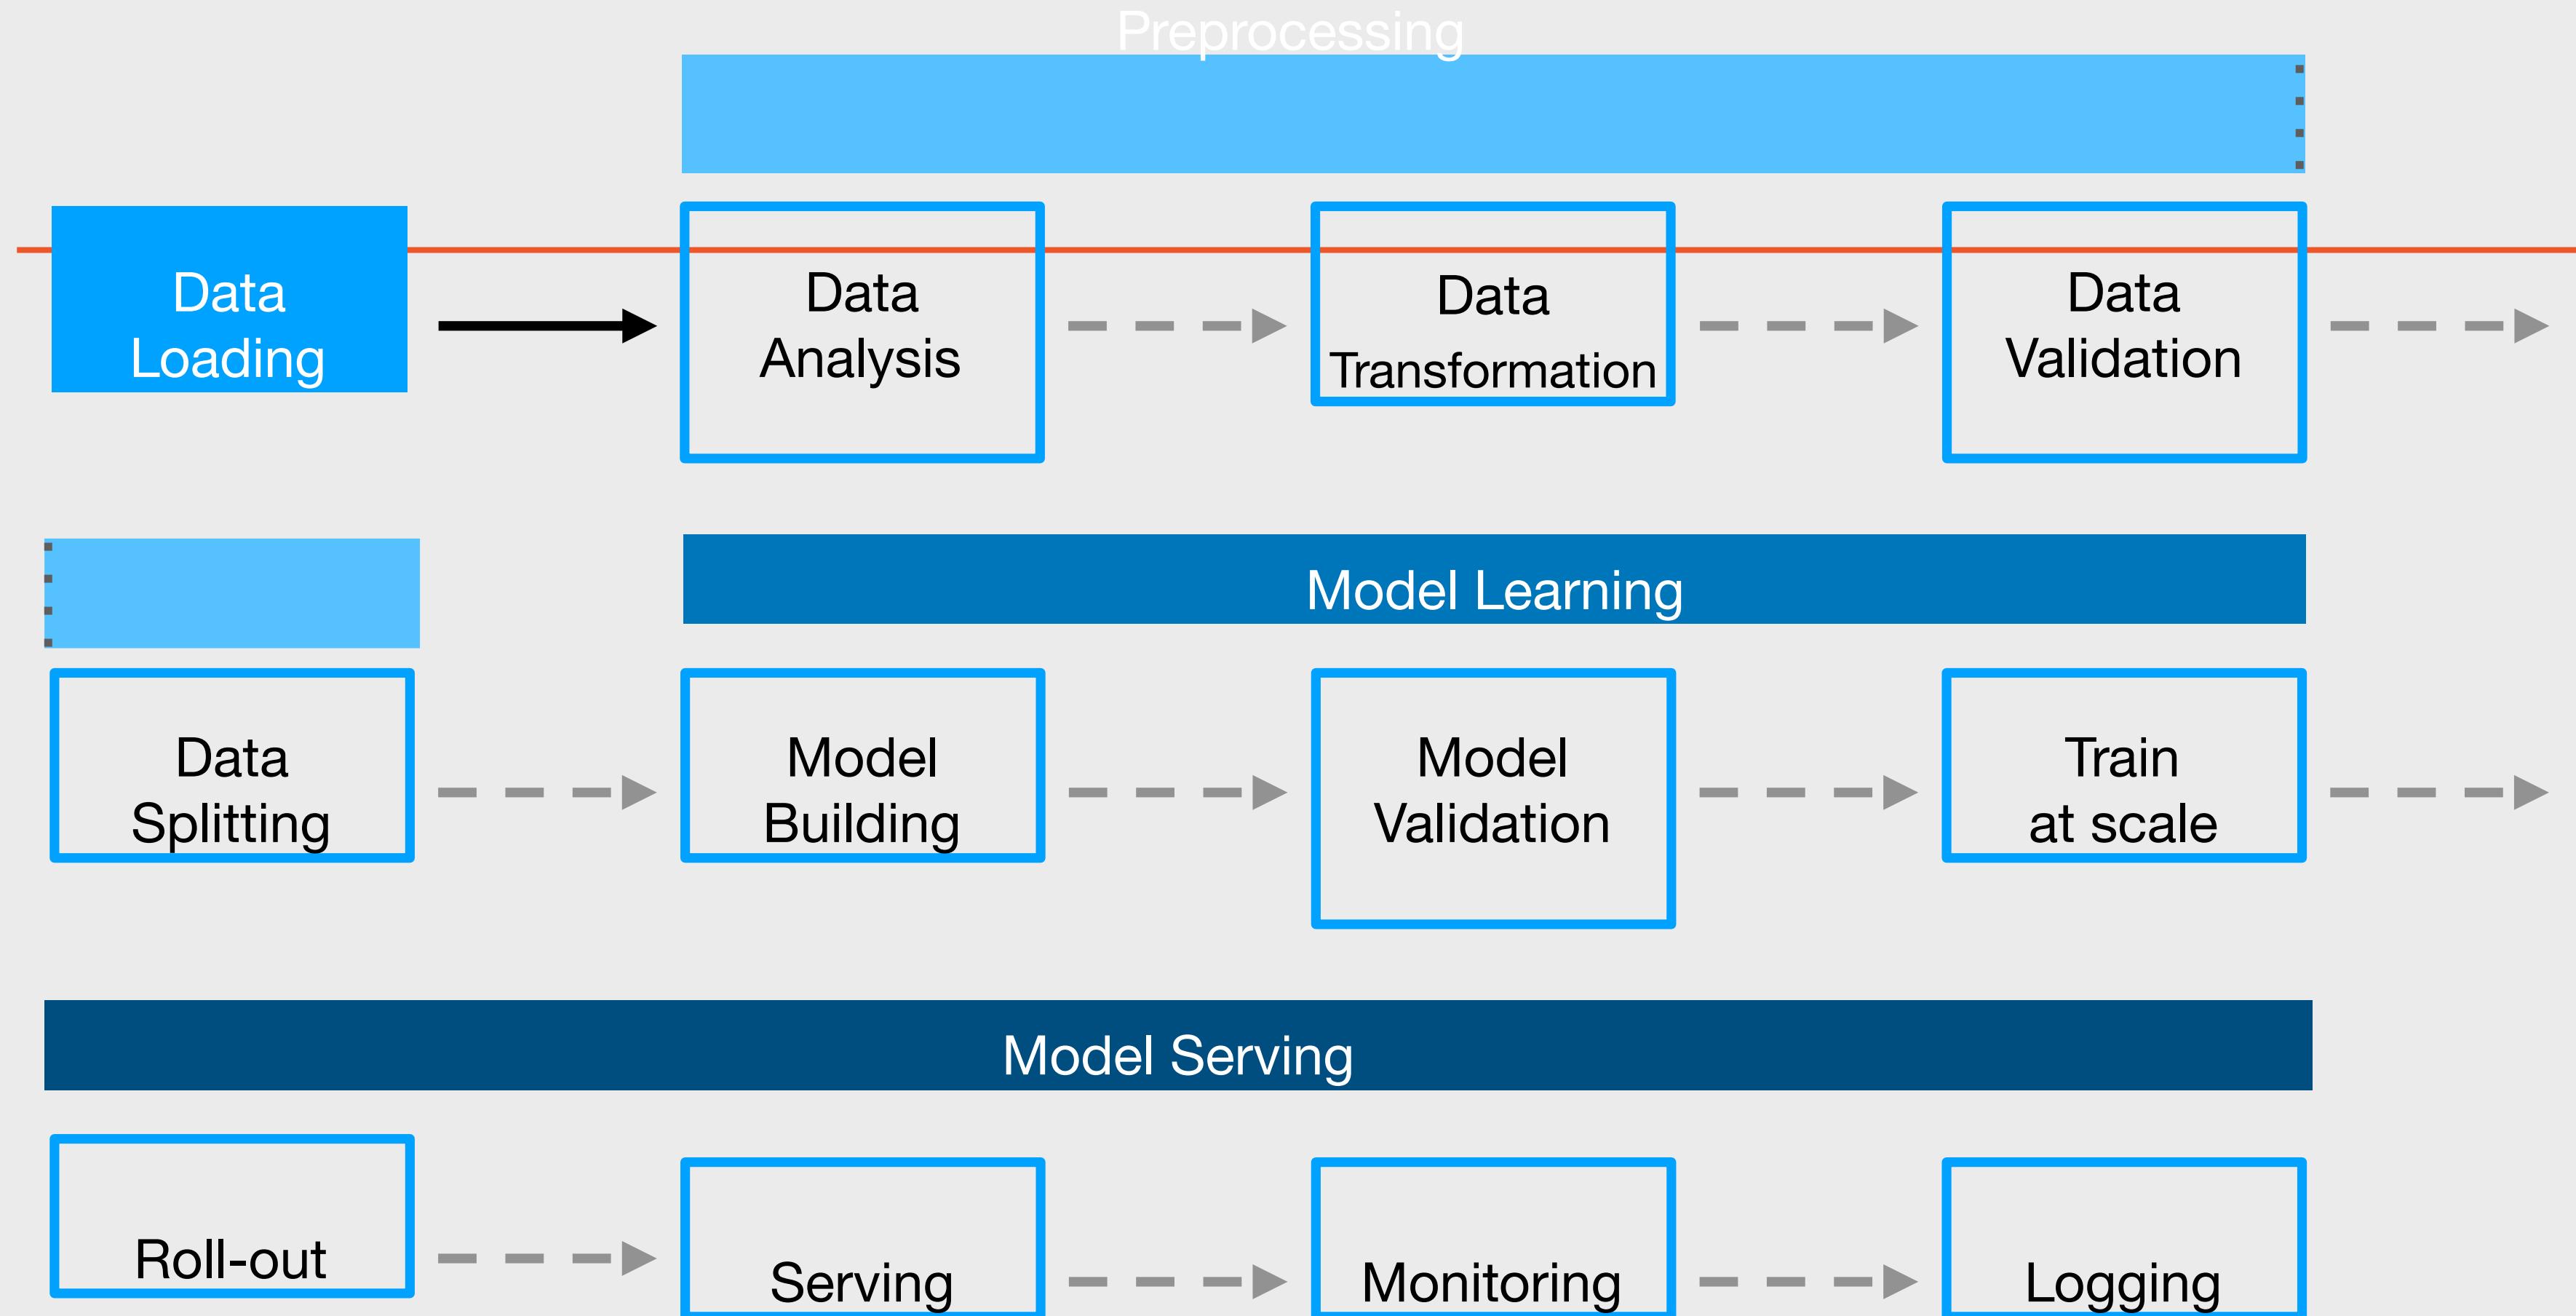
with a large team

Production Pipeline



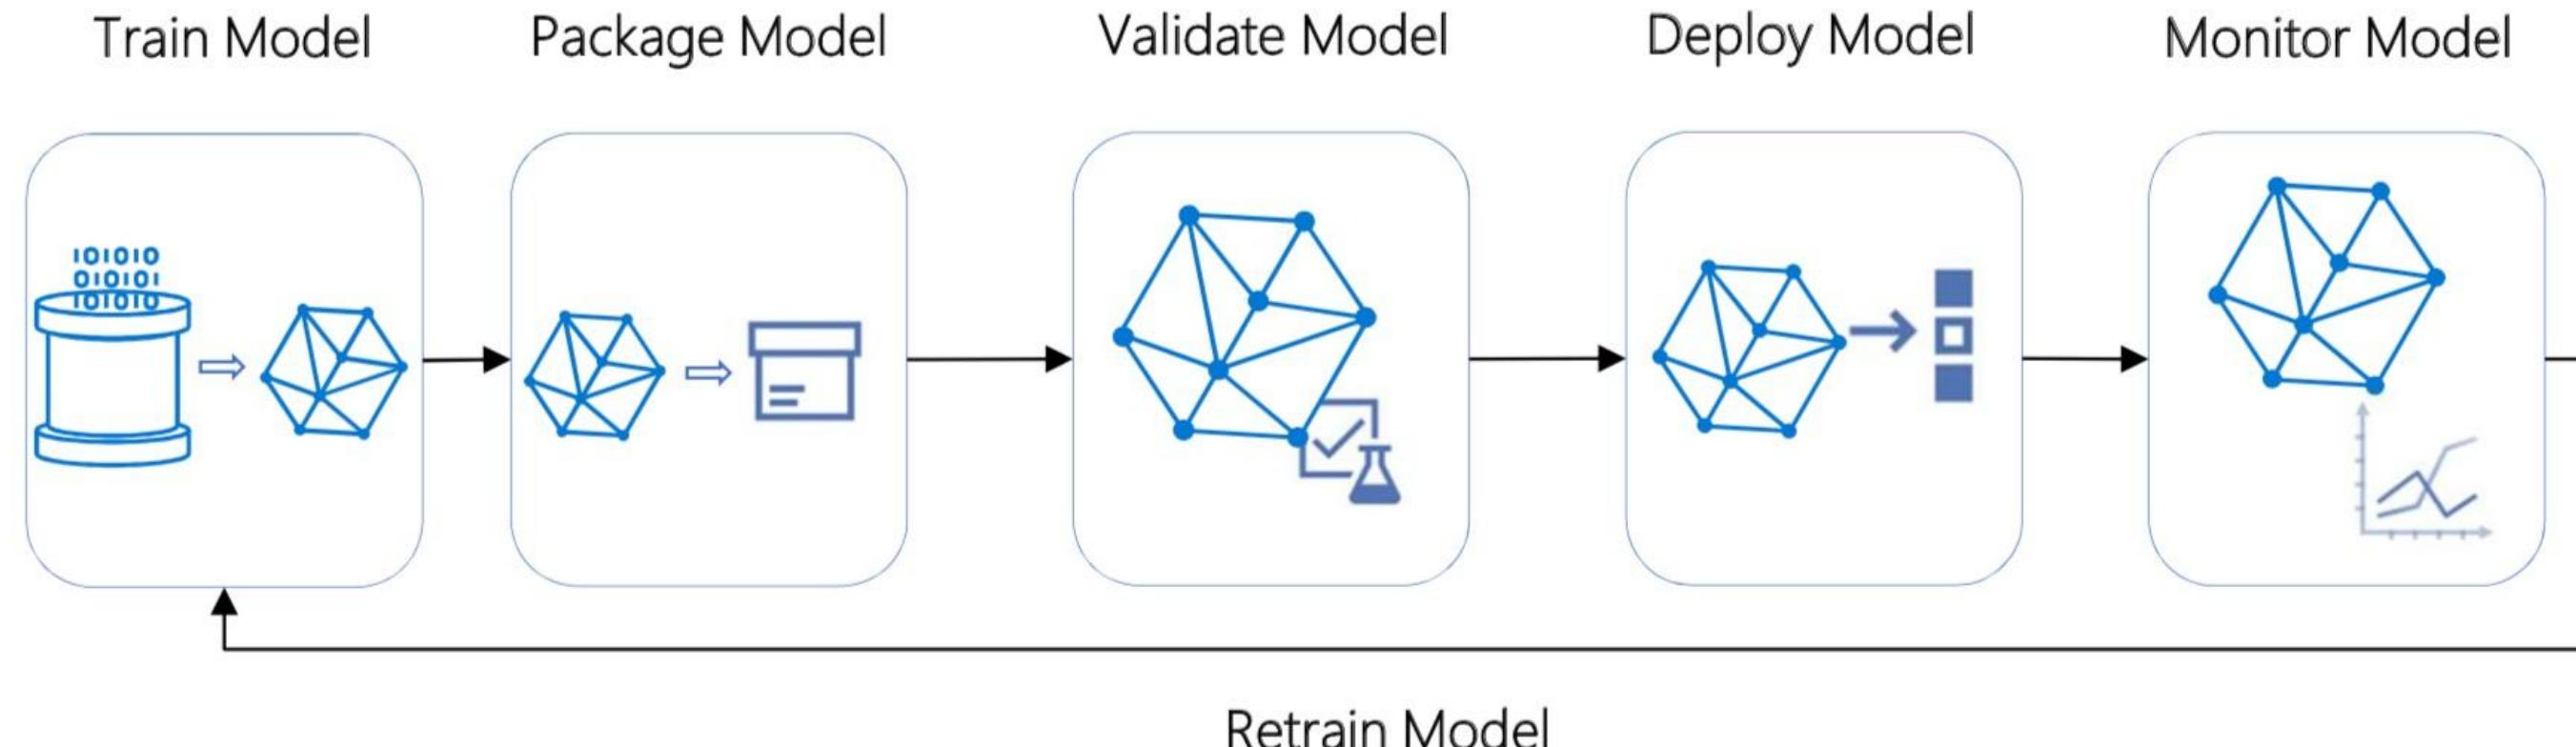
“toy” ML Pipeline



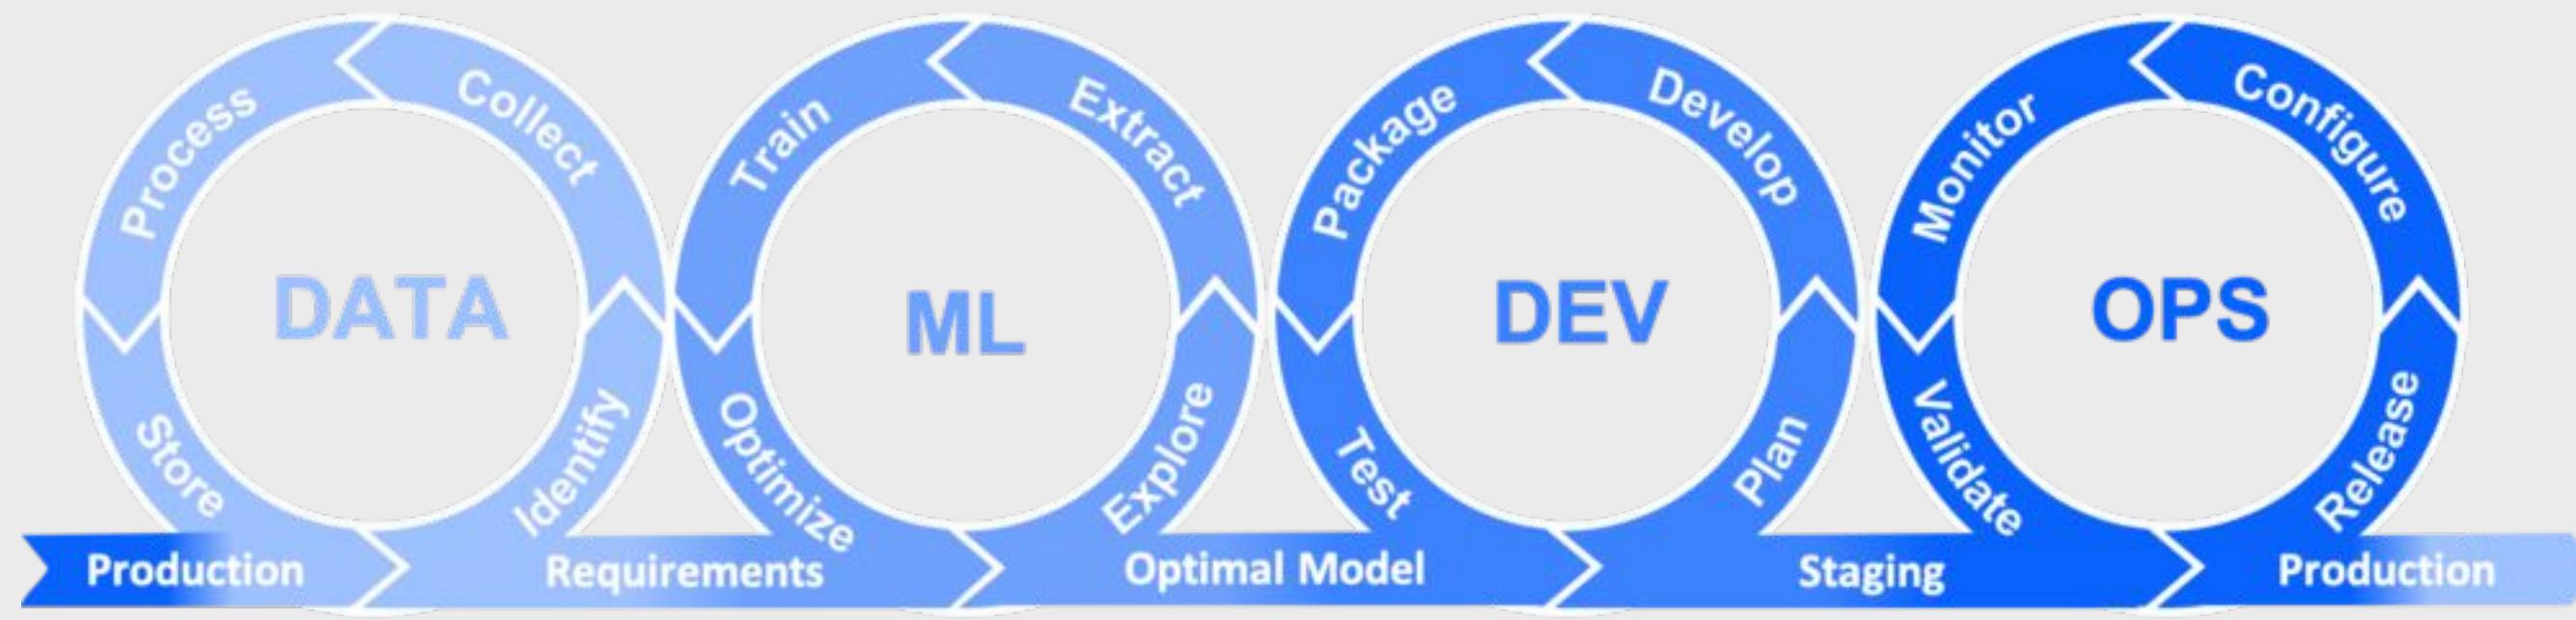


What is the E2E ML lifecycle?

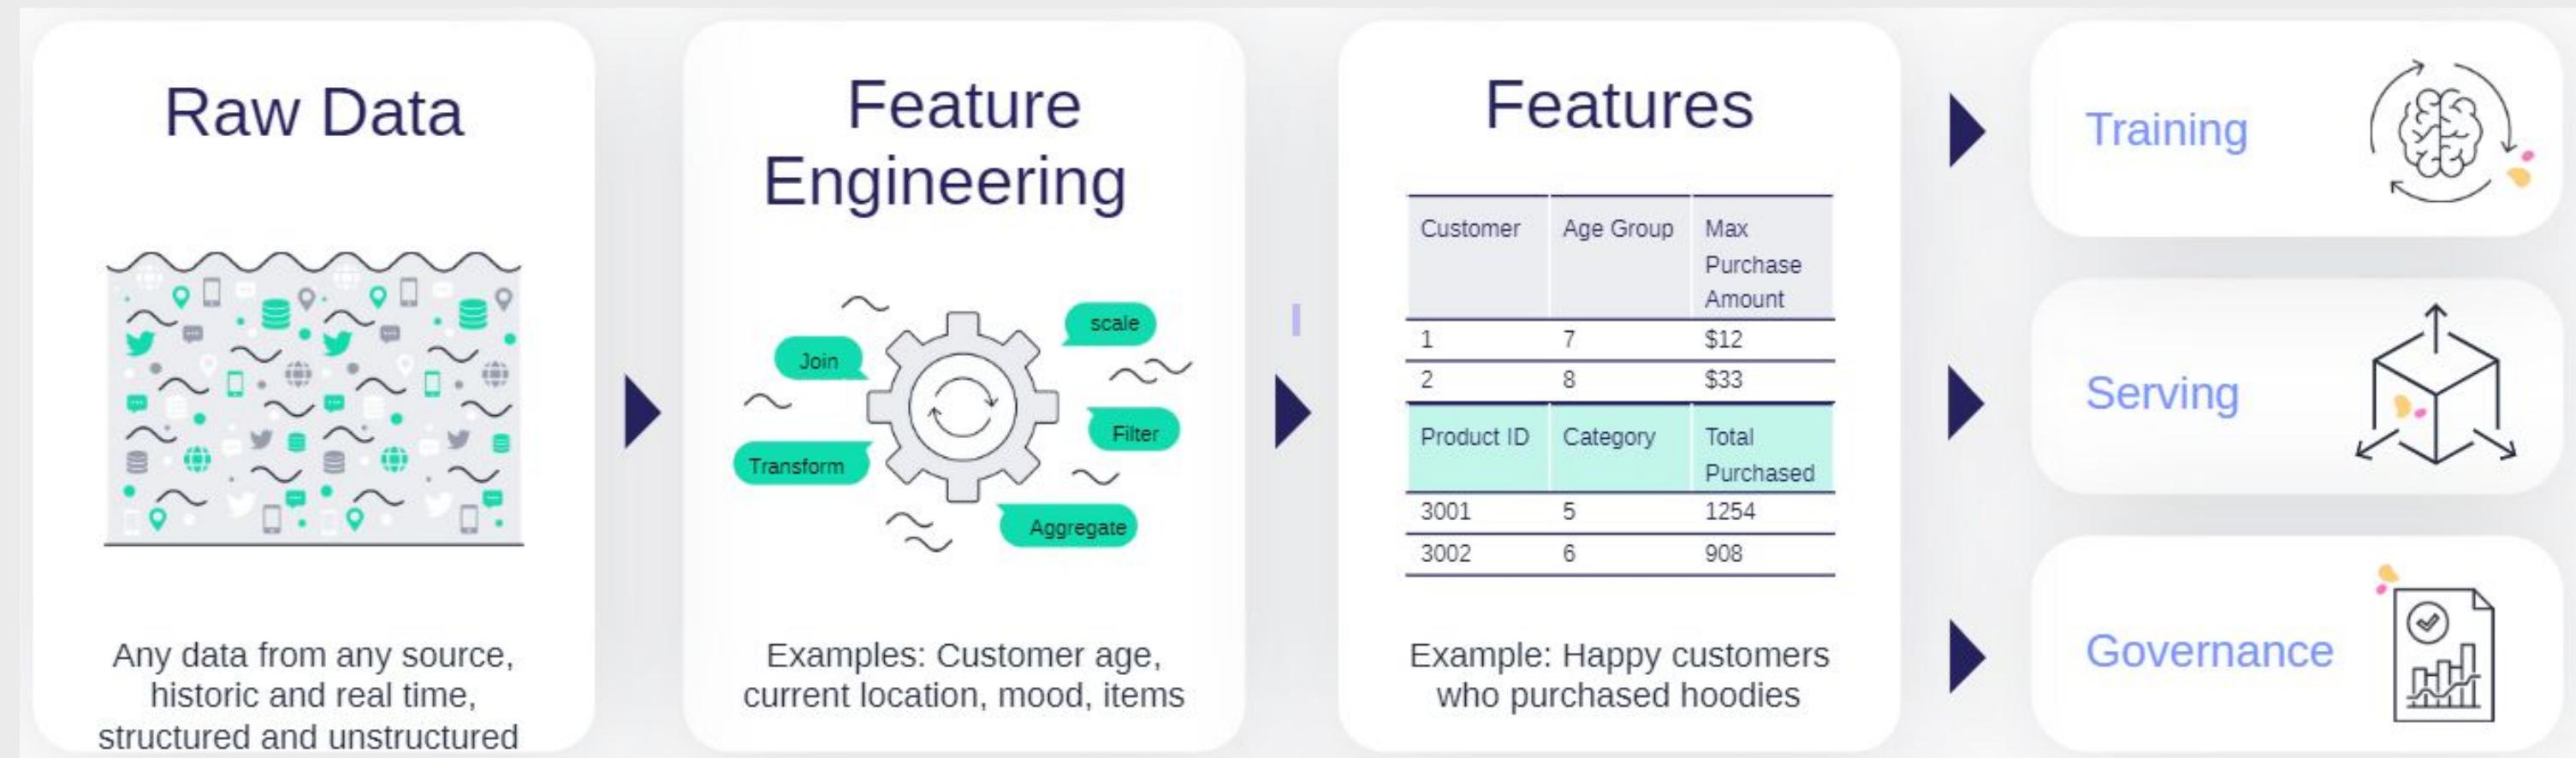
- **Develop & train model** with reusable ML pipelines
- **Package model** using containers to capture runtime dependencies for inference
- **Validate model behavior** – functionally, in terms of responsiveness, in terms of regulatory compliance
- **Deploy model** - to cloud & edge, for use in real-time / streaming / batch processing
- **Monitor model** behavior & business value, know **when to replace / deprecate a stale model**



Introduction



Introduction

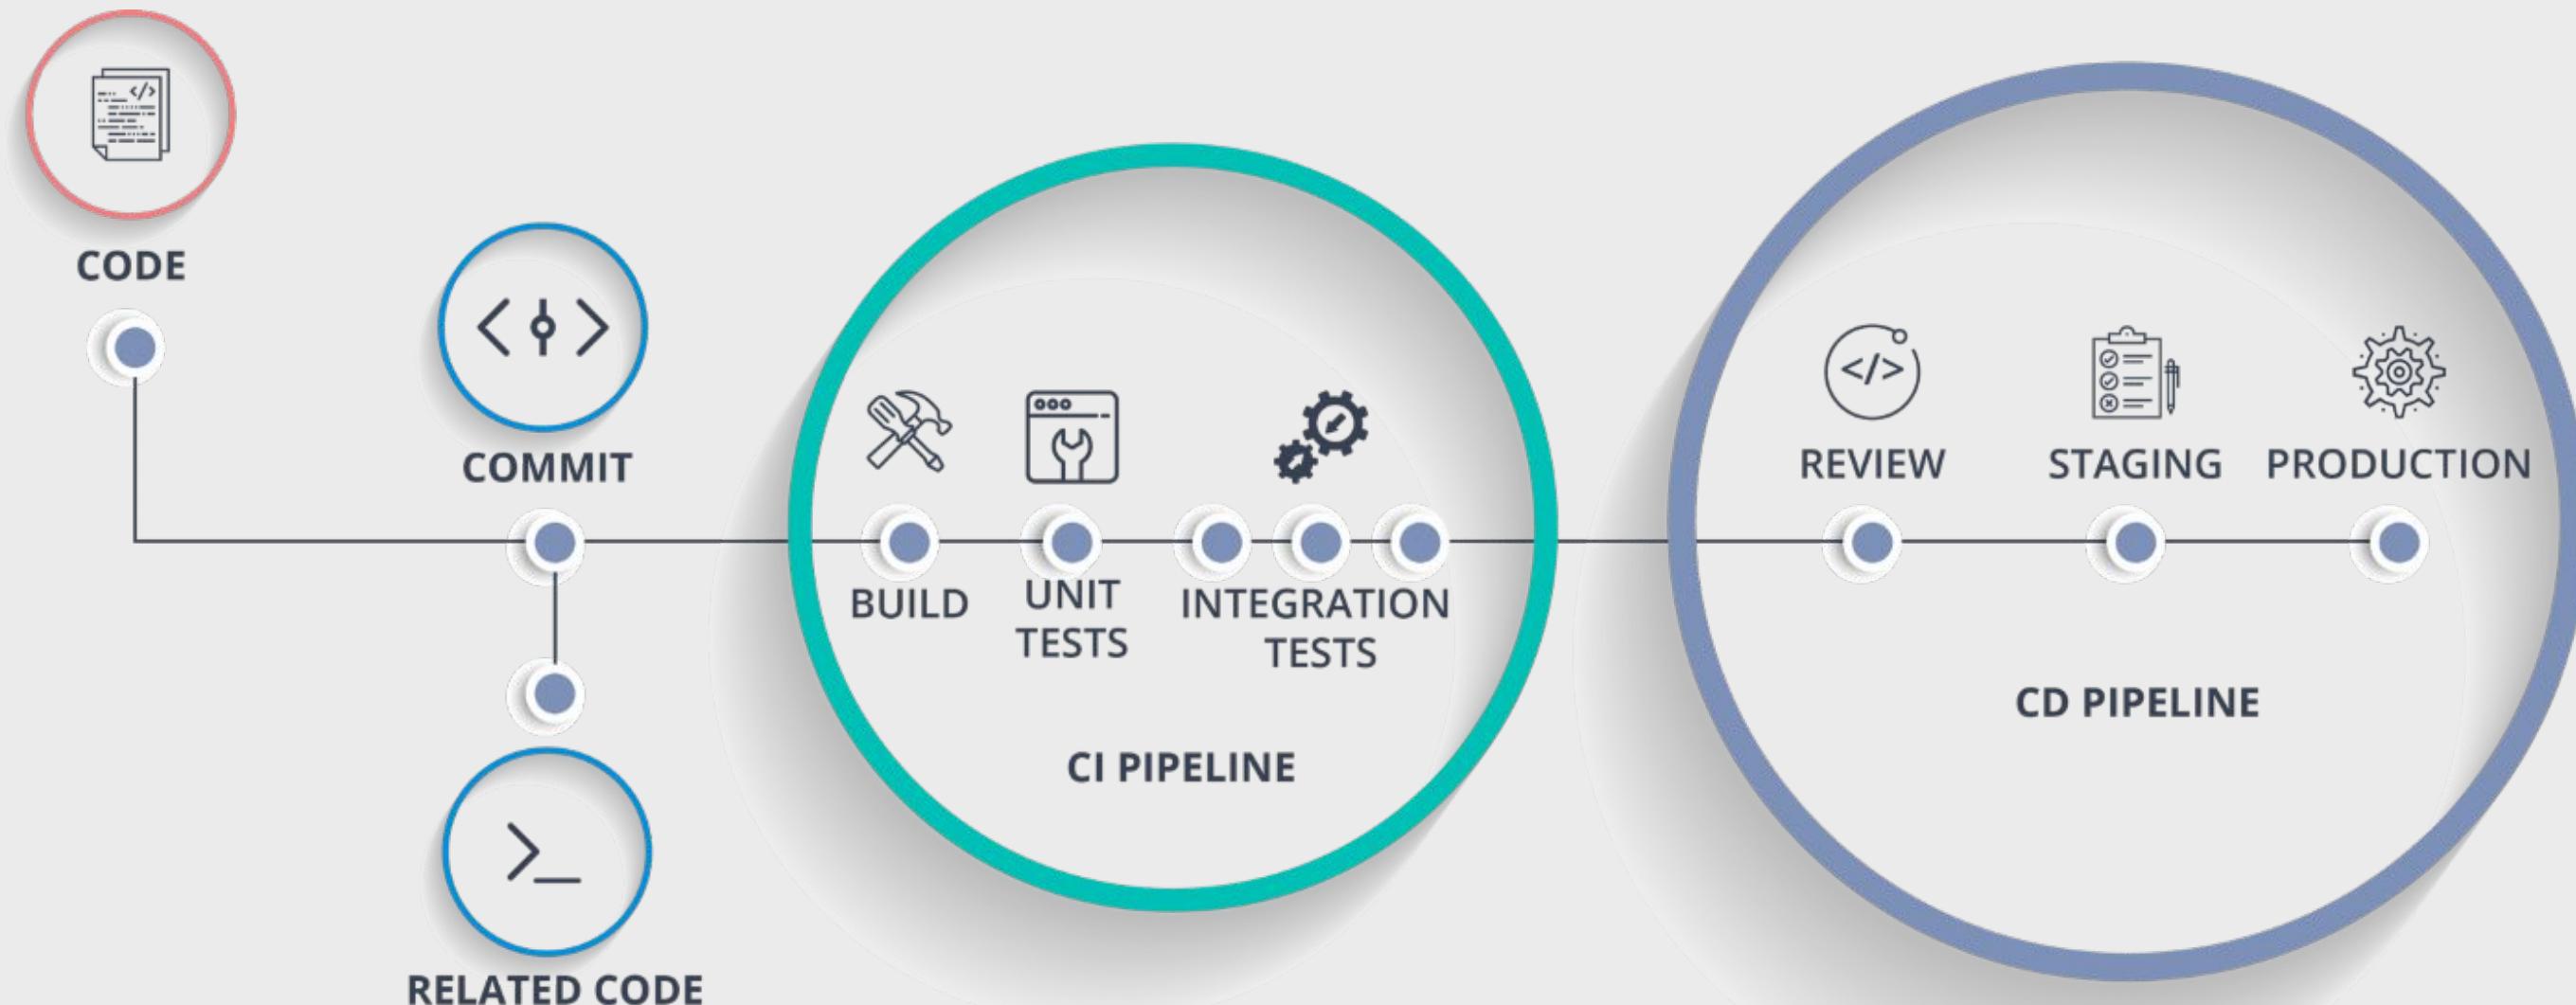


MLOps Challenges

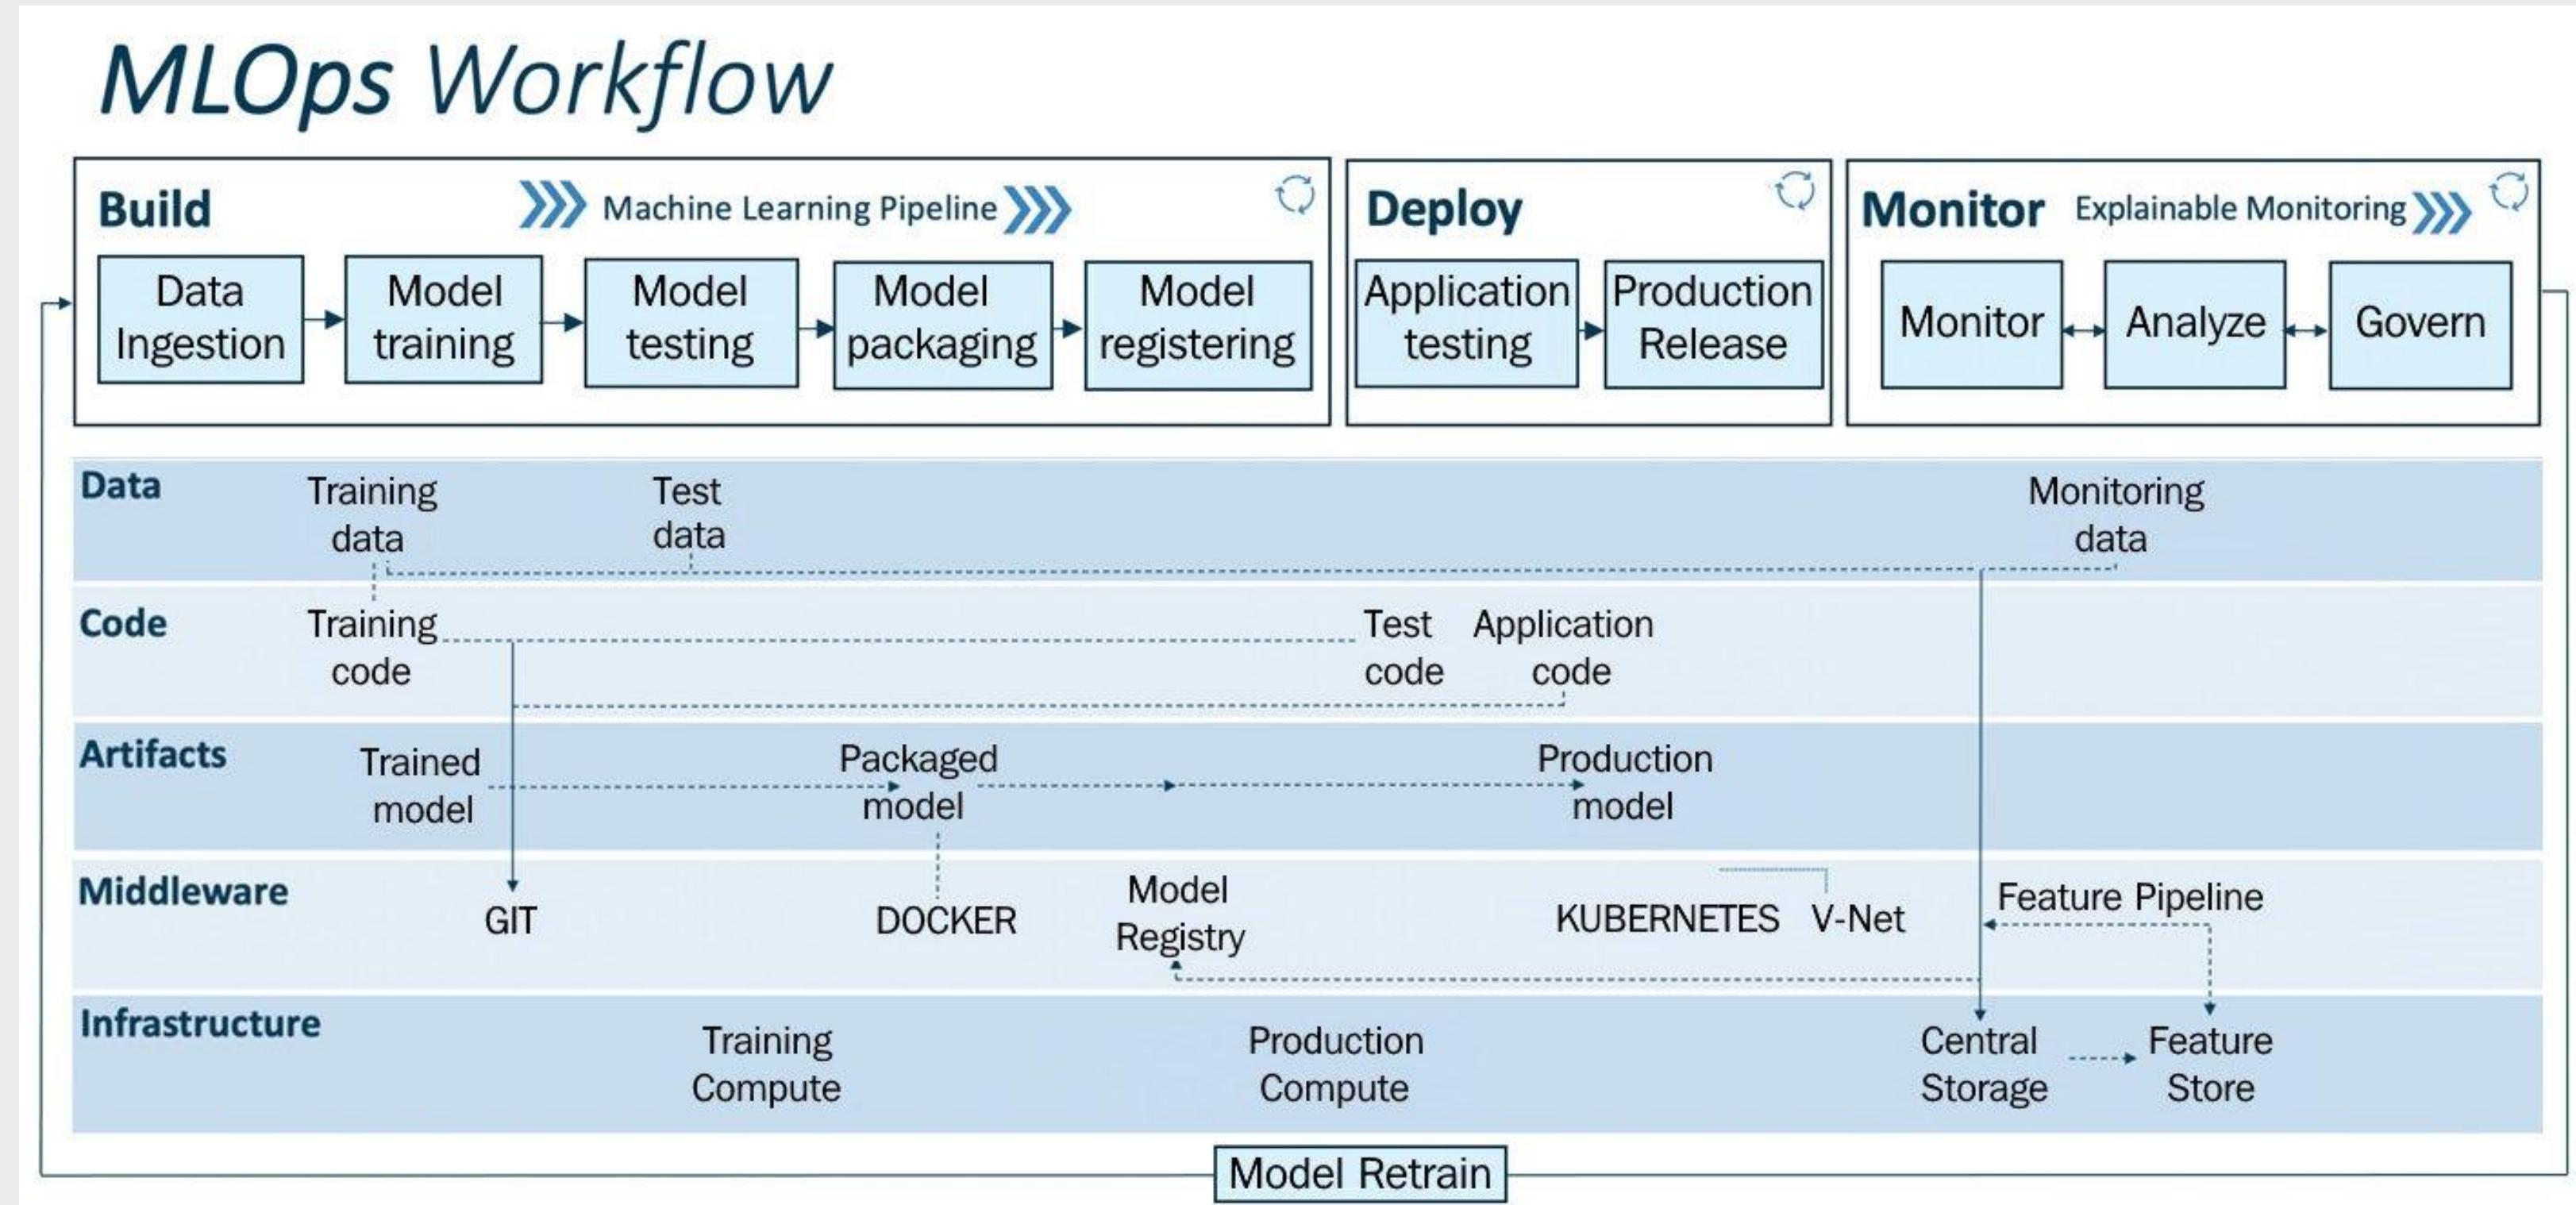
- **Process and model** vast amount of data
 - In **real time** and **batch**
- Data acquisition from **heterogeneous sources**
- Communication protocols **interoperability**
- **Security and privacy**
 - Authentication, authorization and encryption



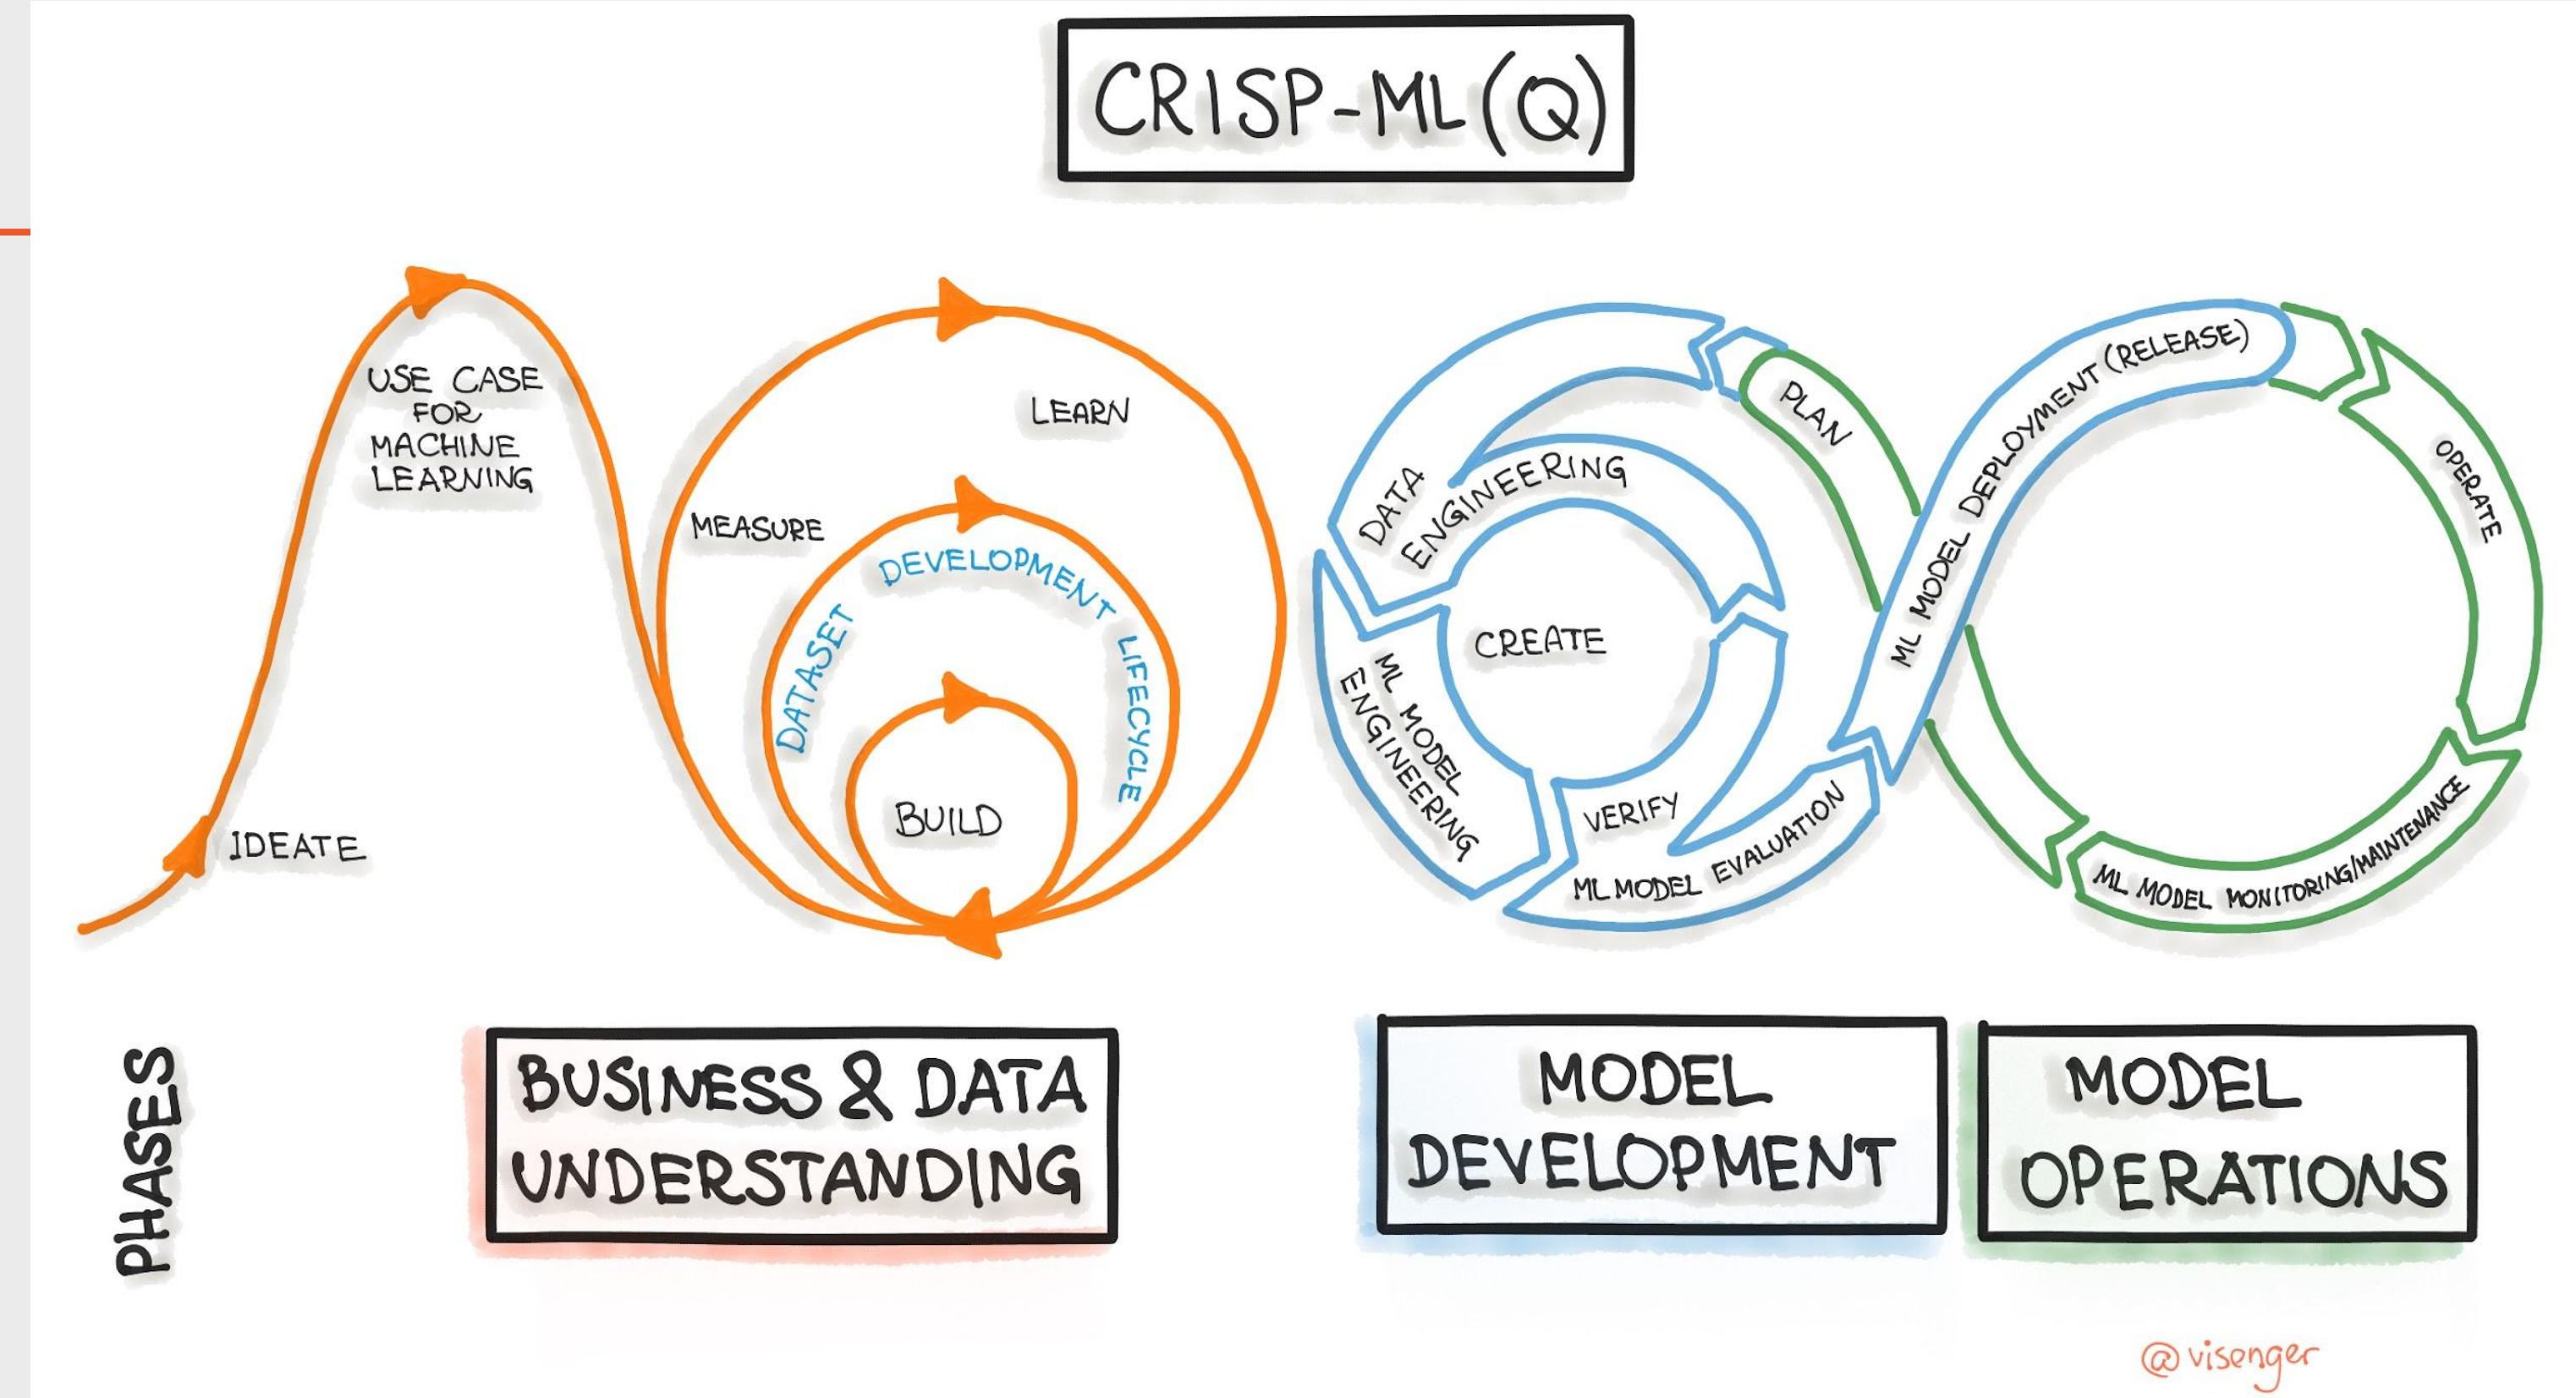
Introduction



Introduction to all the steps



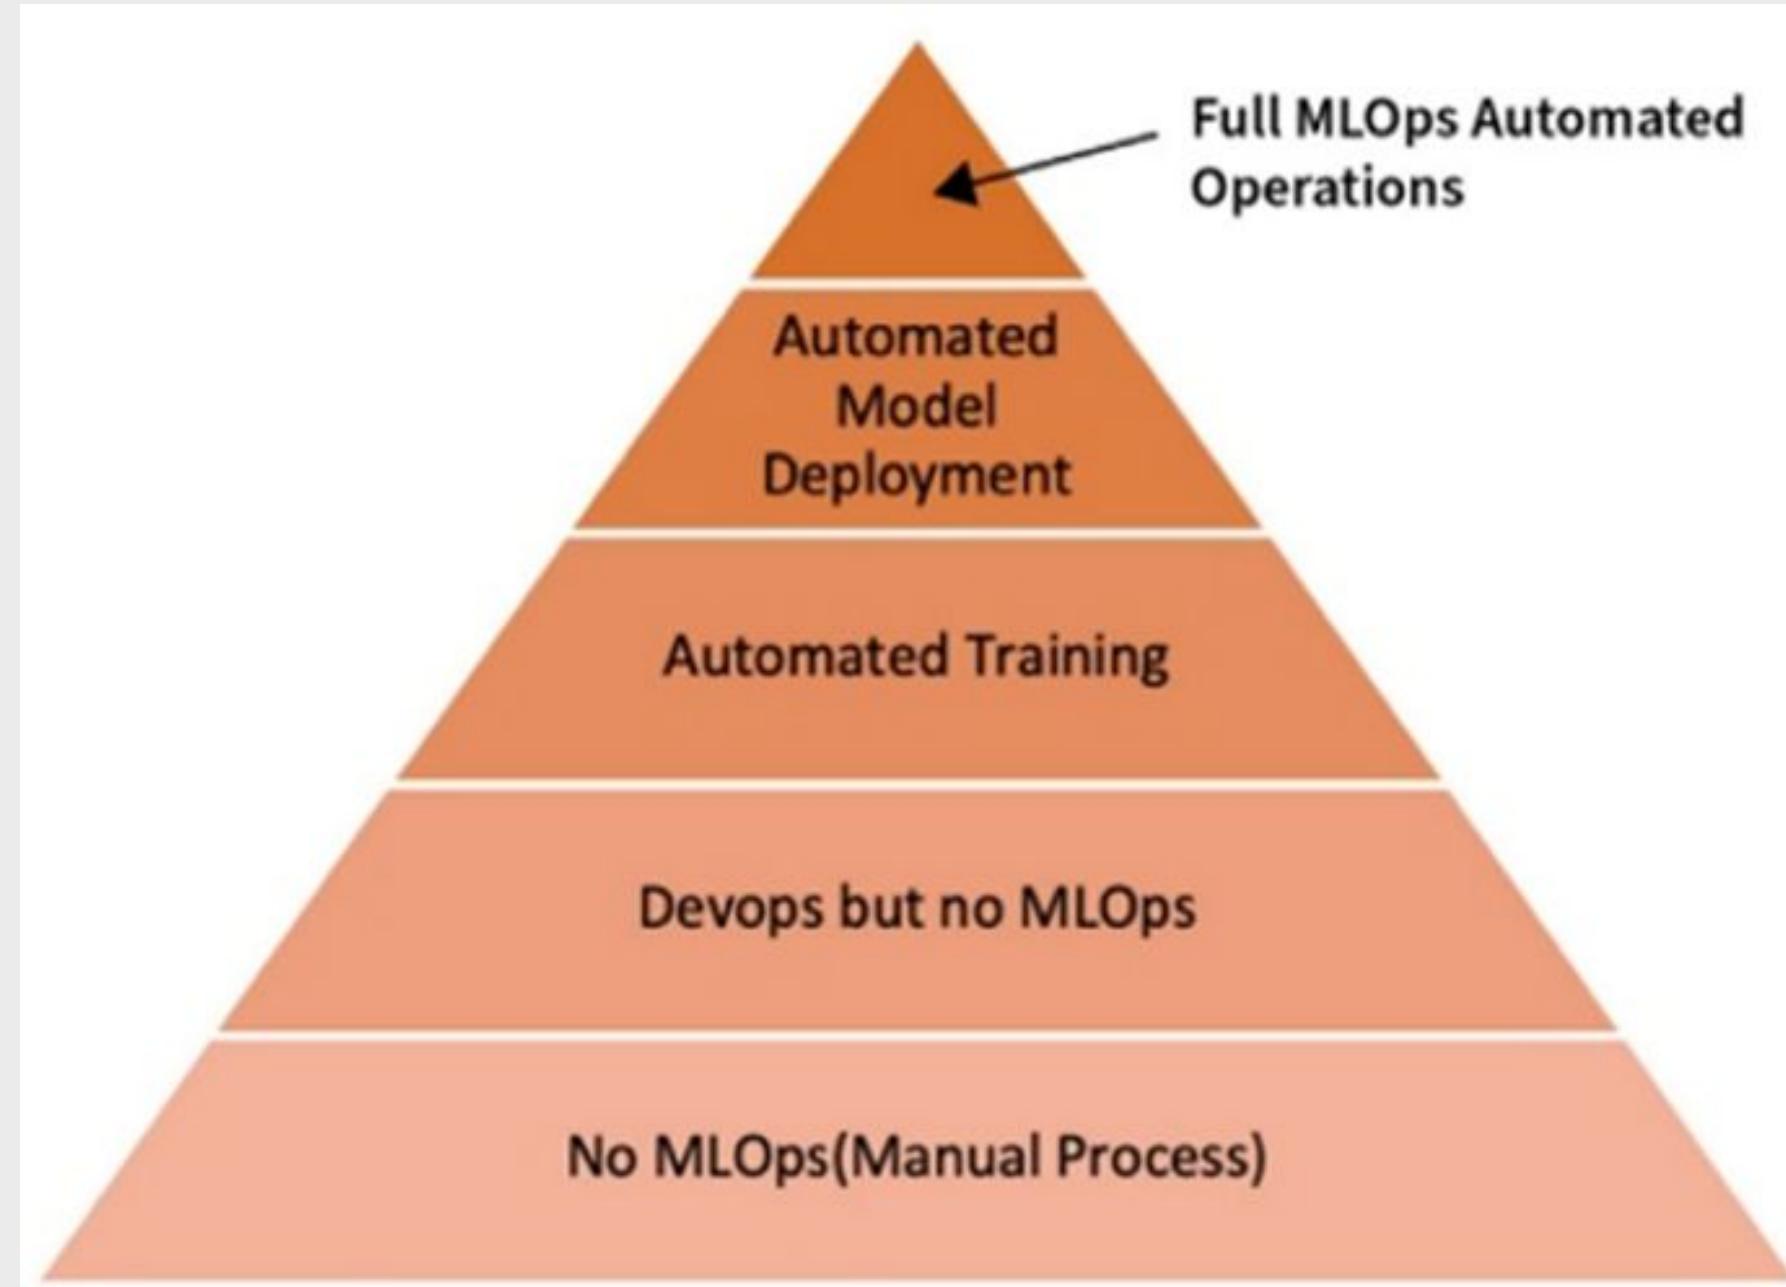
Crisp-ML



<https://arxiv.org/abs/2003.05155>

<https://arxiv.org/abs/2003.05155>

Maturity levels of MLOps



<https://learn.microsoft.com/en-us/azure/architecture/example-scenario/mlops/mlops-maturity-model>
<https://blog.onesaitplatform.com/en/2022/04/12/mlops-maturity-levels-and-open-source-tools/>

MLOps Challenges

Need for different tools that help along all the ML and data analysis workflow.

Fully integrated environment with NGSI-LD

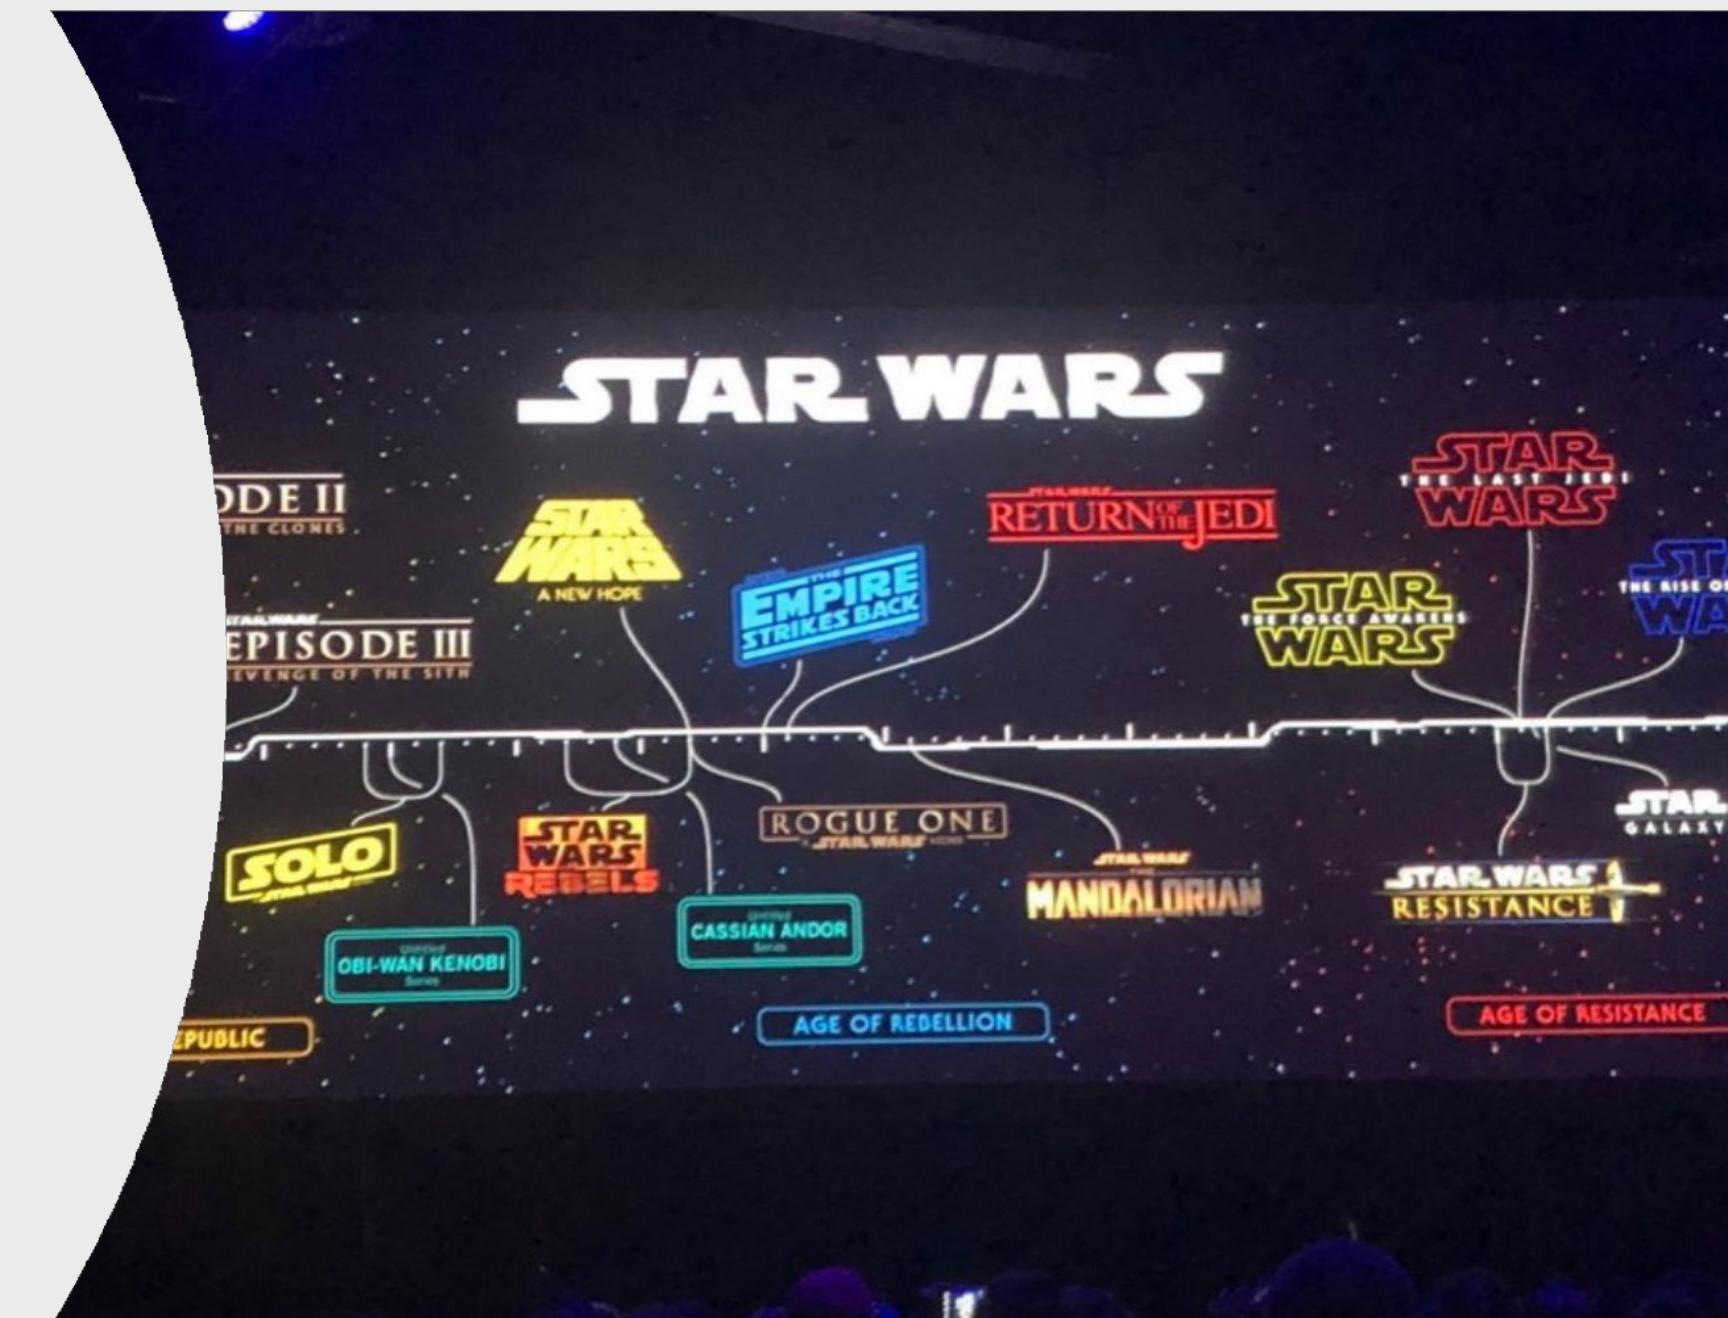
Some tools:

■ Apache Airflow

- Tool to create workflows (similar to NIFI interface) for different tools.
- Google support.

■ MLFlow

- Created by the same authors for Apache Spark.
- Takes care of the full lifecycle for ML models.



Platforms



Prefect and Apache linkis

Alternatives (Orchestrators)

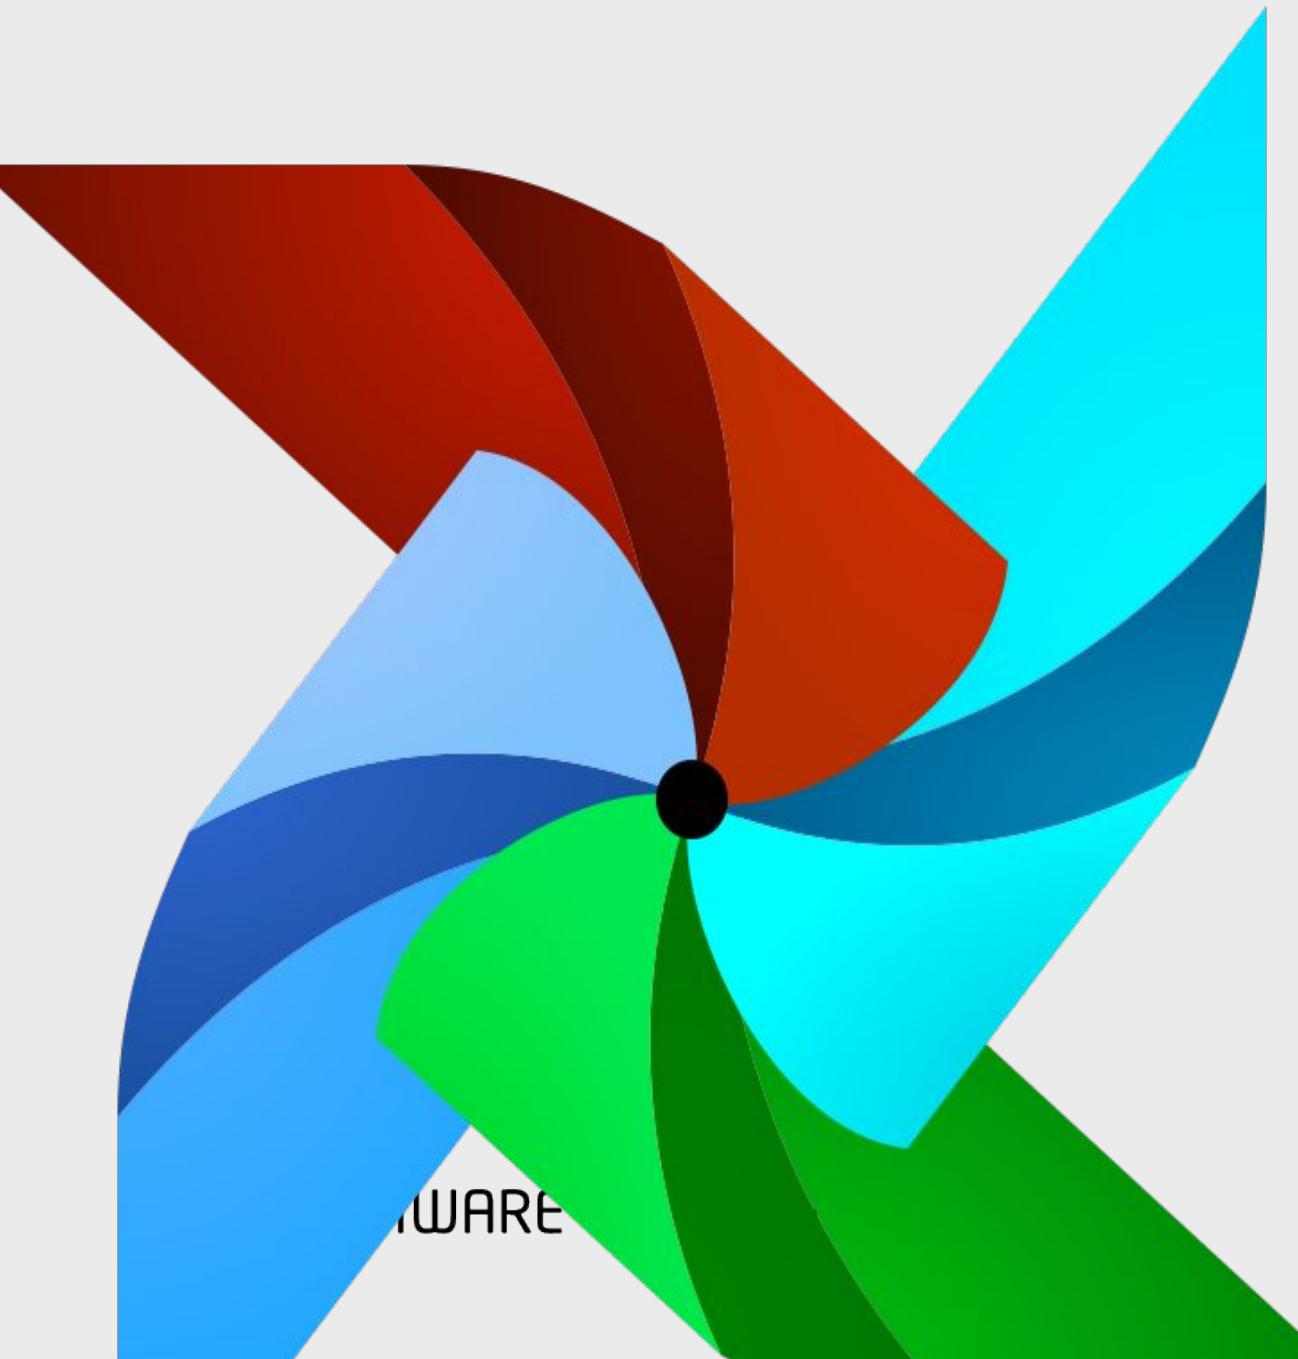


For DS &
ML



Apache Airflow

Vienna, 12-13 June, 2023 | #FIWARESummit



Origin

- Created by Airbnb in October 2014.
- Airflow is the industry standard open-source tool for programmatic workflow orchestration.
- Try to solve the company's increasingly complex workflows.
- Written in Python, and workflows are created via Python scripts.
- Is designed under the principle of "configuration as code".
- Is open source from the start :
 - <https://github.com/apache/airflow>
- Complements other workflow tools:
 - Luigi (<https://github.com/spotify/luigi>) (created at Spotify)
 - Apache Oozie (<https://oozie.apache.org/>)
 - Azkaban (<https://azkaban.github.io>) (created at LinkedIn)

Use cases

- ETL pipelines
- Machine learning pipelines
- Predictive data pipelines: fraud detection, scoring/ ranking, classification, recommender system, etc.
- General job scheduling: DB back-ups
- Anything... automate the garage door?

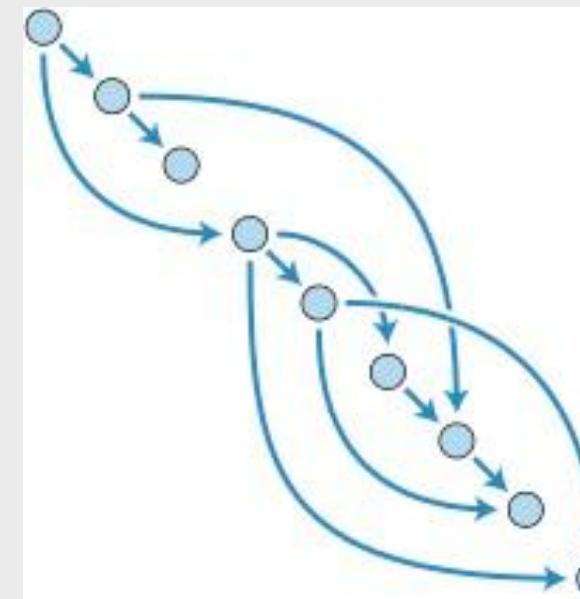
Airflow uses *Operators* as the fundamental unit of abstraction to define tasks, and uses a **DAG** to define workflows using a set of operators

MAD · NOV 23-24 · 2018

DAGS

- At the heart of the tool is the concept of a DAG (Directed Acyclic Graph). A DAG is a series of tasks that you want to run as part of your workflow. This might include something like extracting data via a SQL query, performing some calculations with Python and then loading the transformed data into a new table. In Airflow each of these steps would be written as individual tasks in a DAG.
- Airflow enables you to also specify the relationship between the tasks, any dependencies (e.g. data having loaded in a table before a task is run) and the order in which the tasks should be run.
- A DAG is written in Python and saved as a .py file. The DAG_ID is used extensively by the tool to orchestrate the running of the DAG's.

What is a DAG



- **Directed Acyclic Graph**
- Represents a workflow: set of tasks with a dependency structure
- Each node represents some form of



UI Screens – DAGs List

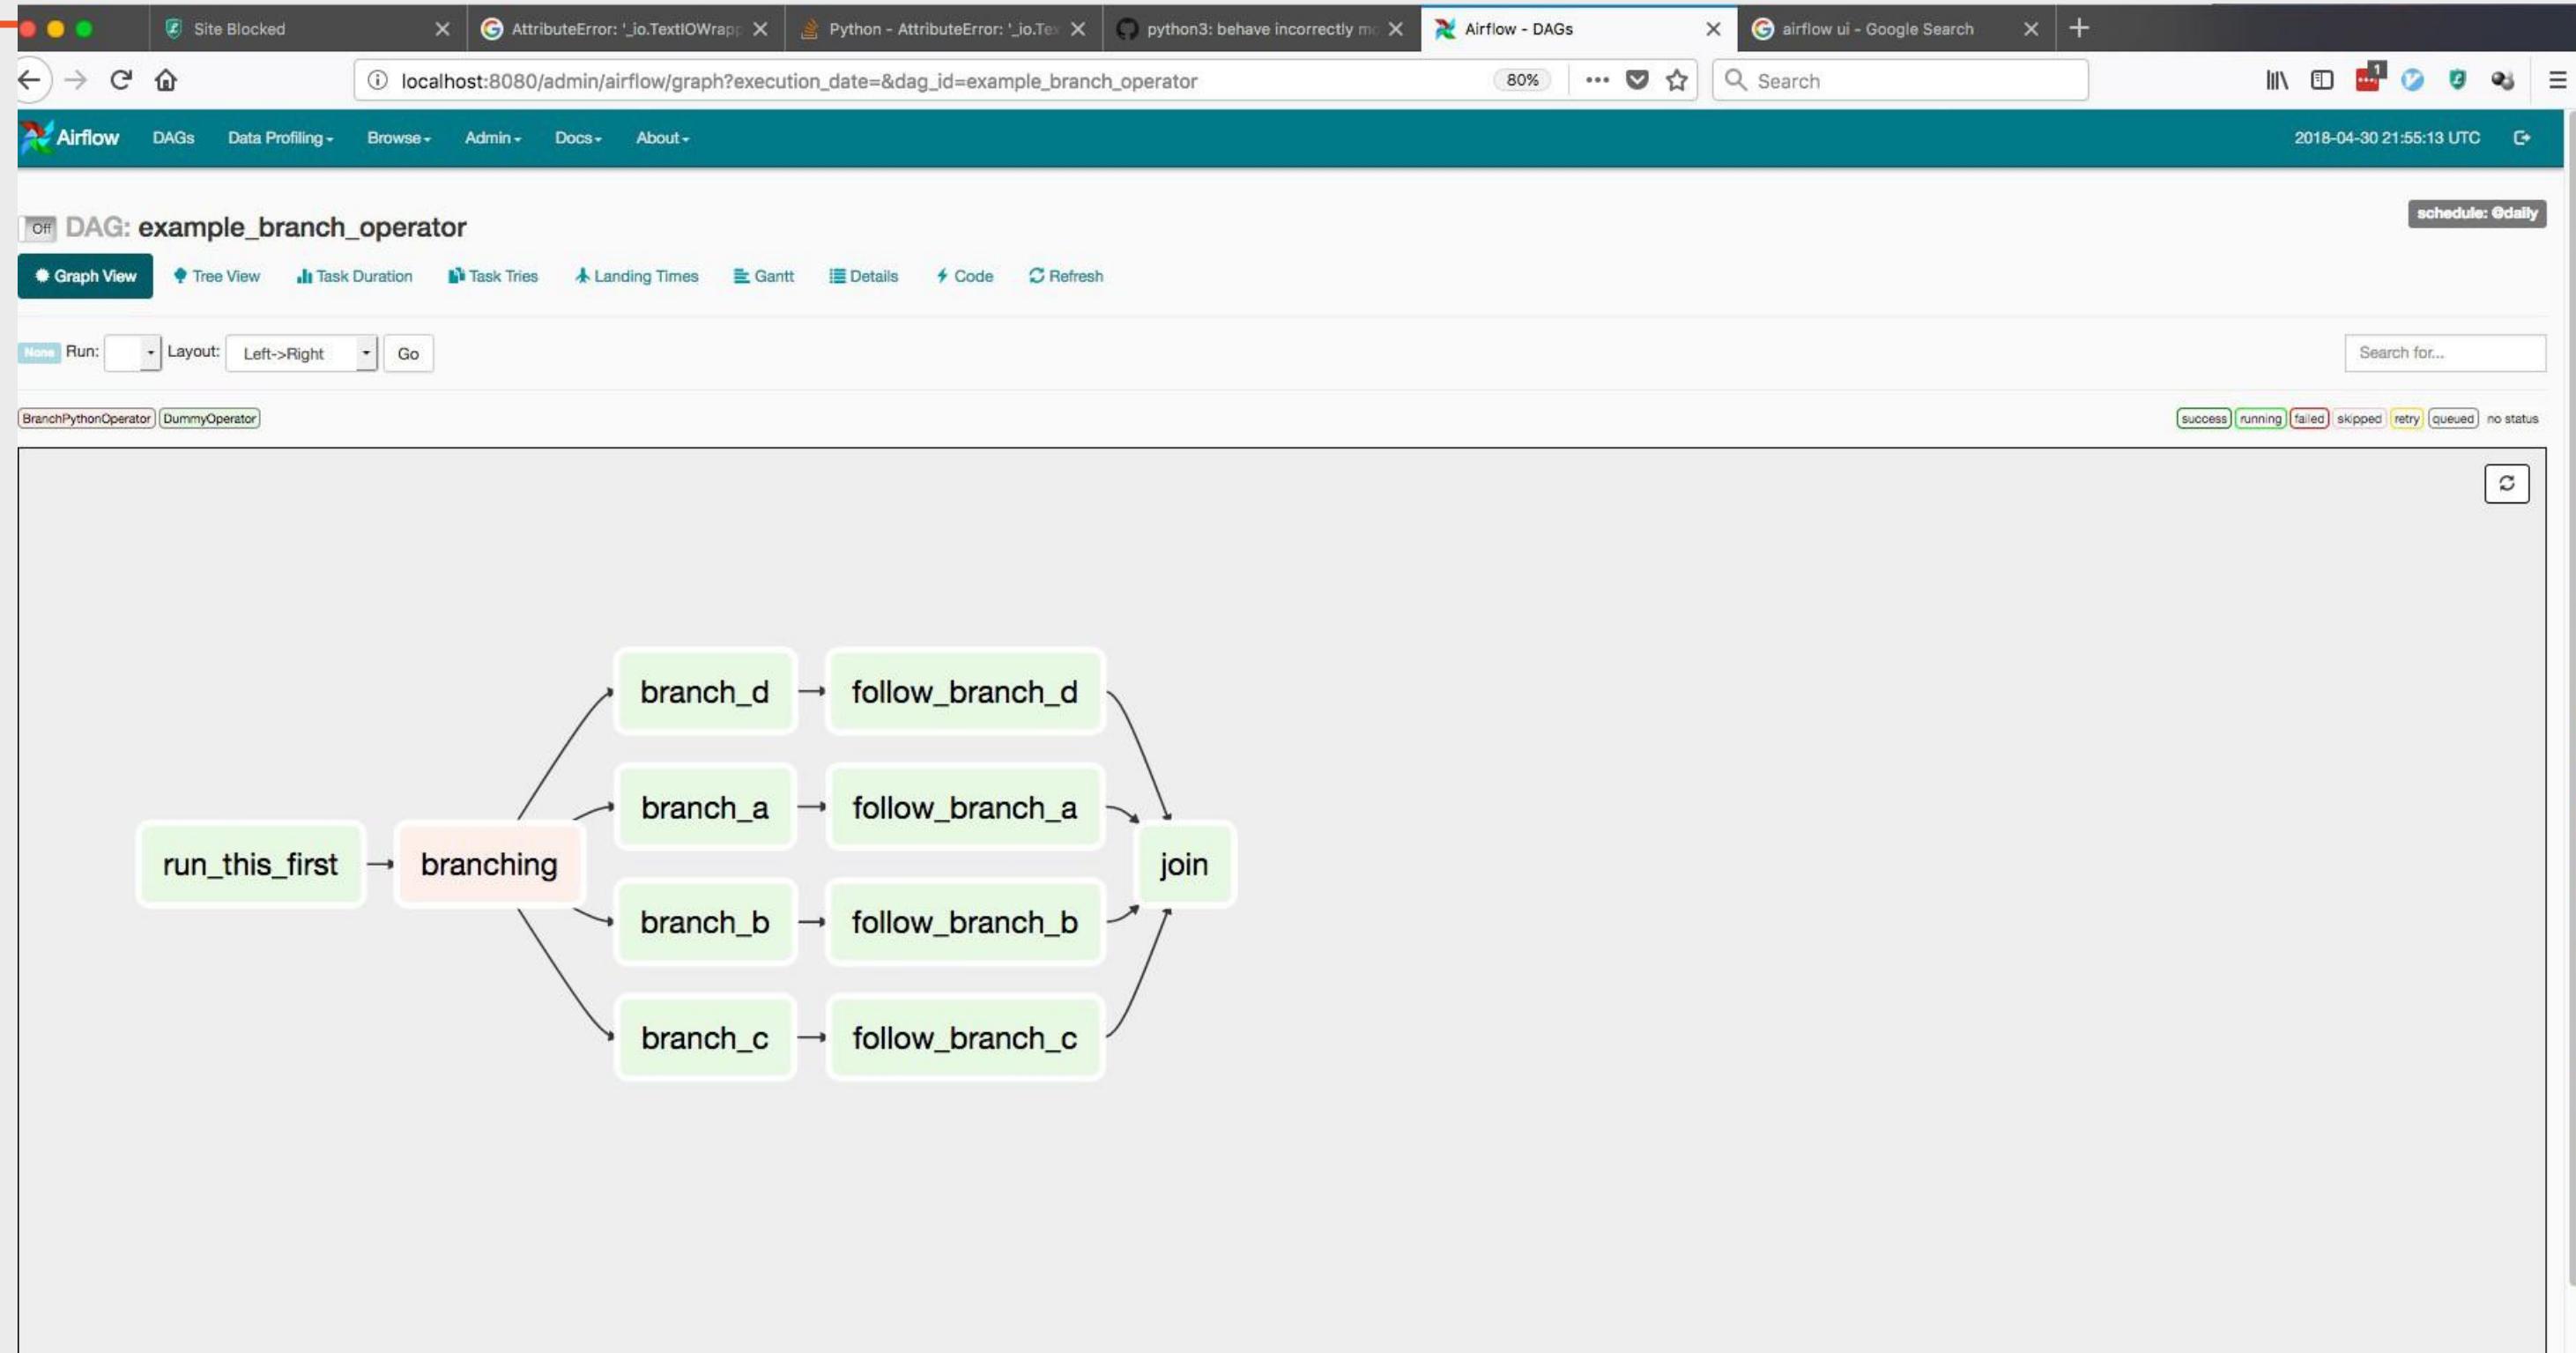
A screenshot of the Airflow web interface showing the DAGs list. The top navigation bar includes links for Airflow, DAGs, Data Profiling, Browse, Admin, Docs, and About, along with a timestamp of 2018-09-22 10:48:00 UTC and a refresh button.

DAGs

		DAG	Schedule	Owner	Recent Tasks	Last Run	DAG Runs	Links
		example_bash_operator	0 0 * * *	airflow				
		example_branch_dop_operator_v3	* /1 * * * *	airflow				
		example_branch_operator	@daily	airflow				
		example_http_operator	1 day, 0:00:00	airflow				
		example_kubernetes_executor	None	airflow				
		example_passing_params_via_test_command	* /1 * * * *	airflow				
		example_python_operator	None	airflow				
		example_short_circuit_operator	1 day, 0:00:00	airflow				
		example_skip_dag	1 day, 0:00:00	airflow				



UI Screens – DAGs View





UI Screens – DAGs Runs, Tree View

Airflow DAGs Data Profiling ▾ Browse ▾ Admin ▾ Docs ▾ About ▾ 2018-09-07 22:15:40 UTC ⌂

On DAG: example_branch_dop_operator_v3 schedule: */1 * * * *

Graph View Tree View Task Duration Task Tries Landing Times Gantt Details Code Refresh Delete

Base date: 2018-09-05 01:04:00 Number of runs: 25 Go

BranchPythonOperator DummyOperator success running failed skipped retry queued no status

[DAG] oper_1 condition oper_2 condition



DAG Code Example

```
import uuid
from datetime import datetime from
airflow import DAG
from airflow.utils.trigger_rule import TriggerRule
from airflow.operators.postgres_operator import PostgresOperator

dag_params = {
    'dag_id': 'PostgresOperator_dag',
    'start_date': datetime(2019, 10, 7),
    'schedule_interval': None
}

with DAG(**dag_params) as dag:

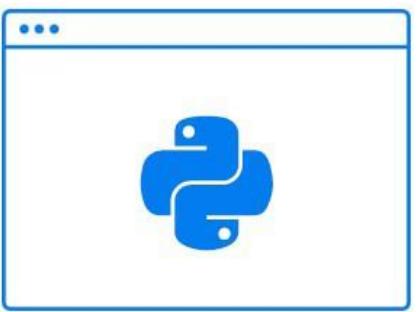
    create_table = PostgresOperator(
        task_id='create_table',
        sql="""CREATE TABLE new_table(
            custom_id integer NOT NULL, timestamp TIMESTAMP NOT NULL, user_id VARCHAR (50) NOT NULL
        );""",
    )

    insert_row = PostgresOperator(
        task_id='insert_row',
        sql='INSERT INTO new_table VALUES(%s, %s, %s)',
        trigger_rule=TriggerRule.ALL_DONE,
        parameters=(uuid.uuid4().int % 123456789, datetime.now(), uuid.uuid4().hex[:10])
    )

    create_table >> insert_row
```

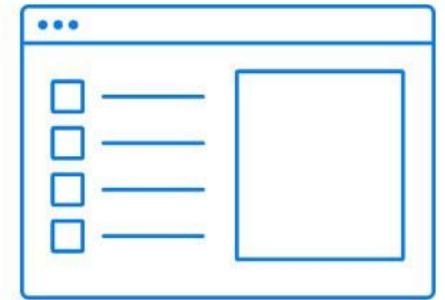


Main Features



Pure Python

No more command-line or XML black-magic! Use standard Python features to create your workflows, including date time formats for scheduling and loops to dynamically generate tasks. This allows you to maintain full flexibility when building your workflows.



Easy to Use

Anyone with Python knowledge can deploy a workflow. Apache Airflow does not limit the scope of your pipelines; you can use it to build ML models, transfer data, manage your infrastructure, and more.



Open Source

Wherever you want to share your improvement you can do this by opening a PR. It's simple as that, no barriers, no prolonged procedures. Airflow has many active users who willingly share their experiences. Have any questions? Check out our buzzing slack.



Apache Airflow Community

<https://github.com/apache/airflow>

Official community Slack:

<https://apache-airflow-slack.herokuapp.com/>

List of committers (maintainers):

<https://people.apache.org/committers-by-project.html#airflow> (about 40 people)

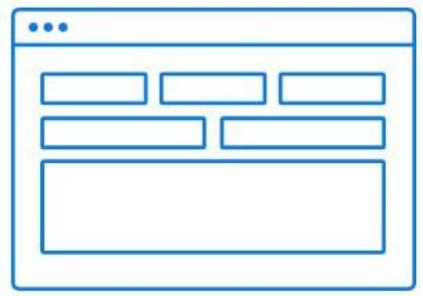
The screenshot shows the GitHub repository page for 'apache / airflow'. The 'Code' tab is selected. On the left, there's a 'Contributors' section with a count of 1,395, showing profile icons for several contributors. Below it is a 'XD' icon. The main area displays a list of recent commits from the 'master' branch:

Author	Commit Message	Time Ago
jhtimmins	Refactor and speed up "DAG:" prefix permission...	9 hours ago
	Improve verification of images with PIP check (#1...	2 days ago
	Refactor and speed up "DAG:" prefix permissions...	9 hours ago
	Fix chart jobs delete policy for improved idempot...	yesterday
	Enable Markdownlint rule MD003/heading-style/...	15 days ago
	Enable Black - Python Auto Formatter (#9550)	29 days ago
	User-friendly output of Breeze and CI scripts (#1...	2 days ago
	Fix typo in docker-context-files/README.md (#1...	29 days ago
	Allow using _CMD / _SECRET to set '[webserver]...	15 hours ago
	Prepare release candidate for backport package...	7 months ago

On the right side of the repository page, there are sections for 'About', 'Readme', 'Apache-2.0 License', and 'Releases'.

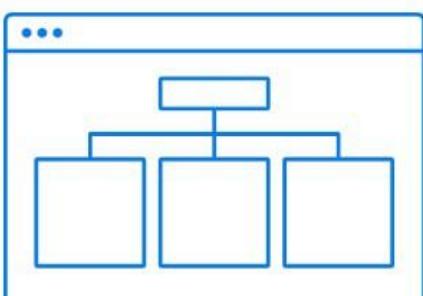


Main Features



Useful UI

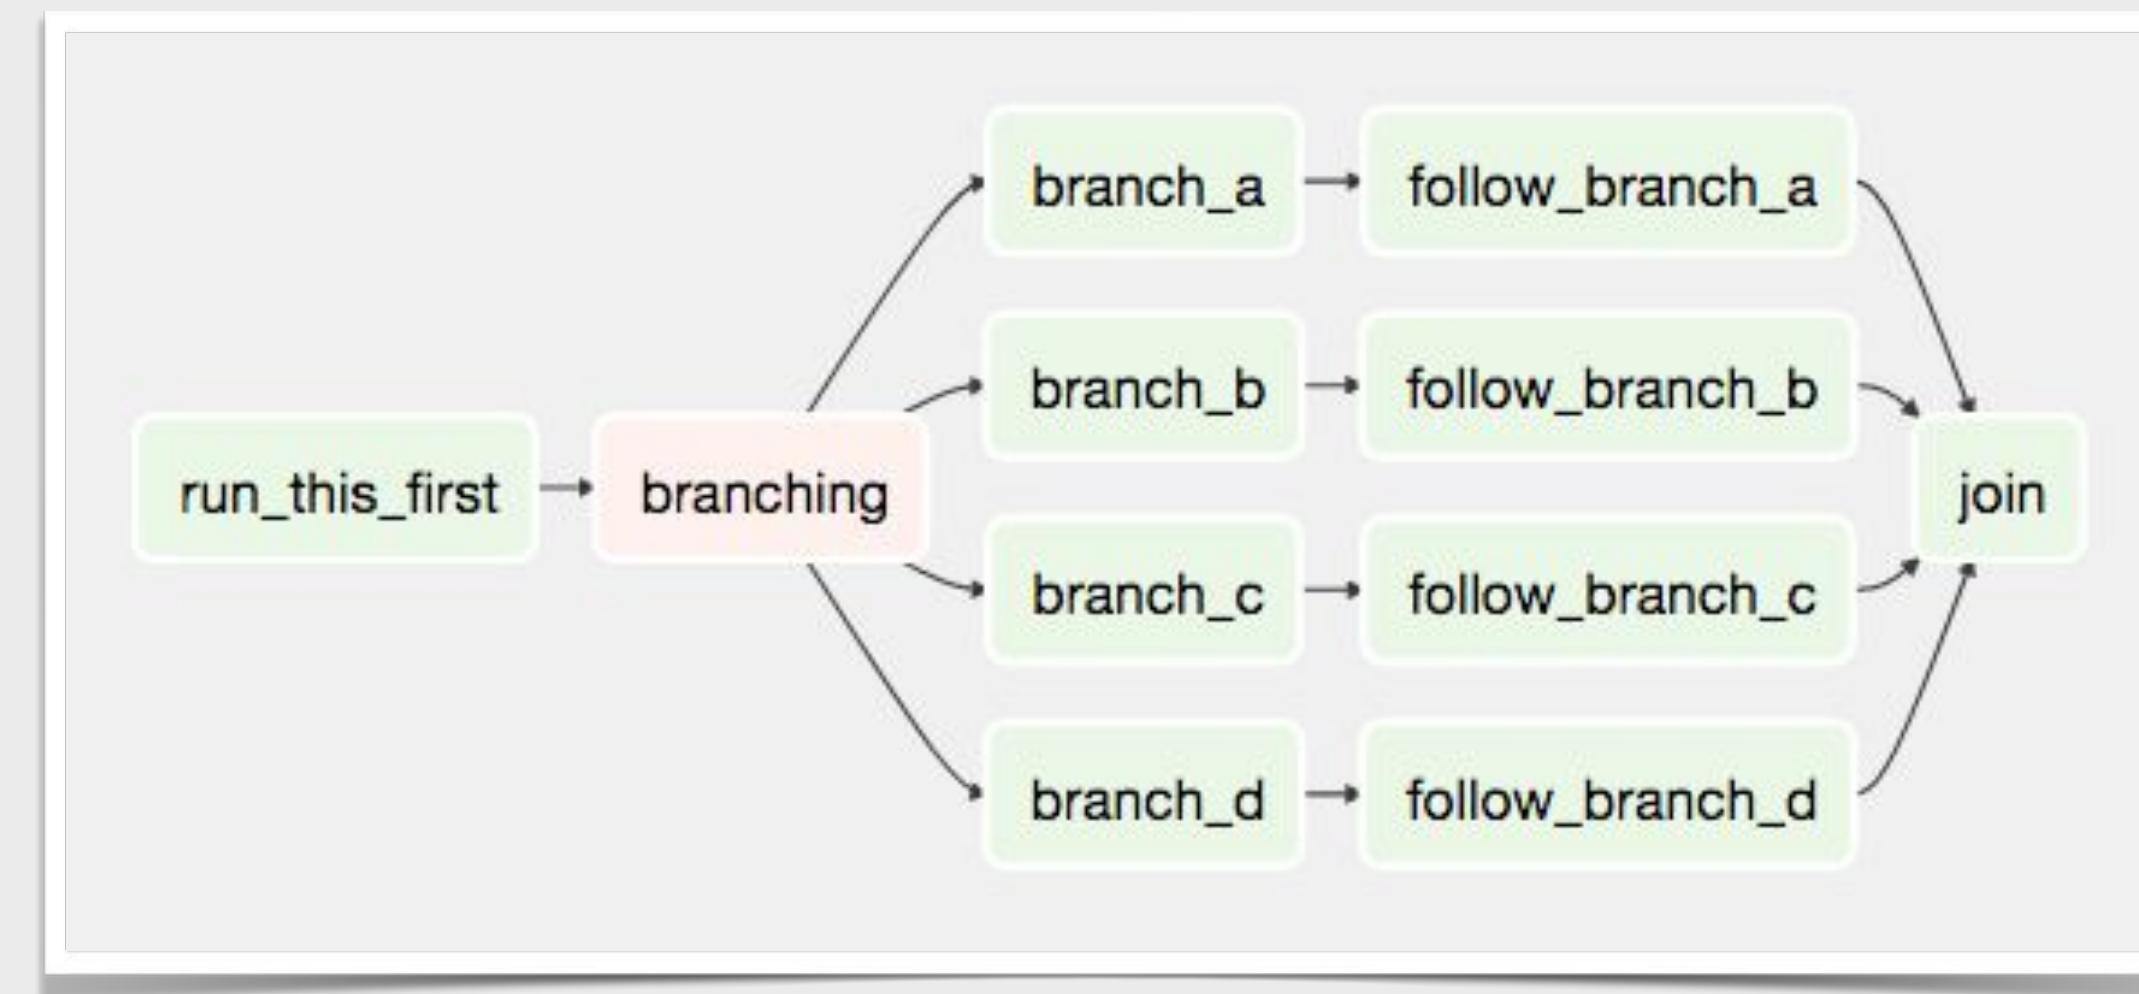
Monitor, schedule and manage your workflows via a robust and modern web application. No need to learn old, cron-like interfaces. You always have full insight into the status and logs of completed and ongoing tasks.



Robust Integrations

Airflow provides many plug-and-play operators that are ready to execute your tasks on Google Cloud Platform, Amazon Web Services, Microsoft Azure and many other third-party services. This makes Airflow easy to apply to current infrastructure and extend to next-gen technologies.

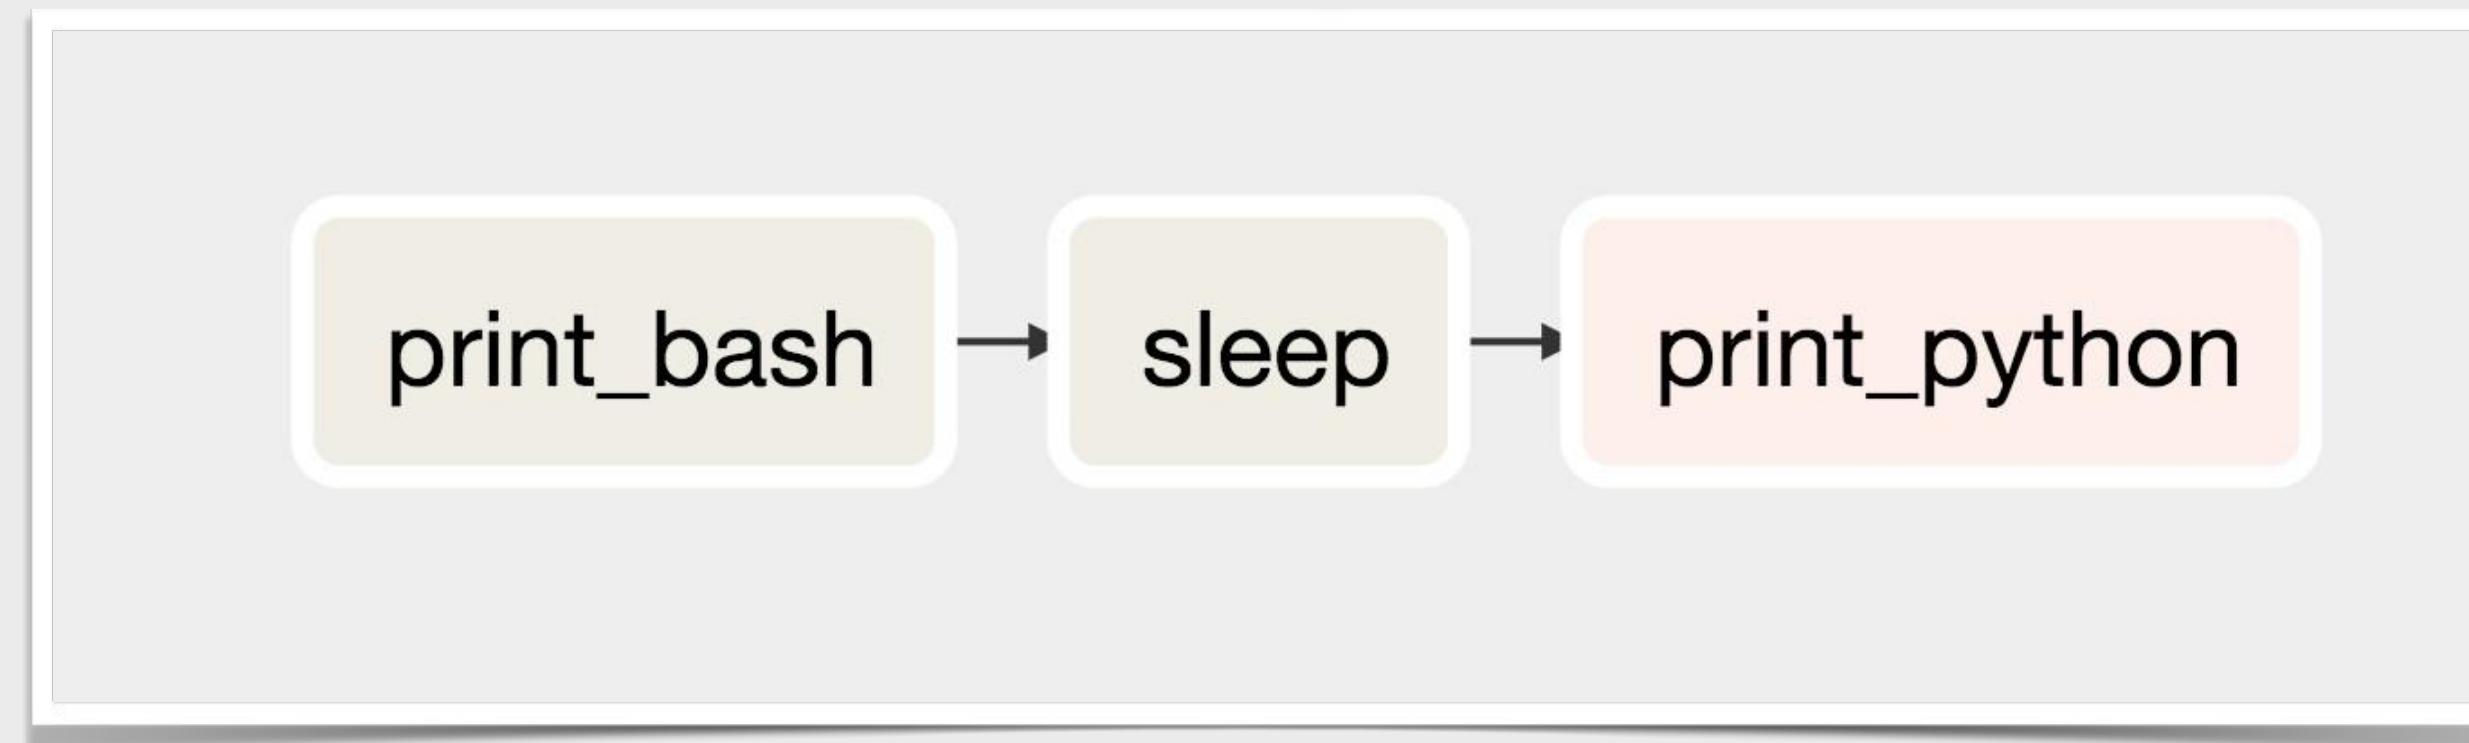
What does it look like?



err...



How it's made



```
import datetime
from airflow import DAG
from airflow.operators.bash_operator import BashOperator
from airflow.operators.python_operator import PythonOperator

def print_commit(): print(' commit
    conf') 🙌

with DAG('commit_dag', schedule_interval='@weekly') as dag:
    print_bash = BashOperator(
        task_id='print_bash', bash_command='echo " sleep = 🙌 commit conf"')
    sleep = BashOperator(
        task_id='sleep', bash_command='sleep 5')
    print_python = PythonOperator(
        task_id='print_python', python_callable=print_commit)

    print_bash >> sleep >> print_python
```

Airflow UI

MAD · NOV 23-2

DAGS

Search:

	i	DAG	Schedule	Owner	Recent Tasks i	Last Run i	DAG Runs i	Links
		commit_dag	@weekly	jriaza	3	2018-11-21 15:57	5	

Showing 1 to 1 of 1 entries

<< < 1 > >>

[Hide Paused DAGs](#)

Airflow UI

The screenshot shows the Airflow UI interface for the DAG: commit_dag. The top navigation bar includes links for Airflow, DAGs, Data Profiling, Browse, Admin, Docs, and About, along with the current timestamp (2018-11-21 17:18:32 UTC) and a refresh button.

The main header displays "On DAG: commit_dag" and "schedule: @weekly". Below this, there are several navigation tabs: Graph View (selected), Tree View, Task Duration, Task Tries, Landing Times, Gantt, Details, and Code. A "Refresh" button is also present.

Below the tabs, there are filters for "success" (highlighted in green), "Base date: 2018-11-21 15:57:09", "Number of runs: 25", "Run: manual_2018-11-21T15:57:08.523284+00:00", and "Layout: Left->Right". A "Go" button and a search bar ("Search for...") are also visible.

At the bottom, there are buttons for BashOperator and PythonOperator, and a status legend: success (green), running (green), failed (red), skipped (pink), retry (yellow), queued (yellow), and no status (grey). A small diagram shows the sequence of tasks: print_bash → sleep → print_python.

Airflow UI

Airflow DAGs Data Profiling ▾ Browse ▾ Admin ▾ Docs ▾ About ▾ 2018-11-21 17:21:02 UTC ⏪

On DAG: commit_dag schedule: @weekly

Graph View Tree View Task Duration Task Tries Landing Times Gantt Details Code Refresh

Base date: 2018-11-21 15:57:08 Number of runs: 25 Go

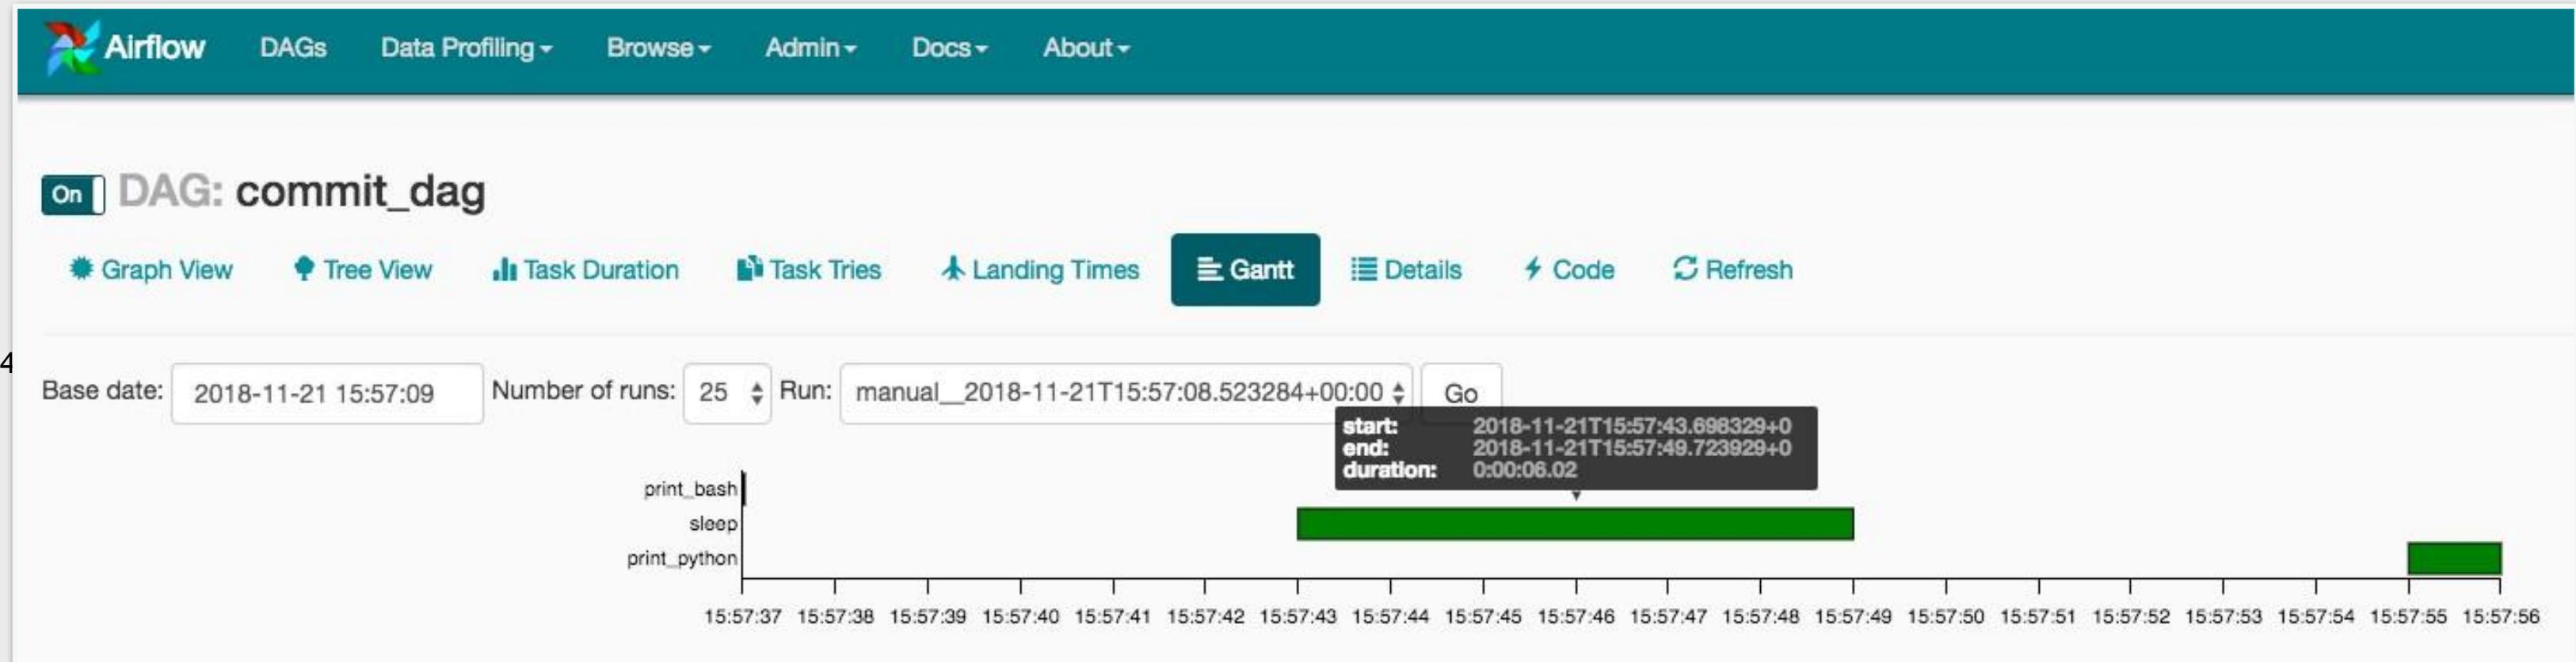
MAD · NOV 23-24 · 2

BashOperator PythonOperator success running failed skipped retry queued no status

```
graph TD; DAG([DAG]) --> print_python1((print_python)); print_python1 --> sleep1((sleep)); sleep1 --> print_bash1((print_bash))
```

The screenshot displays the Airflow UI for the DAG: commit_dag. At the top, there's a navigation bar with links for Airflow, DAGs, Data Profiling, Browse, Admin, Docs, and About, along with a timestamp of 2018-11-21 17:21:02 UTC and a refresh button. Below the navigation is a header with the DAG name "commit_dag" and its schedule "schedule: @weekly". A "Graph View" tab is active. Below the header are buttons for "Tree View", "Task Duration", "Task Tries", "Landing Times", "Gantt", "Details", "Code", and "Refresh". Underneath these are filters for "Base date" (2018-11-21 15:57:08), "Number of runs" (25), and a "Go" button. A legend indicates task types: BashOperator (light blue circle) and PythonOperator (pink circle). Status indicators include green for success, yellow for running, red for failed, pink for skipped, yellow for retry, grey for queued, and white for no status. The main area shows a graph of the DAG with four nodes: "[DAG]", "print_python", "sleep", and "print_bash". The "sleep" node has a self-loop arrow. To the right is a Gantt chart for November, showing 25 runs where all tasks have been successful.

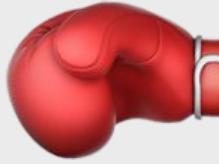
Airflow UI



Operators

- **Action:** perform an action locally or make a call to an external system to perform another action
- **Transfer:** move data from one system to another
- **Sensor:** wait for and detect some condition in a source system

Action operators



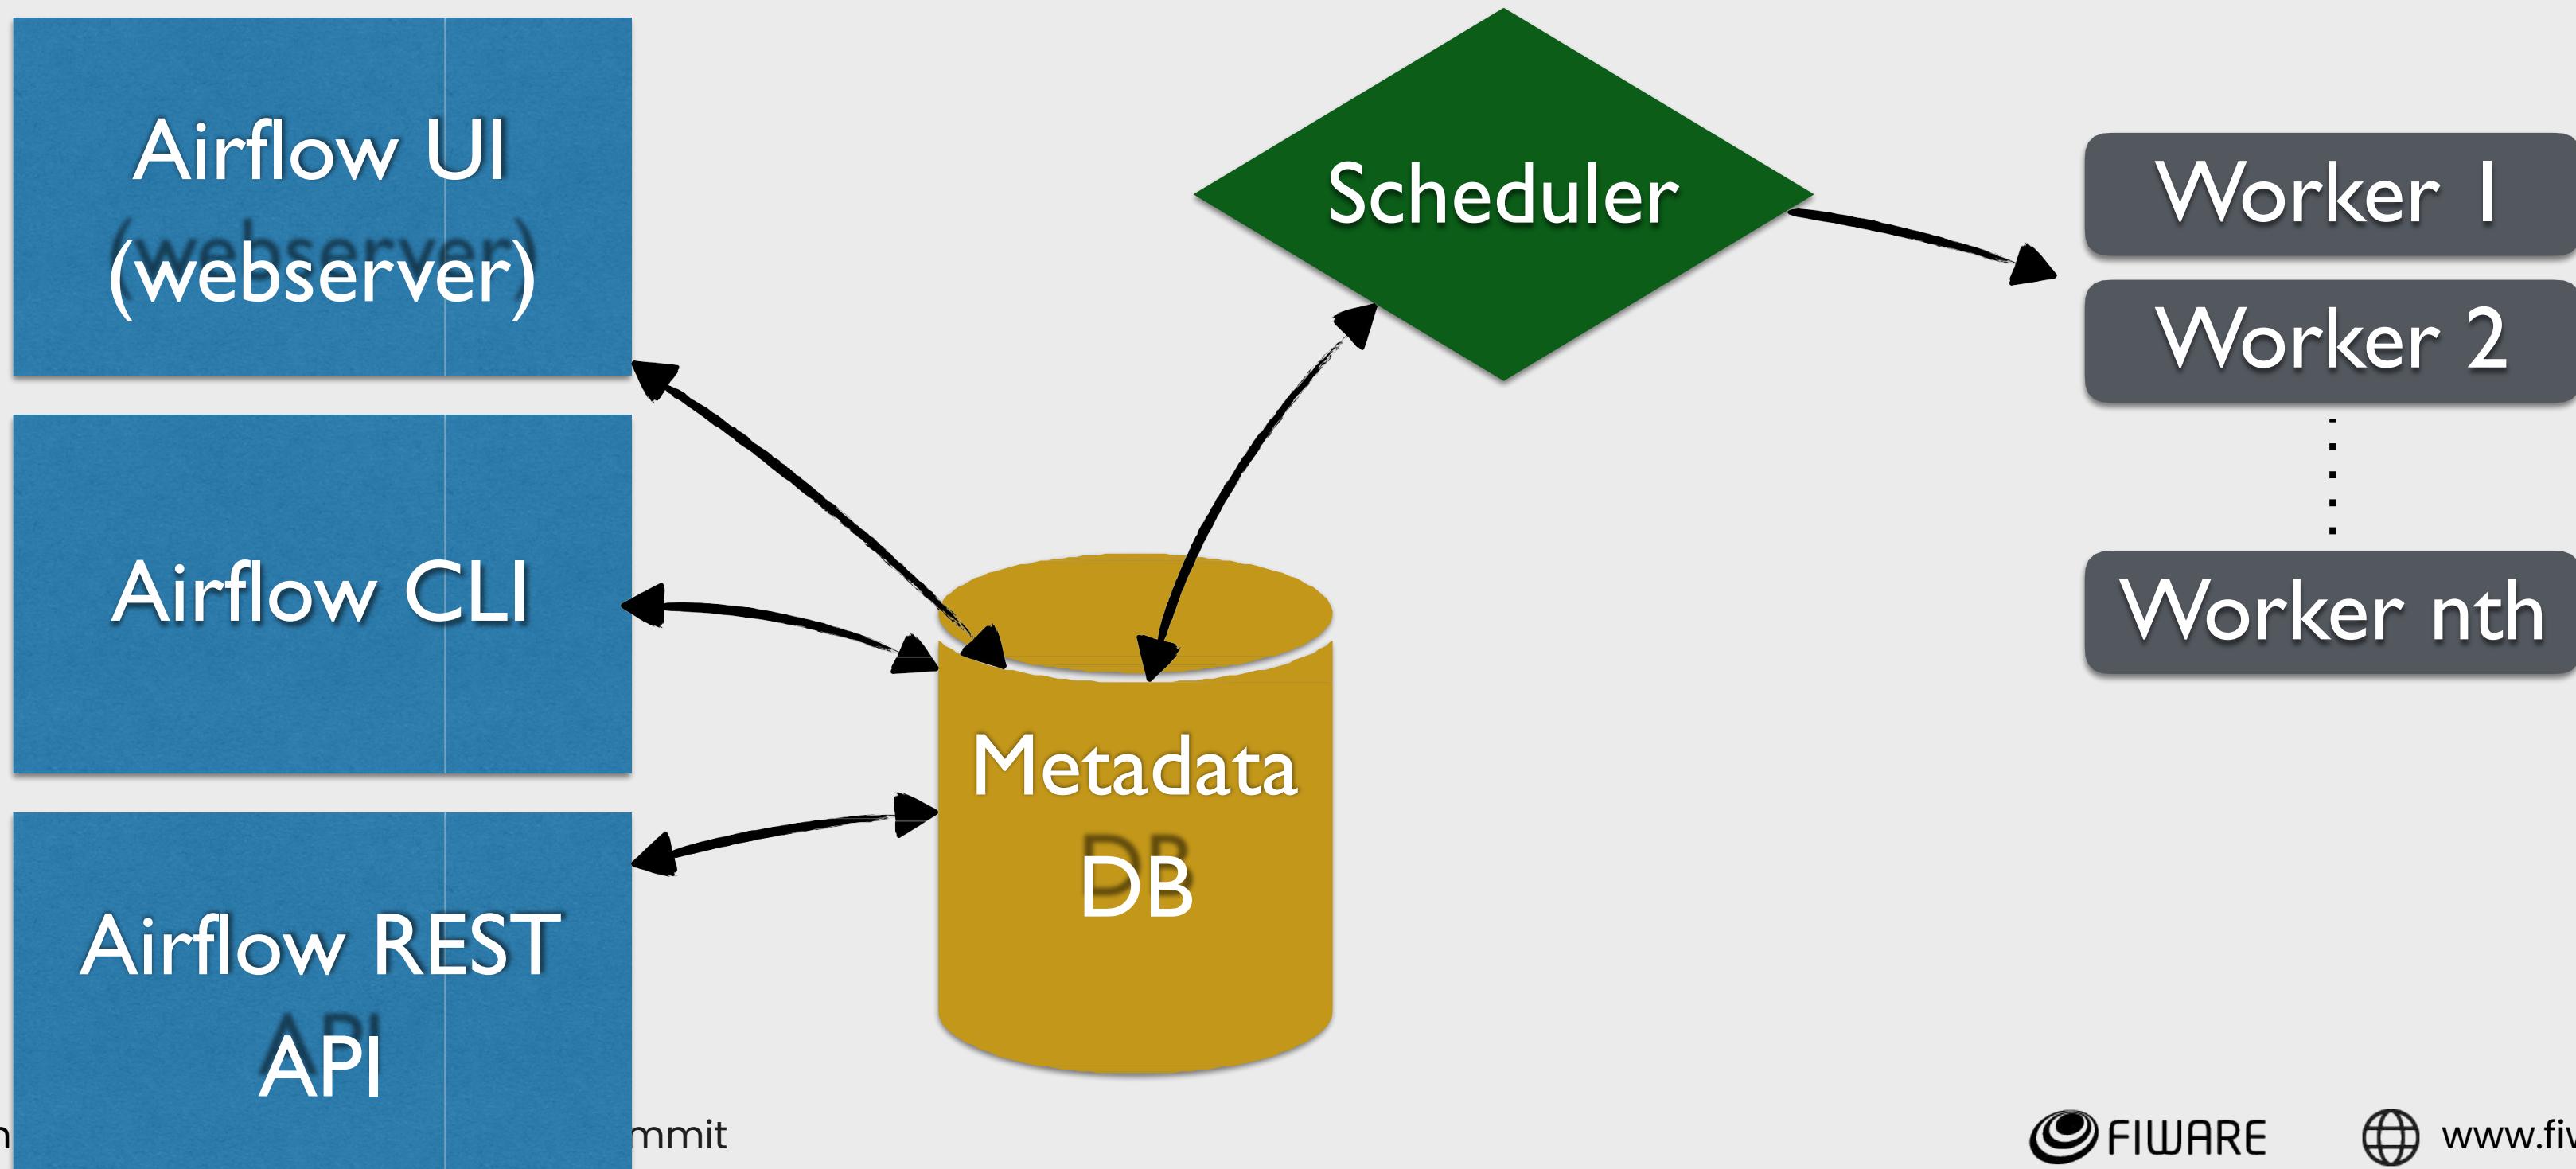
- Perform an action such as executing a Python function or submitting a Spark Job
- **Built-in** BashOperator, PythonOperator, DockerOperator, EmailOperator, ...

Sensor operators

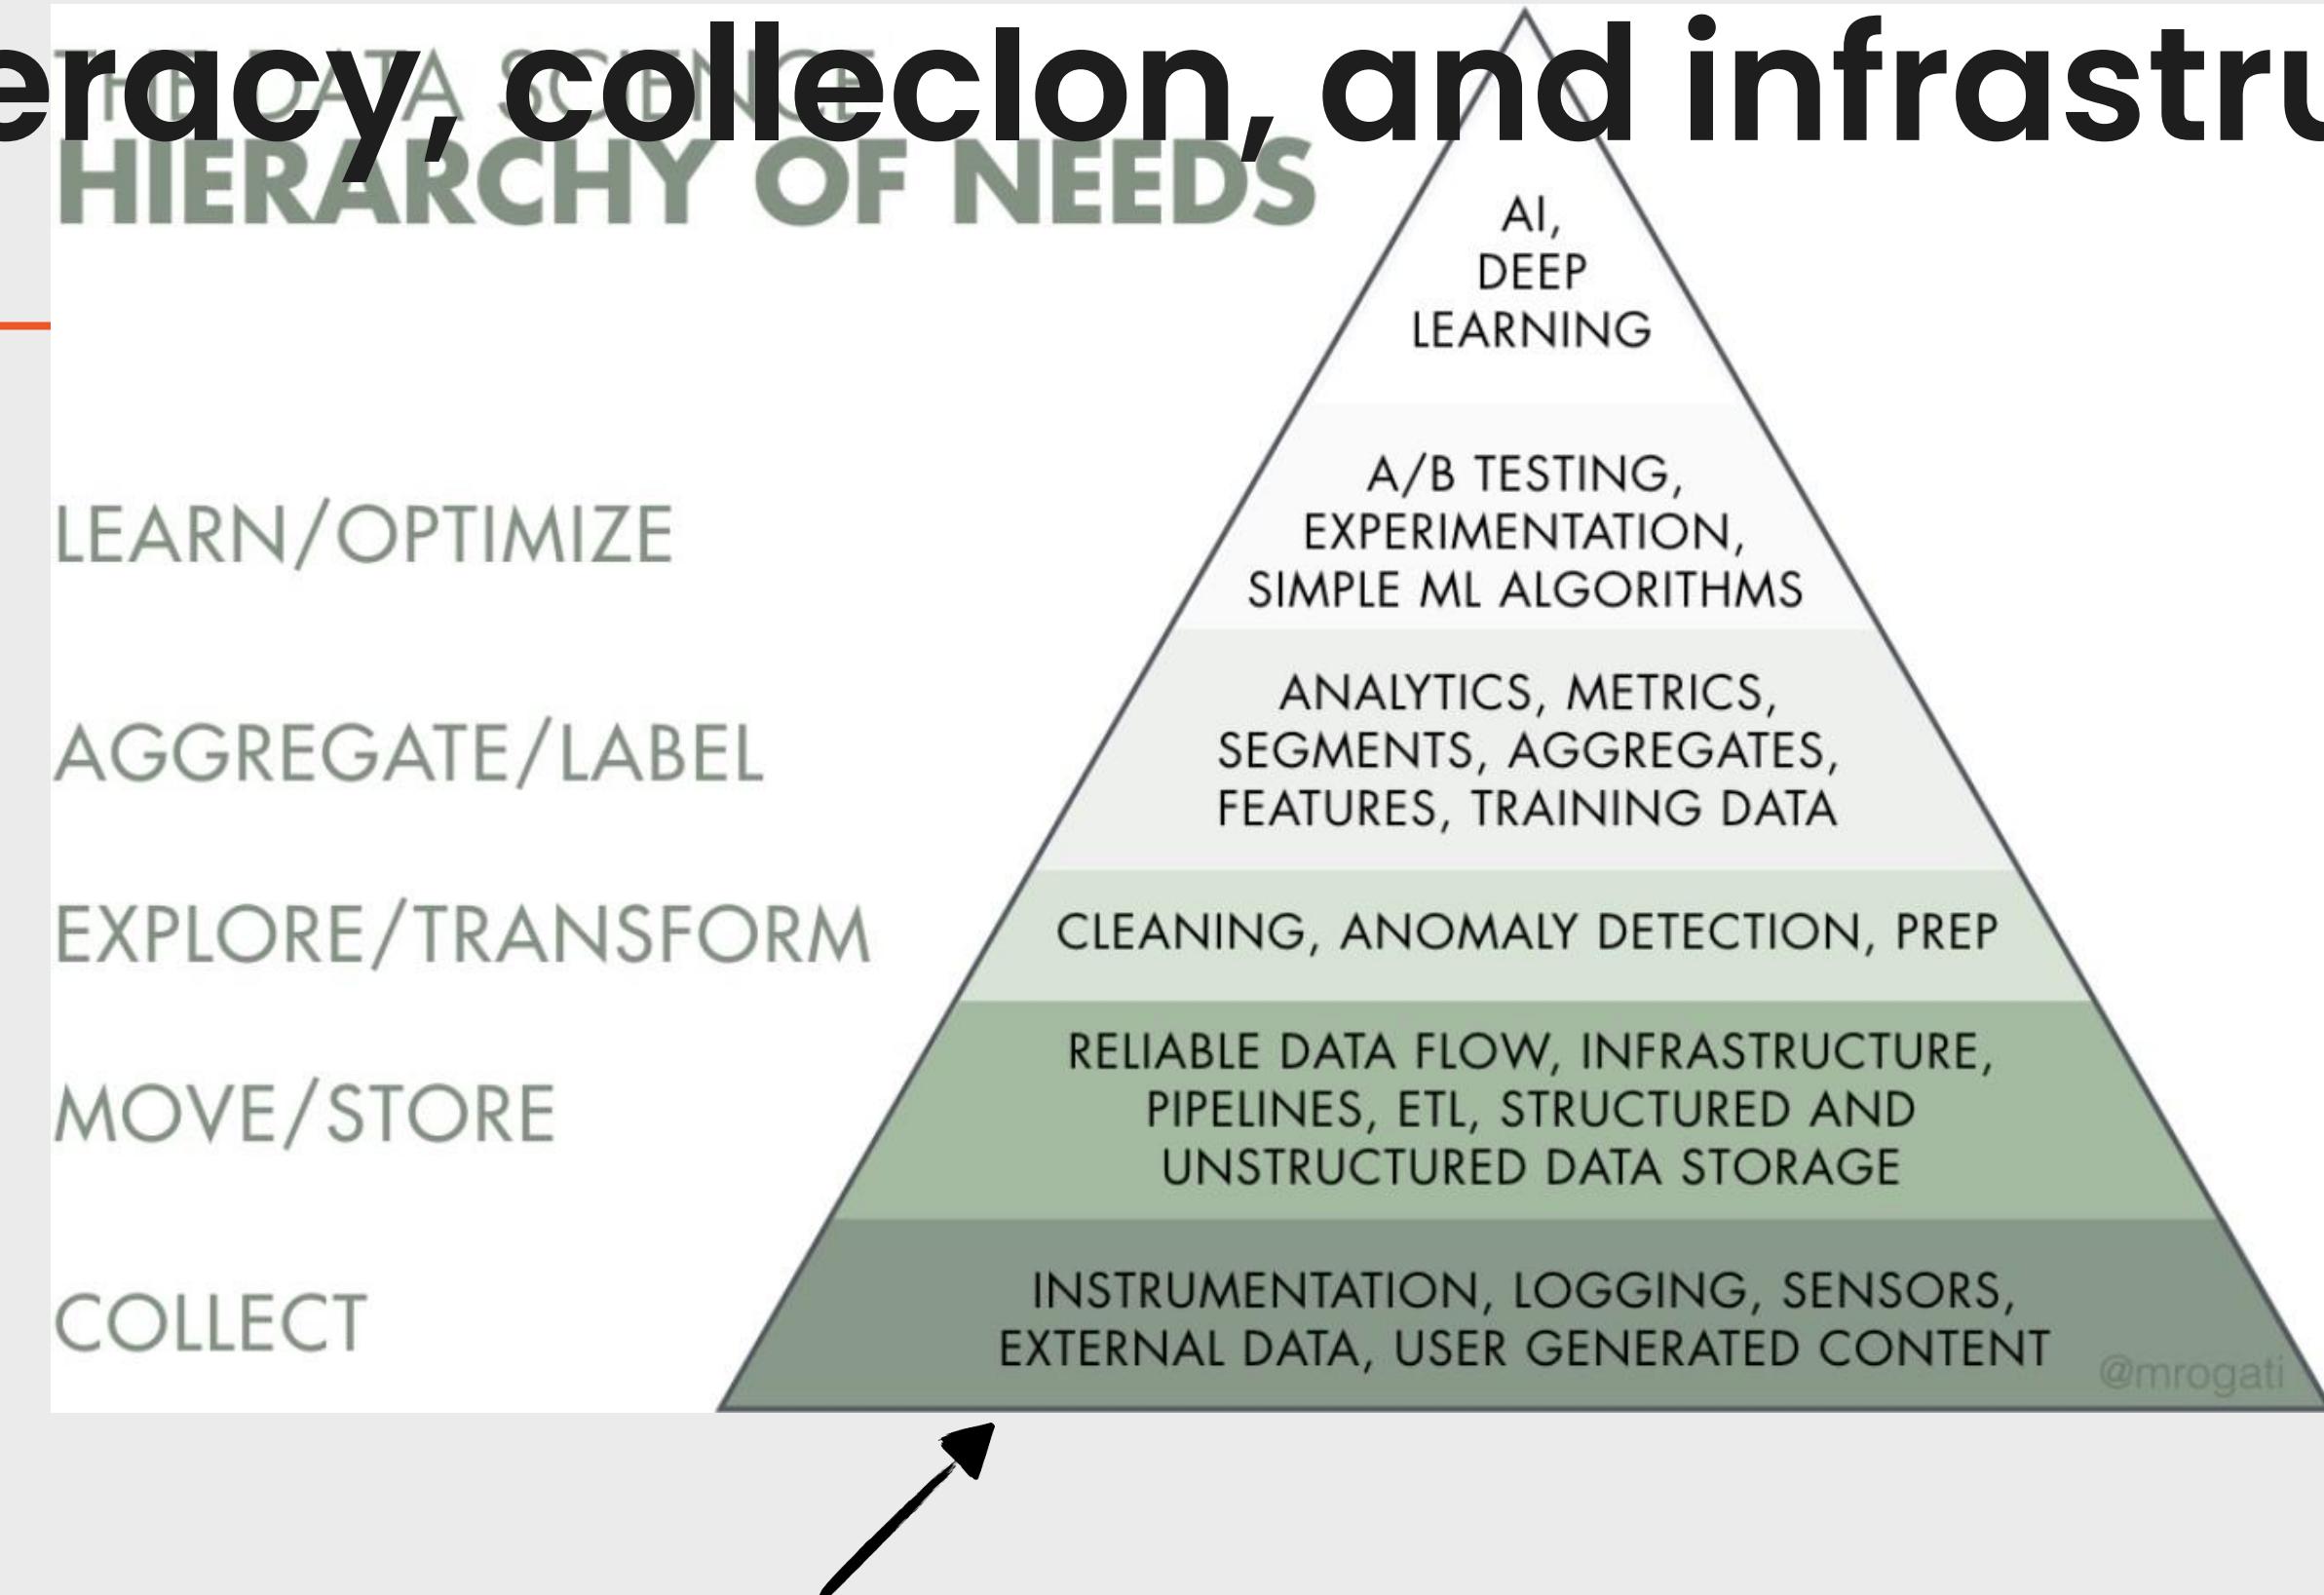


- Triggers downstream tasks in the dependency graph when a certain criteria is met. For example, checking for a certain file has become available on S3 before using it downstream
- **Built-in** HiveParXXonSensor, HcpSensor, S3KeySensor,

The big picture



Data literacy, collection, and infrastructure



Airflow deployment

MAD · NOV 23-24 · 2018



<https://github.com/idealista/airflow-role>

Astronomer

- <https://www.astronomer.io>
- Tool to create Airflow deployments easily.

```
curl -sSL https://install.astronomer.io | sudo bash  
astro dev init #(genera scaffold)  
astro dev start  
Web interface en puerto 8080
```



The Maturity of the MLOPS Process

- MLOps is an ML engineering culture and practice that aims at unifying ML system development (Dev) and ML system operation (Ops).

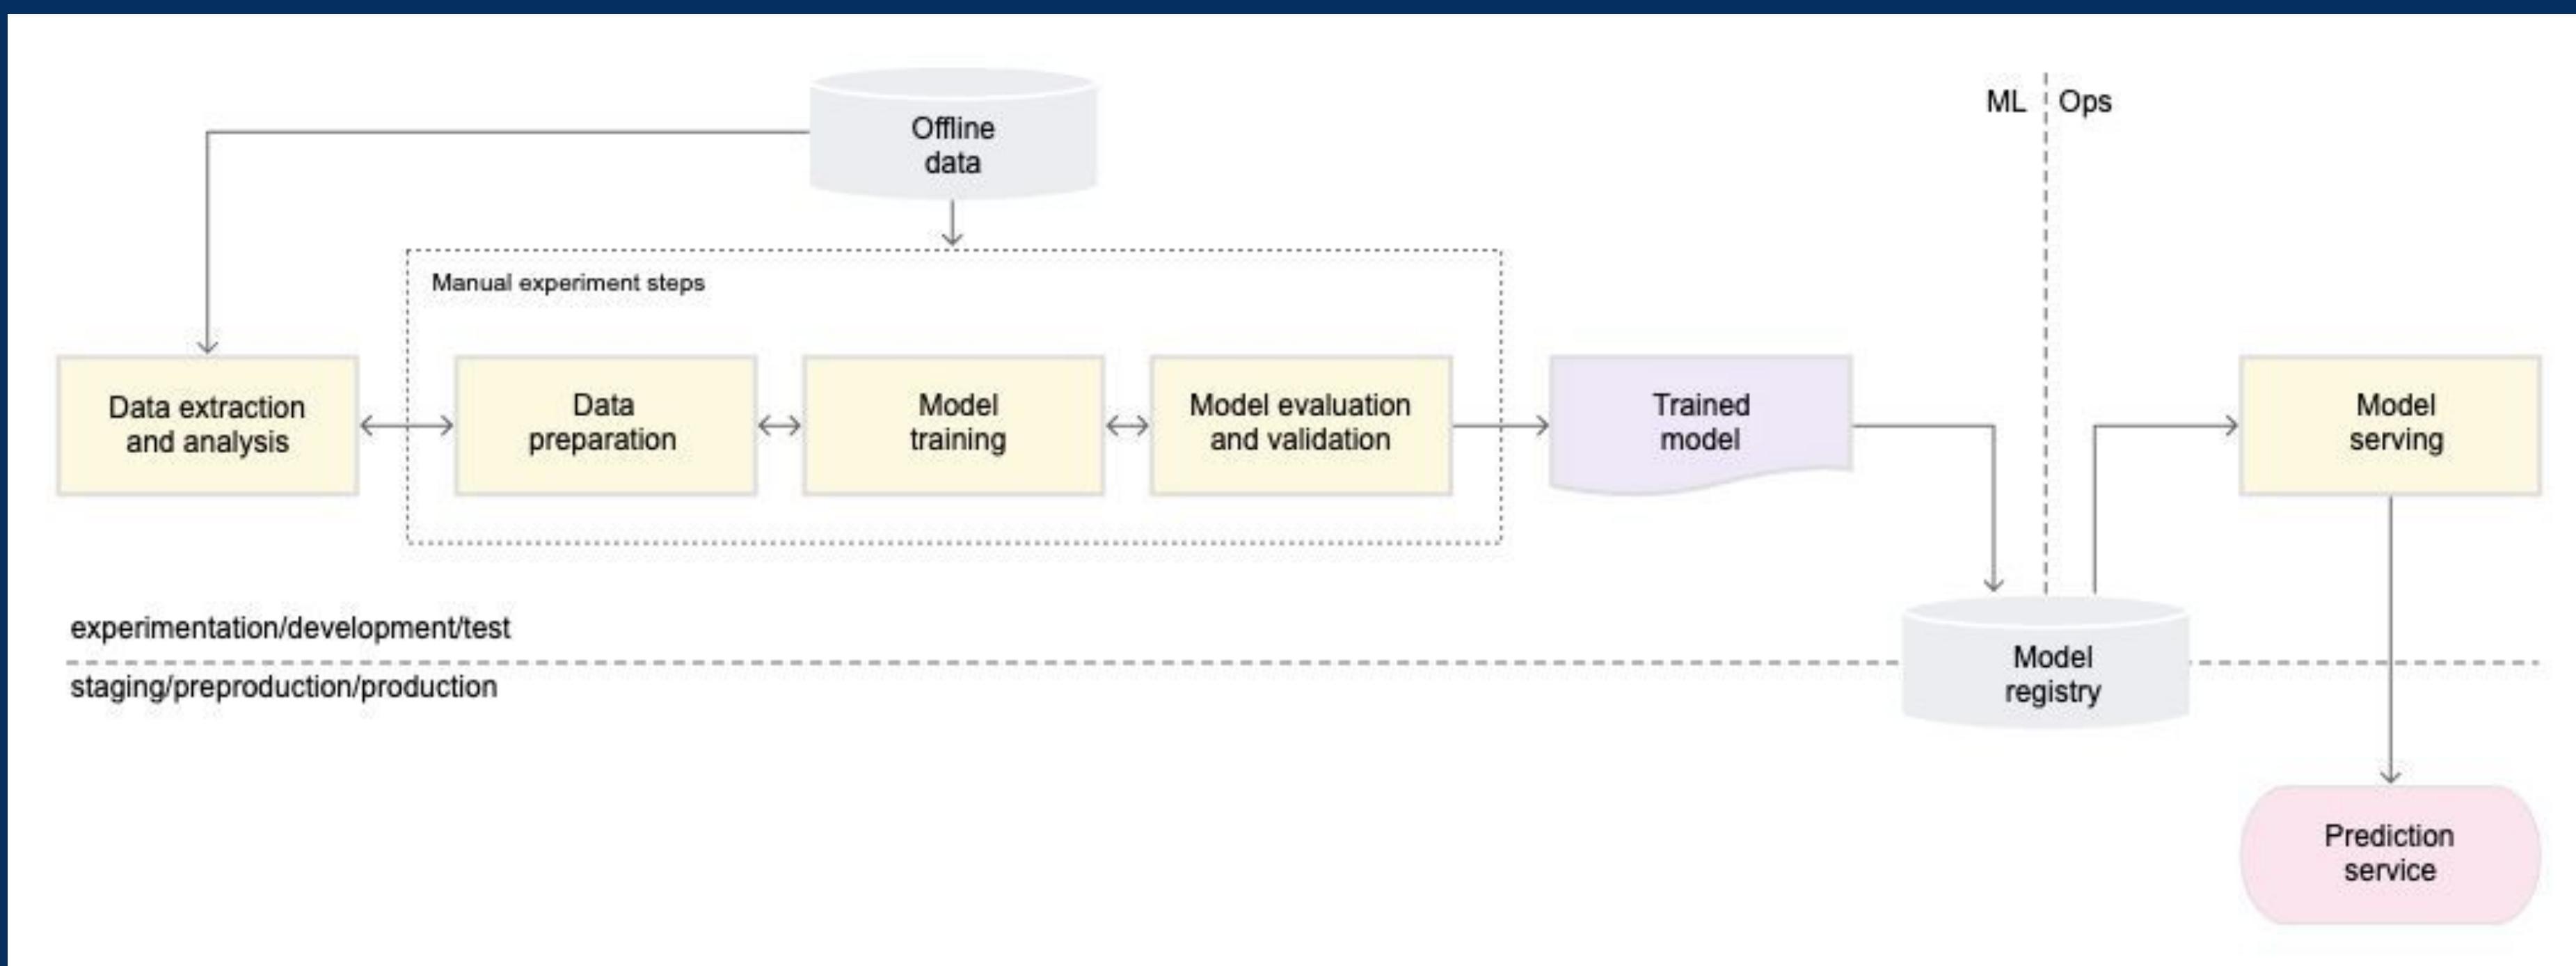
level 0 Manual process

level 1 ML pipeline automation

level 2 CI/CD pipeline automation

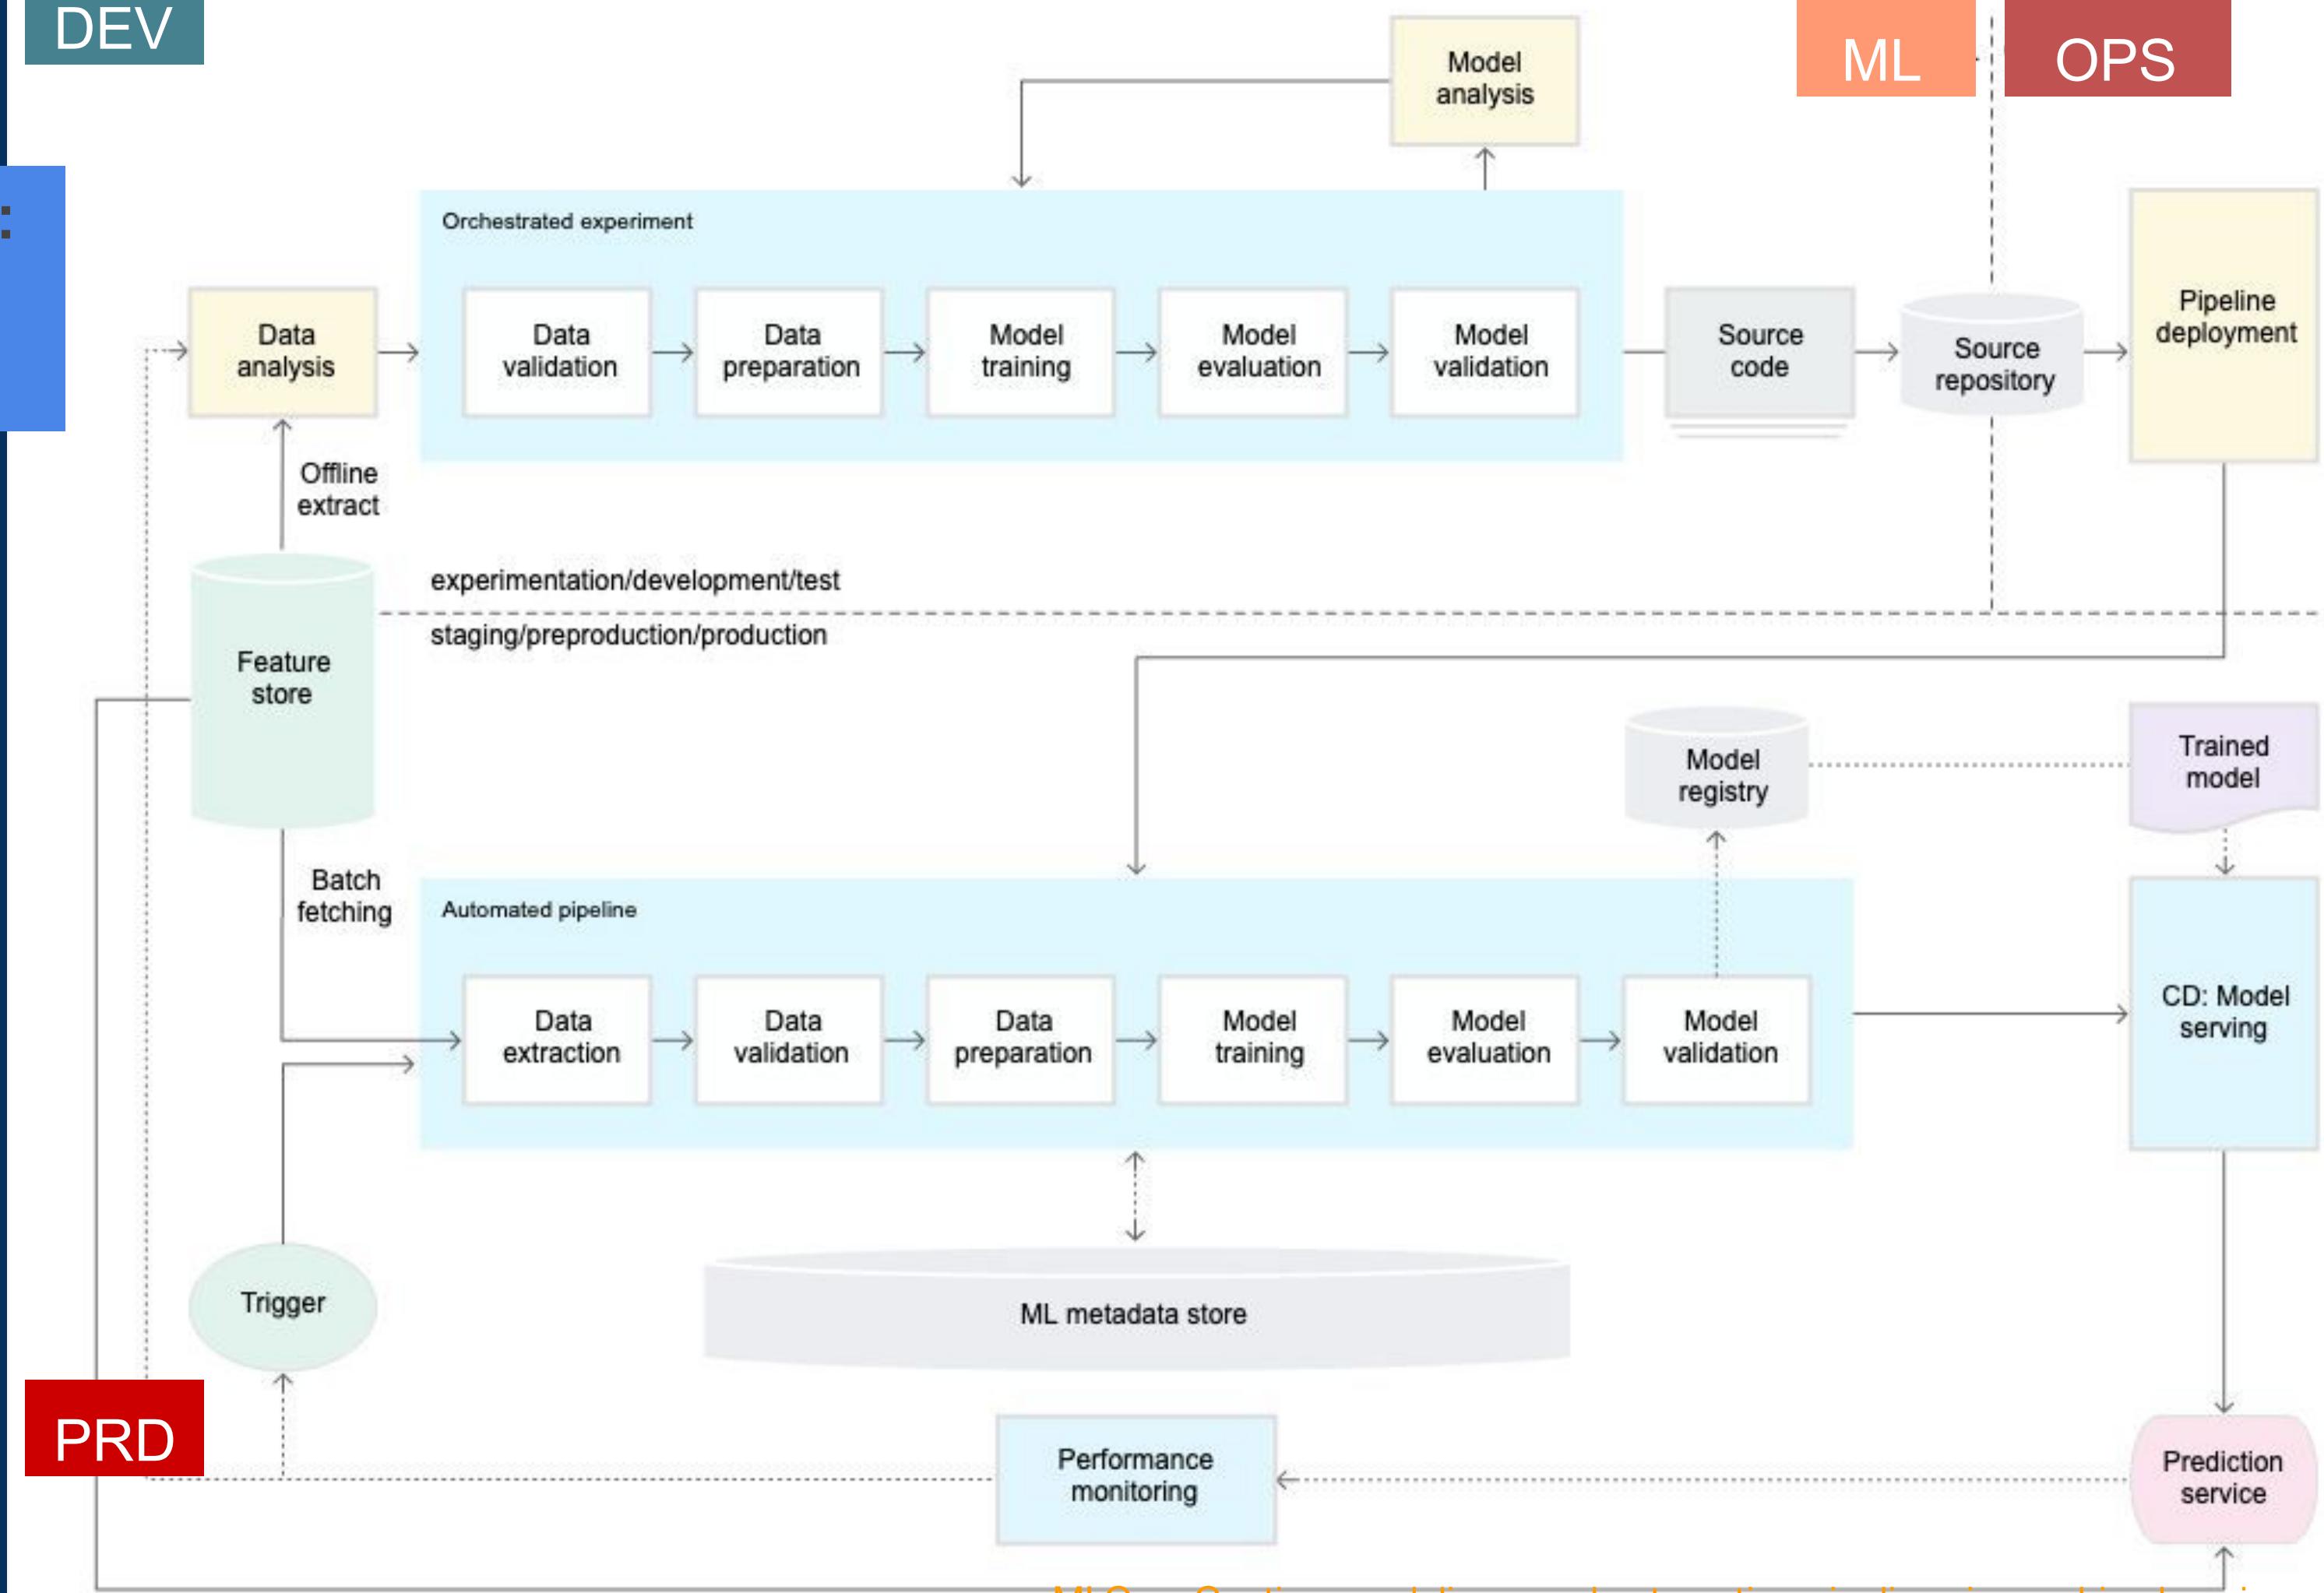
- Practicing MLOps means that you advocate for **automation** and **monitoring** at all steps of ML system construction, including integration, testing, releasing, deployment and infrastructure management.

MLOps level 0: Manual process



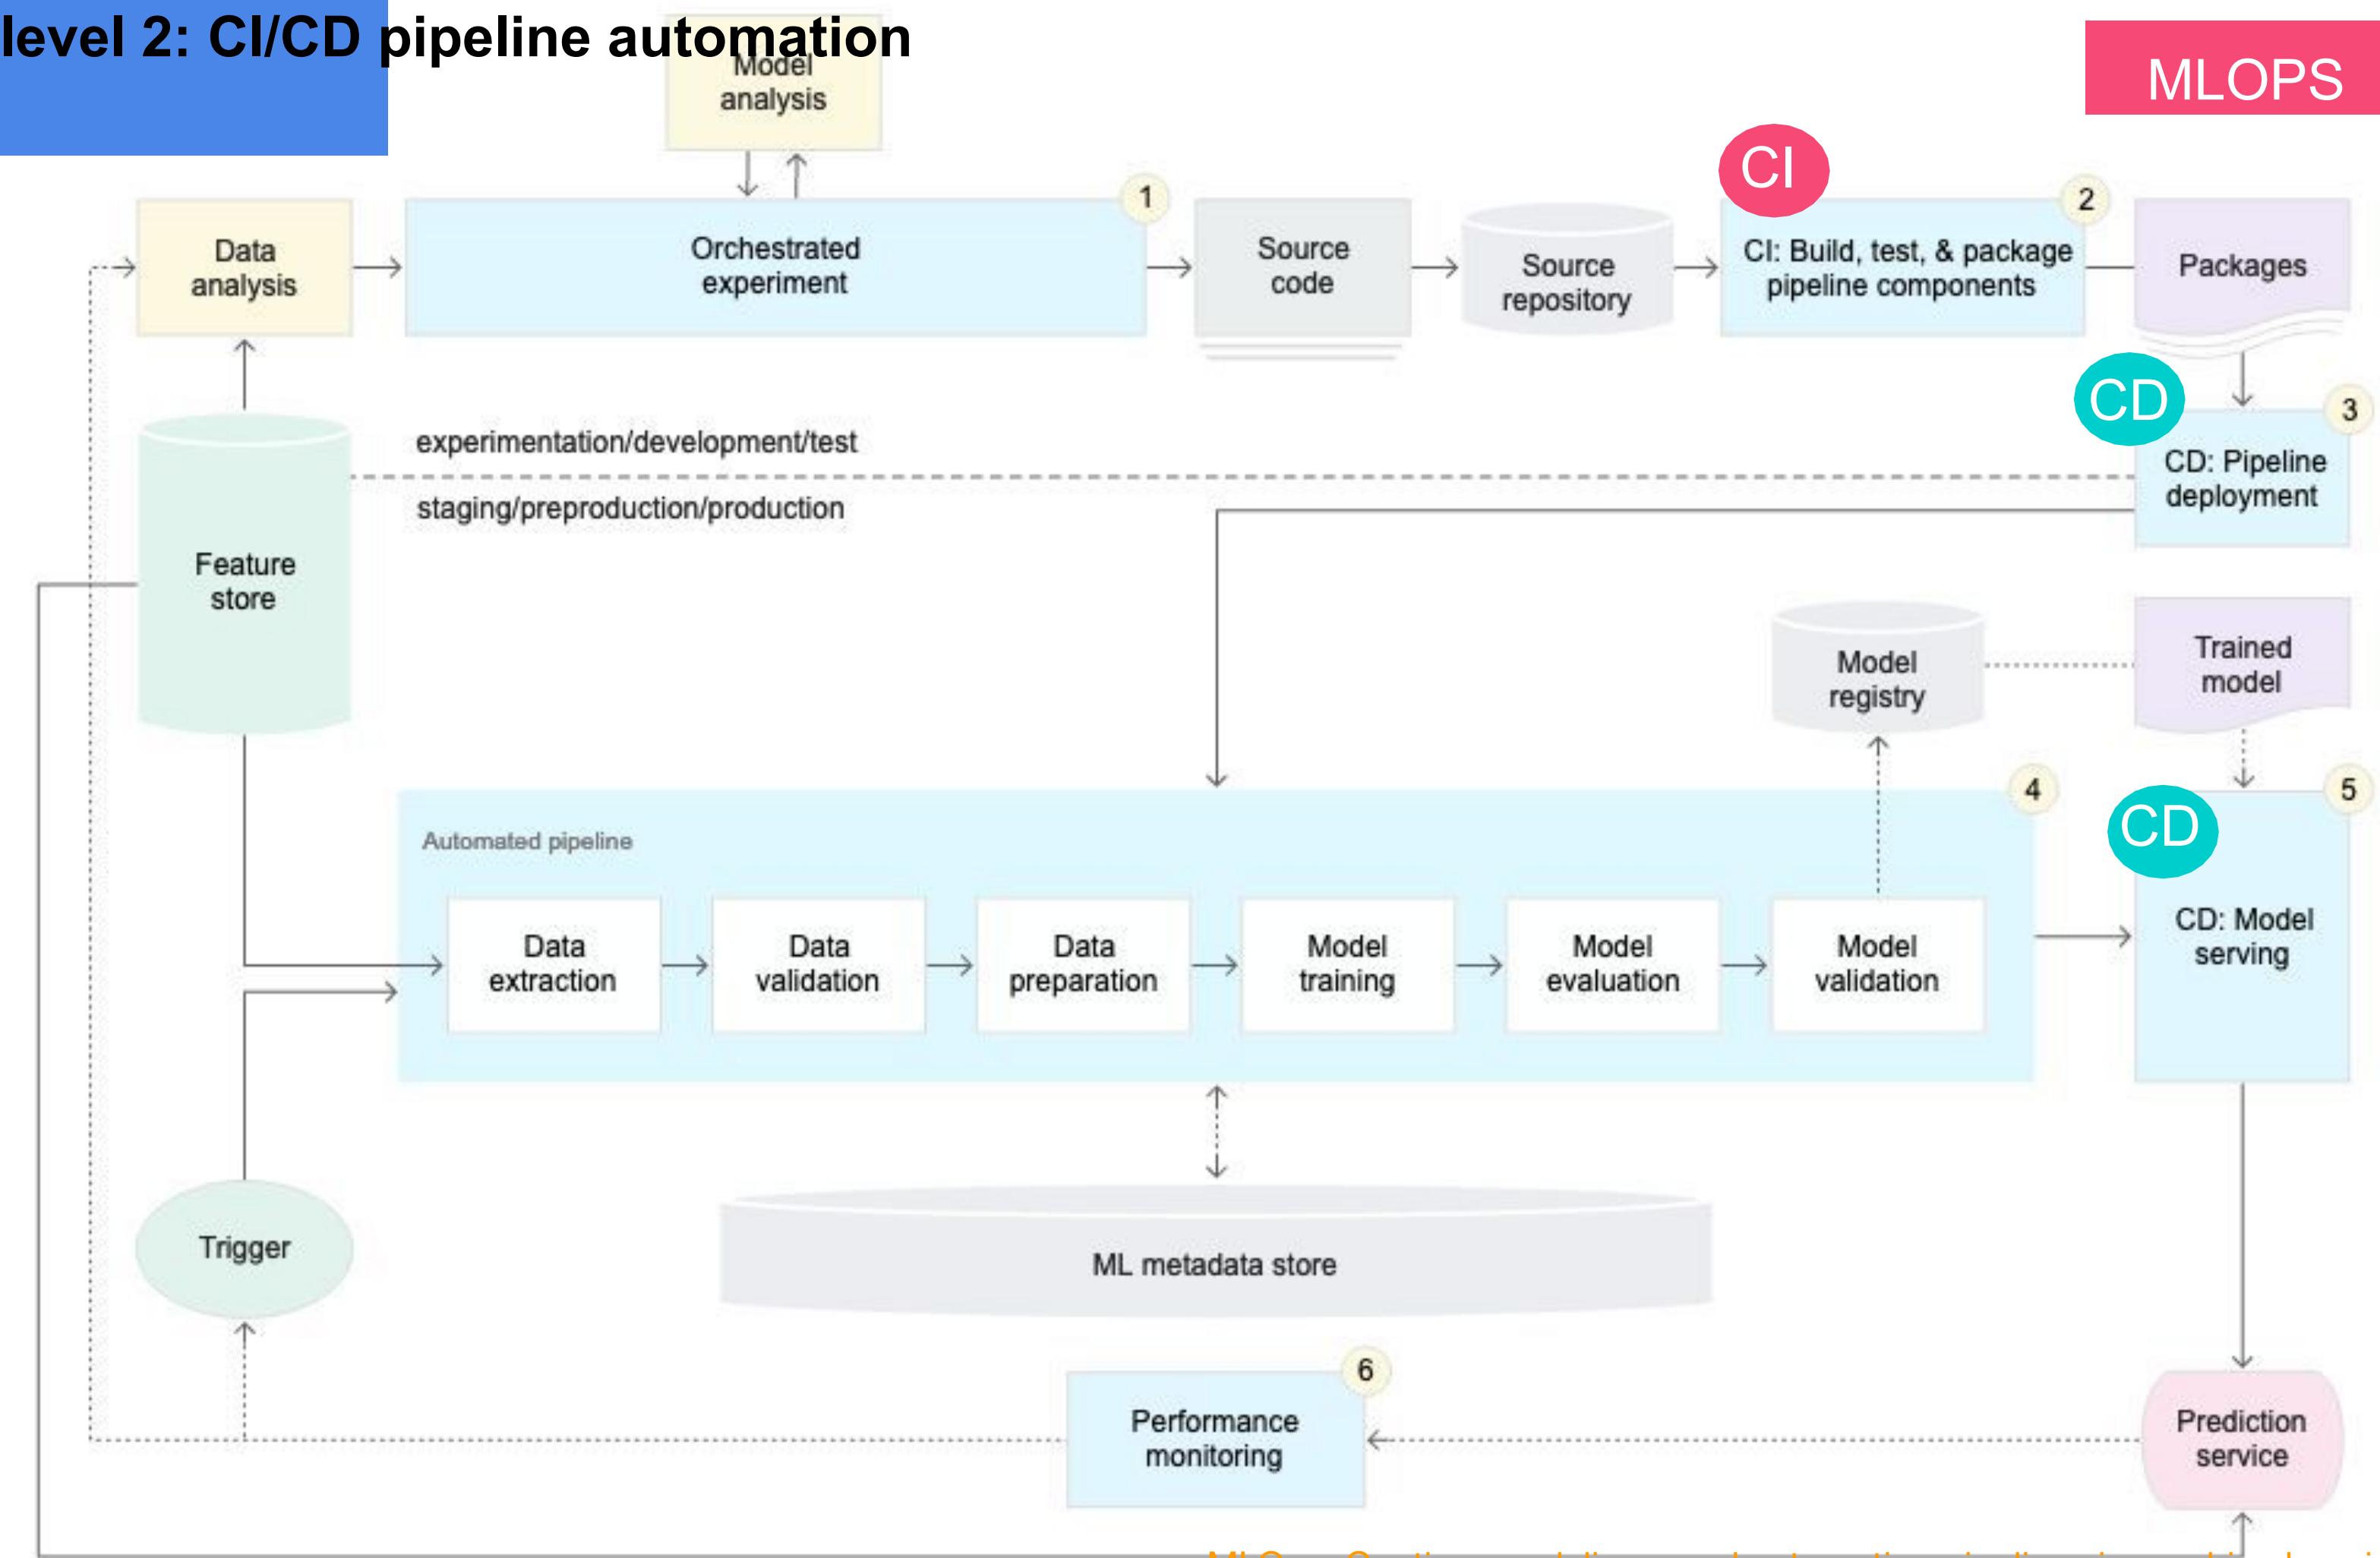
DEV

MLOps level 1: ML pipeline automation



MLOps level 2: CI/CD pipeline automation

MLOPS



Experiment Tracking

Model development & Post-Deployment

- Prove value of experiment
 - Need baseline to show and compare
- Collaborate
 - Need to refer and access models and artifacts from other members
- Reproduce work
 - Need same parameters and model of ex-run

What we should log/track in ML

Log day-to-day work in ML life cycle

- Hyper parameters
- Training/modeling performances
- Model
 - Type
 - Building environment
 - Modeling version
- and so on

Evaluation metrics

- Mean Absolute Error (MAE)
- Mean Squared Error (MSE)
- Root Mean Squared Error (RSME)
- R-squared (r²)
- ...

parameters

- Convolutional filter
- Kernel_size
- Max pooling
- Dropout
- Dense
- Batch_size
- Epochs
-



- Since 2018 from DataBricks (**Main contributor**)
- An open platform for the machine learning lifecycle
- Python Library; runs locally and on the cloud
- Built-in UI for experiment visualization
- Logging integrations for major frameworks: scikit-learn, PyTorch, TF,...

mlflow Components

MLflow is an open source platform to manage the ML lifecycle, including experimentation, reproducibility, deployment, and a central model registry. MLflow currently offers four components:

MLflow Tracking

Record and query experiments: code, data, config, and results

[Read more](#)

MLflow Projects

Package data science code in a format to reproduce runs on any platform

[Read more](#)

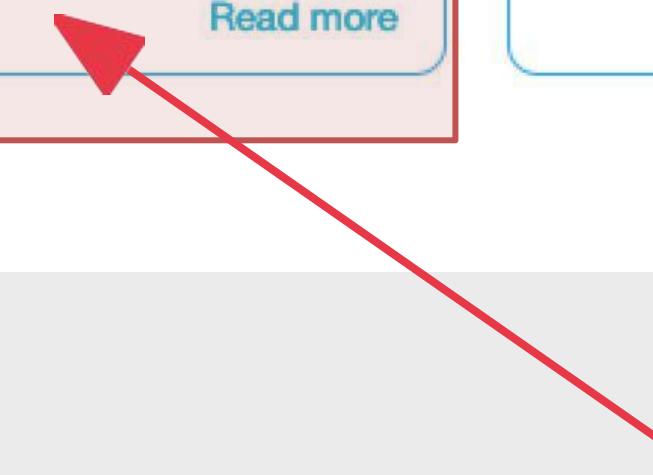
MLflow Models

Deploy machine learning models in diverse serving environments

[Read more](#)

Model Registry

Store, annotate, discover, and manage models in a central repository

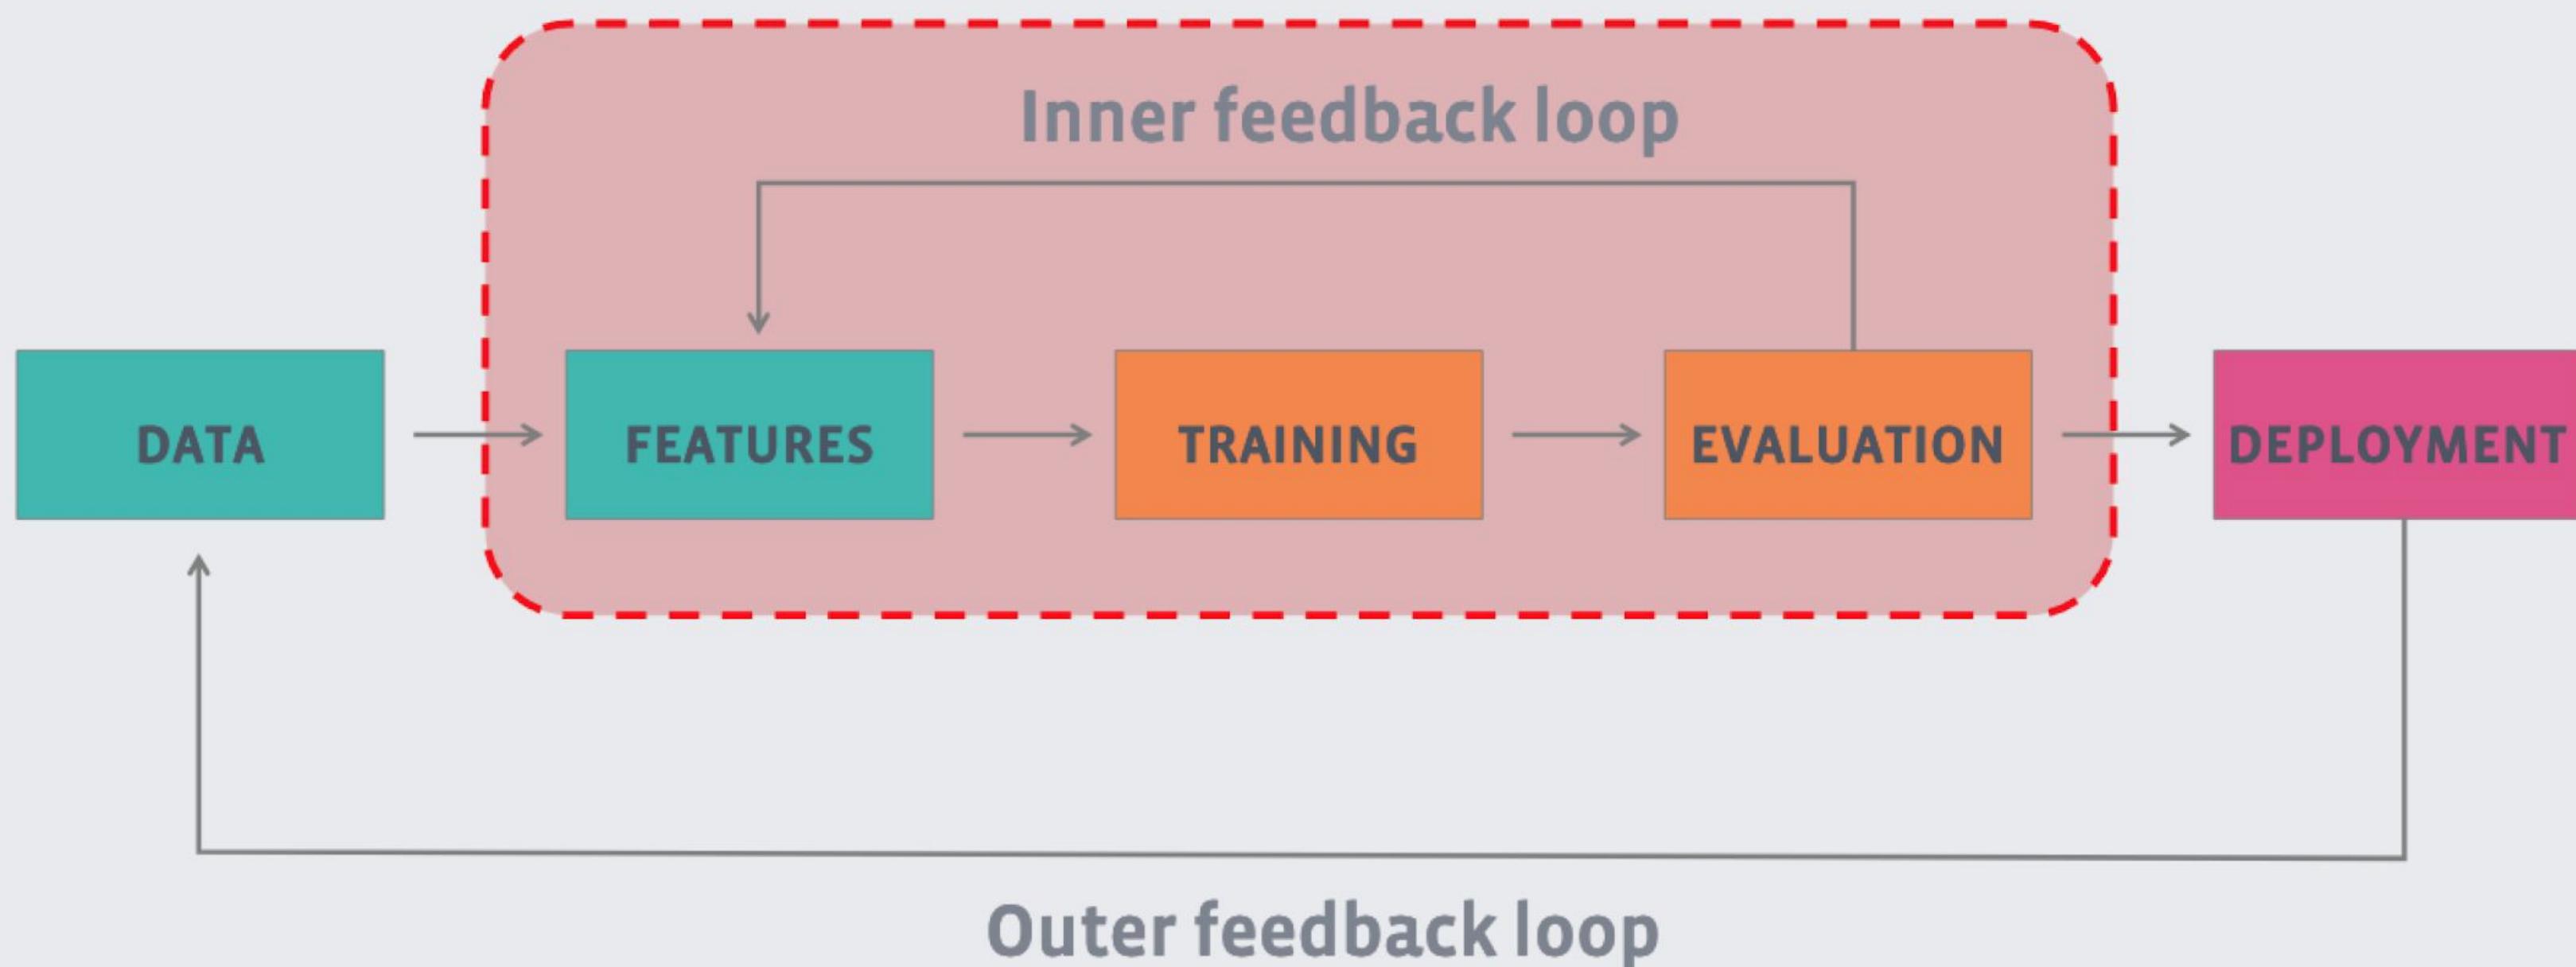
[Read more](#)

Main focus of this sharing!

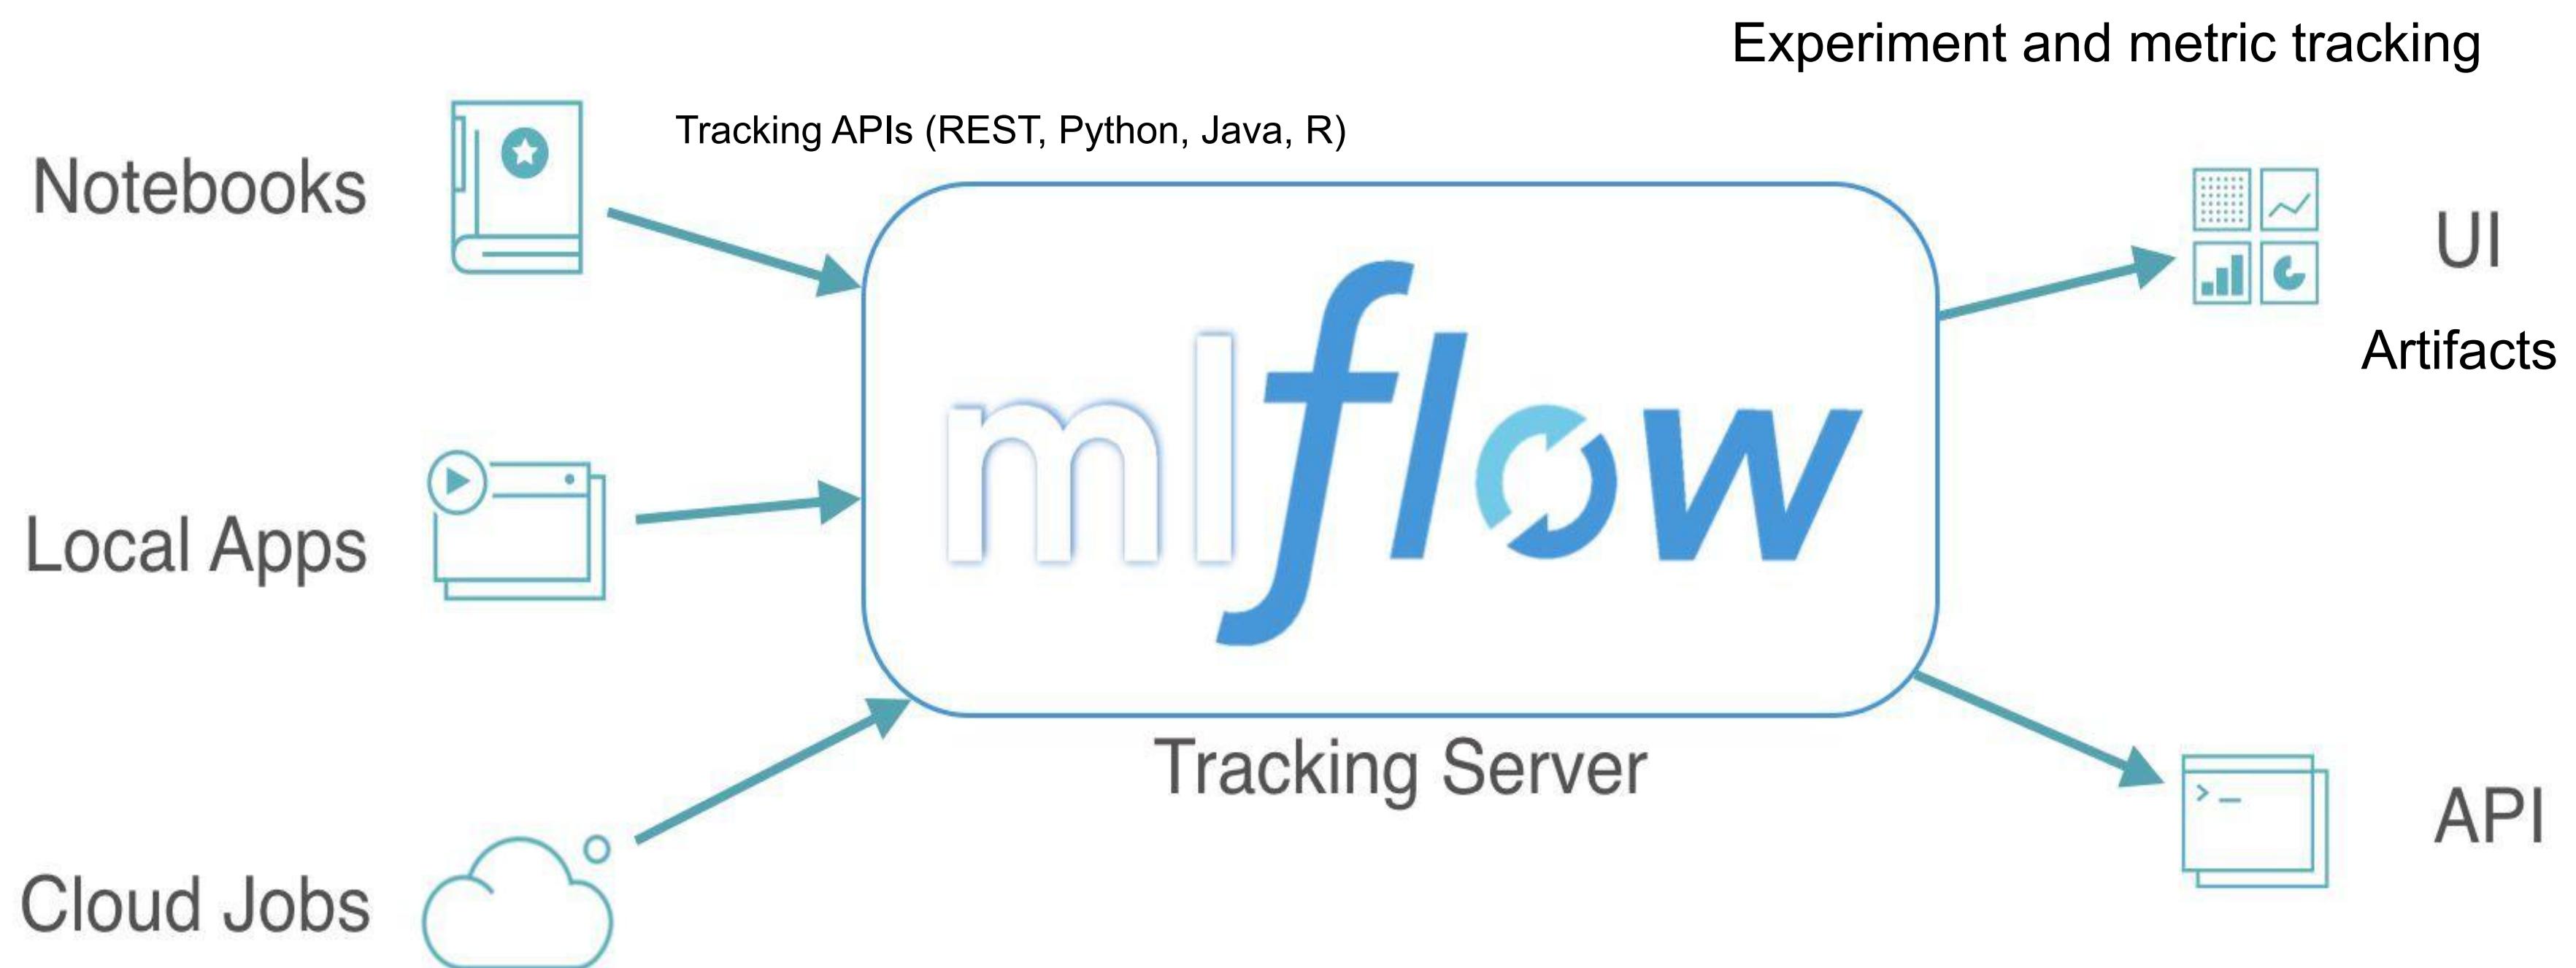
56

MLflow Components

Experiment tracking

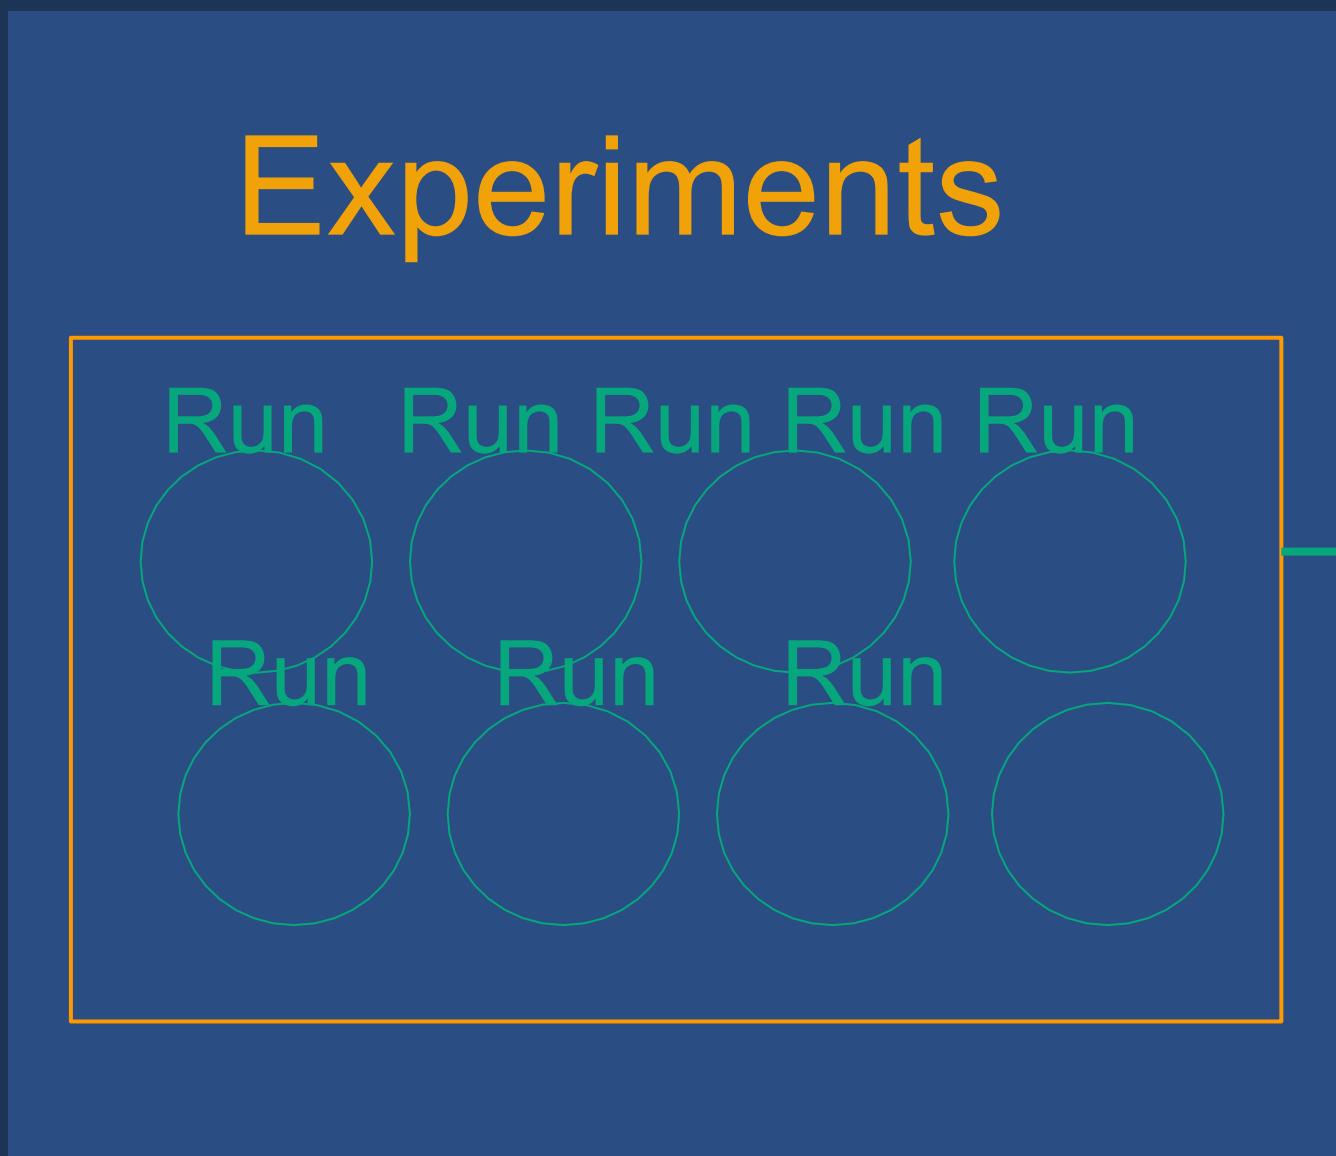


Experiment/Production pipeline



Terminology

Project



Experiments

Run Run Run Run Run
Run Run Run

Entity

- Code version
- Start and end time
- Source
- (Hyper) Parameters
- Metrics
- Tags/Notes

Artifacts

- Output files
 - a. Images
 - b. Pickled models
 - c. Data files...



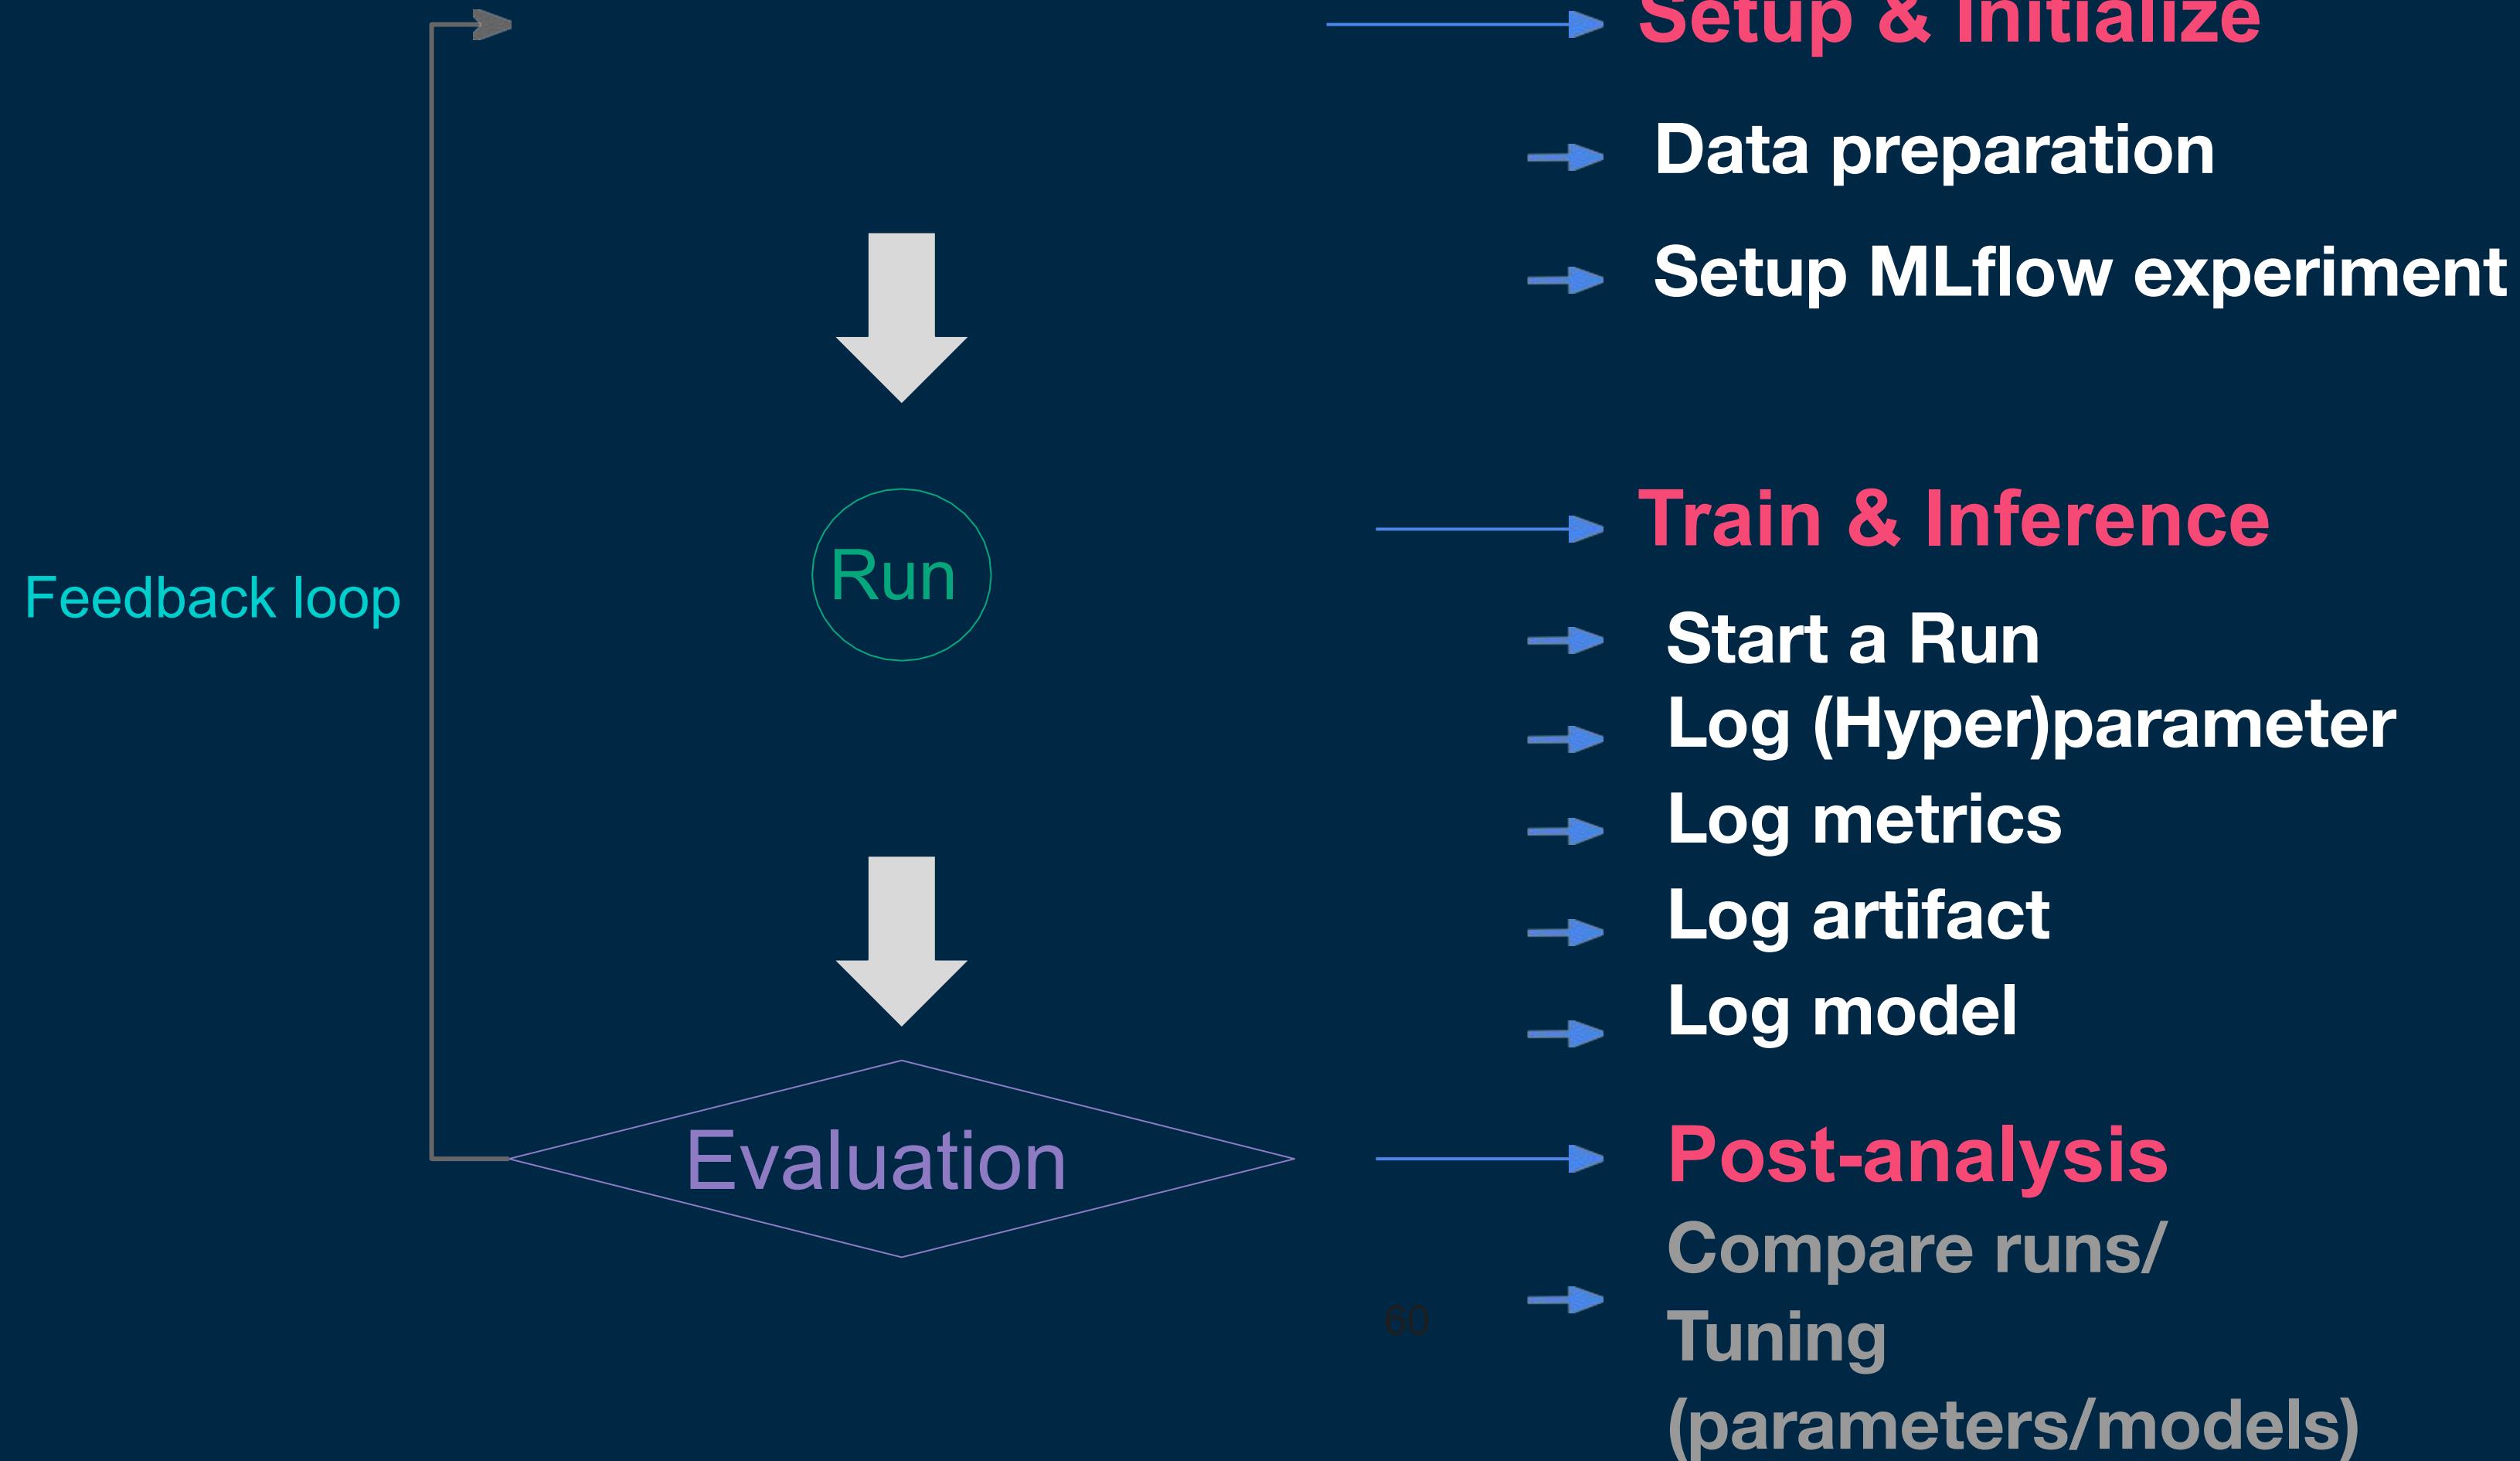
File storage

- Amazon S3
- Azure Blob Storage
- Google Cloud Storage
- FTP server
- SFTP Server
- NFS
- HDFS

Backend stores

- File store
- Database

Experiments



MLflow Tracking for ML Development

Tracking API

- start_Run()
- log_param()
- log_metric()
- log_artifact()
- end_Run()

Output

- Parameters
- Metrics
- Output file
 - Artifact
- Code version
- ...

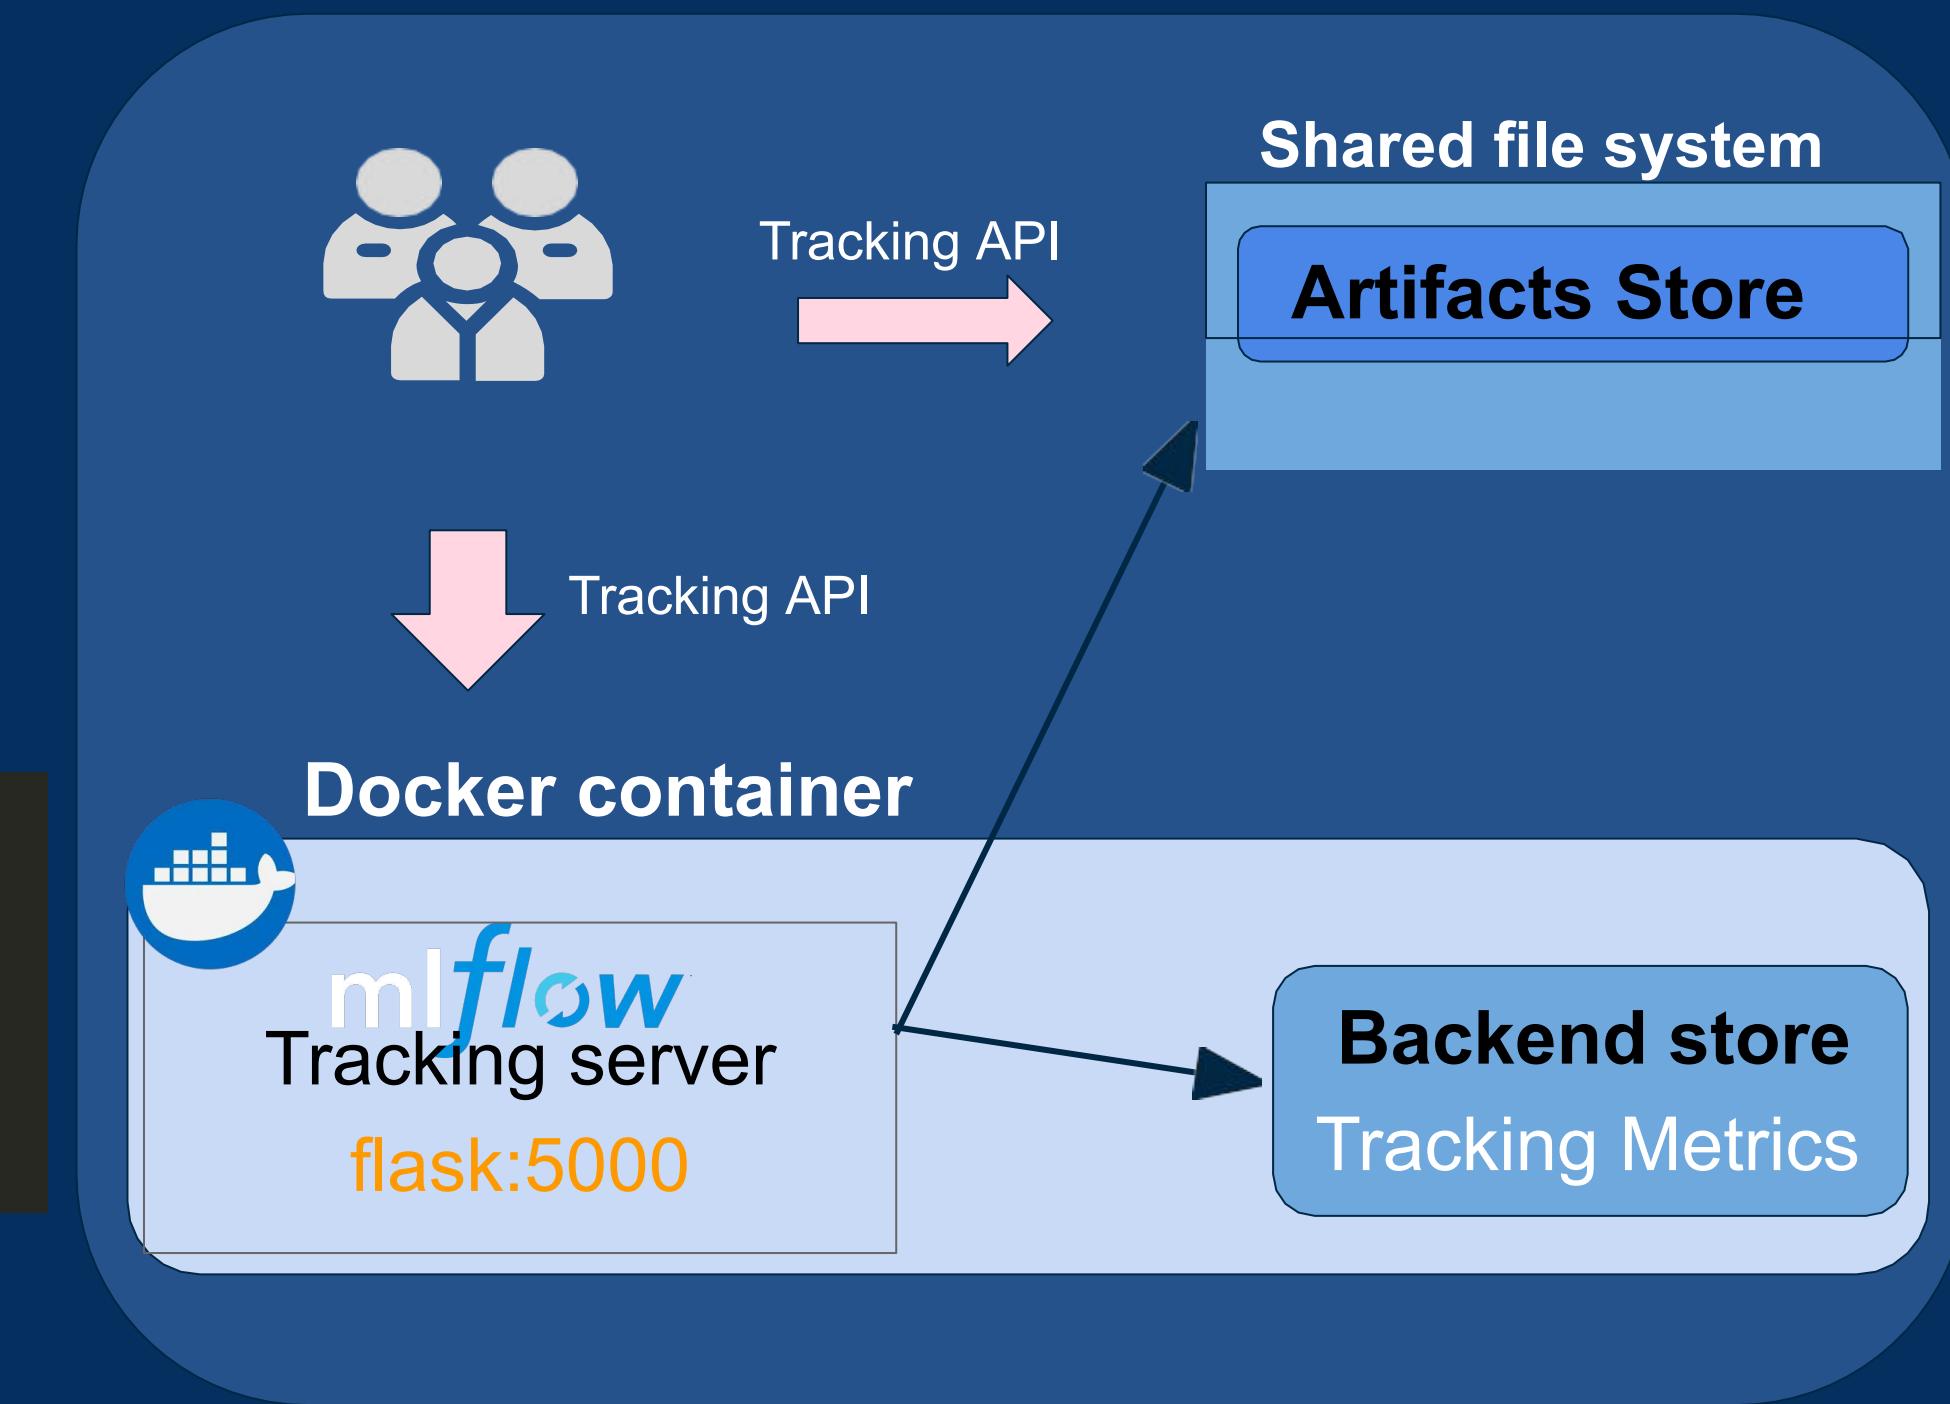
- log_model

• Model

Example Architecture

```
mlflow server \  
  --backend-store-uri $FILE_STORE \  
  --default-artifact-root $ARTIFACT_STORE \  
  --host $SERVER_HOST \  
  --port $SERVER_PORT
```

```
docker run -d -p 5000:5000 \  
  -v /tmp/artifactStore:/tmp/mlflow/artifactStore \  
  --name mlflow-tracking-server \  
  suci/mlflow-tracking
```



Tracking- Experiments

Setup & Initialize

```
# Setup & Initialize MLflow experiment  
experiment_name = "PyconTW 2020 Demo "  
tracking_server = "http://localhost:5000"  
  
mlflow.set_tracking_uri(tracking_server)  
mlflow.set_experiment(experiment_name)
```

#System Env setting

#backend-store-uri
\$FILE_STORE
#default-artifact-root
\$ARTIFACT_STORE
#host
\$SERVER_HOST
#port
\$SERVER_PORT

Tracking- Run

Train & Inference

```
with mlflow.start_run() as run:
```

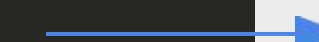
```
    # Log a parameter (key-value pair)
    log_param("param1", randint(0, 100))
```

```
    # Log a metric; metrics can be updated throughout the run
```

```
    log_metric("metricsA", random())
    log_metric("metricsA", random() + 1)
    log_metric("metricsA", random() + 2)
    log_metric("metricsB", random() + 2)
```

```
    # Log an artifact (output file)
```

```
    with open("outputs/test.txt", "w") as f:
        f.write("hello world! Run id:{}".format(type(mlflow.active_run().info)))
    log_artifacts("outputs")
```



Log (Hyper)parameter



Log metrics



Log artifact

Experiments



demo tracking api

experiment

Experiment ID : 0

Artifact Location : /tmp/mlflow/artifactStore/0

demo tracking api



▼ Notes

None

Search Runs: metrics.rmse < 1 and params.model = "tree" and tags.mlflow.source.type = "LOCAL"

? State:

Active ▾

Search

Clear

Showing 4 matching runs

Compare

Delete

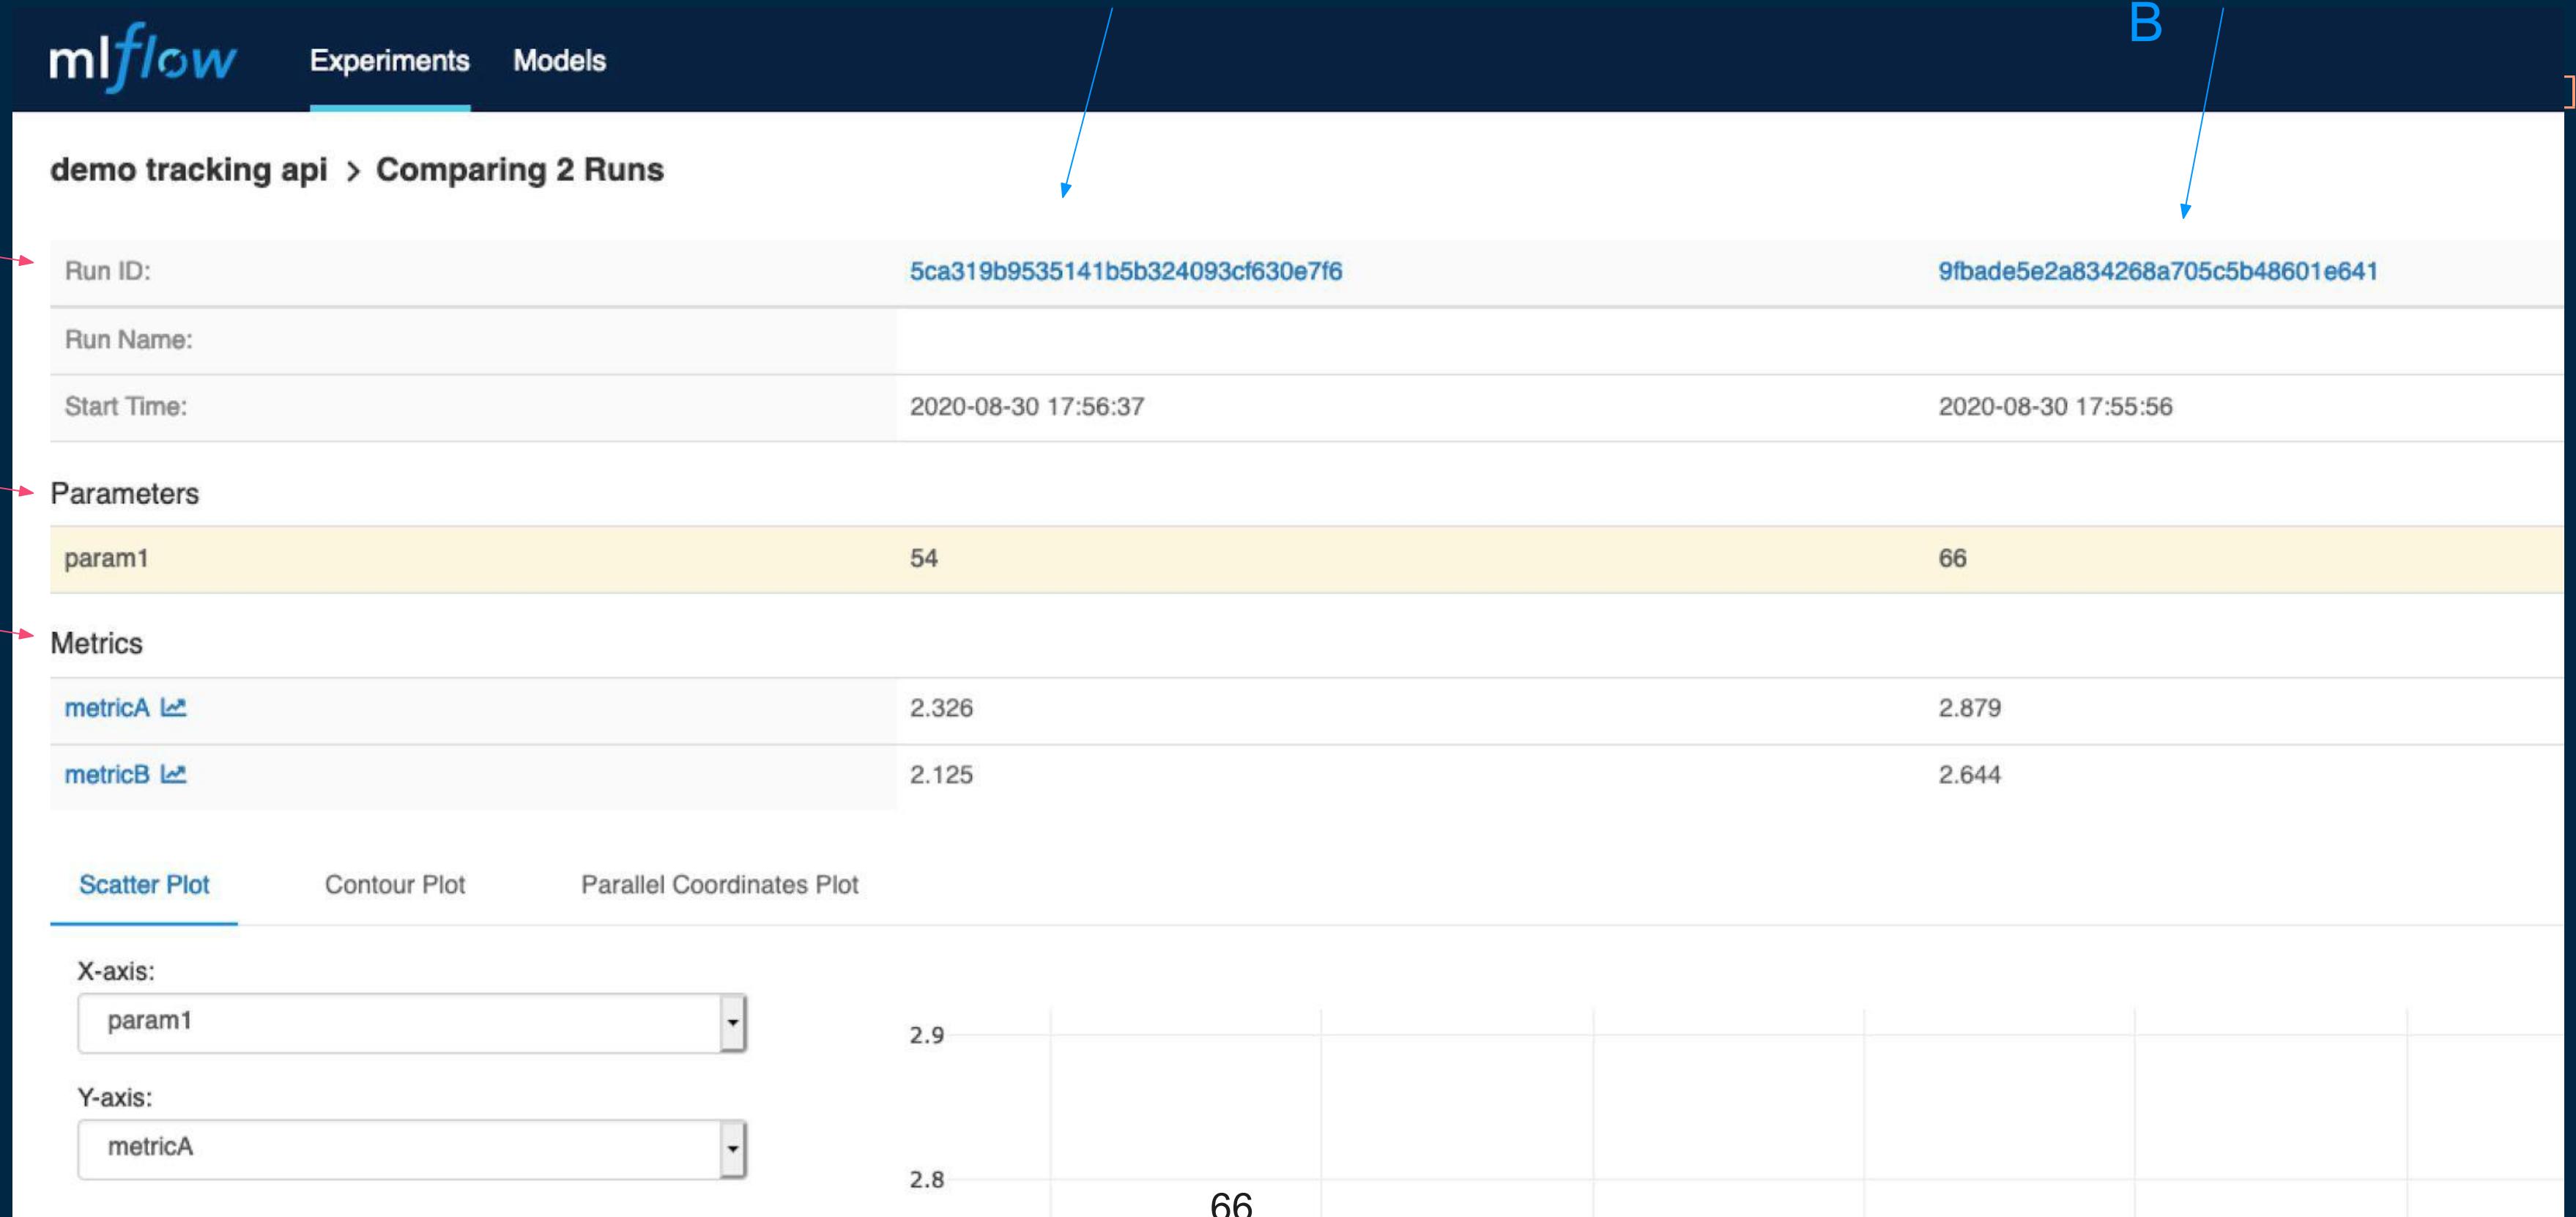
Download CSV



run

	Start Time	Run Name	User	Source	Version	Parameters	Metrics		
						param1	foo	metricA	metricB
<input type="checkbox"/>	2020-08-30 17:56:37	-	shuhs1	pycontw2020_d1	06be11	54	-	2.326	2.125
<input type="checkbox"/>	2020-08-30 17:55:56	-	shuhs1	pycontw2020_d1	06be11	66	-	2.879	2.644
<input type="checkbox"/>	2020-08-30 17:41:27	-	shuhs1	pycontw2020_d1	06be11	10	-	2.044	2.713
<input type="checkbox"/>	2020-08-30 12:09:21	-	shuhs1	pycontw2020_d1	06be11	18	2.468	-	-

Compare Two Runs



Compare Two Runs



demo tracking api > Comparing 2 Runs > metricA

Points: Off

Line Smoothness ②

1

X-axis:

Step

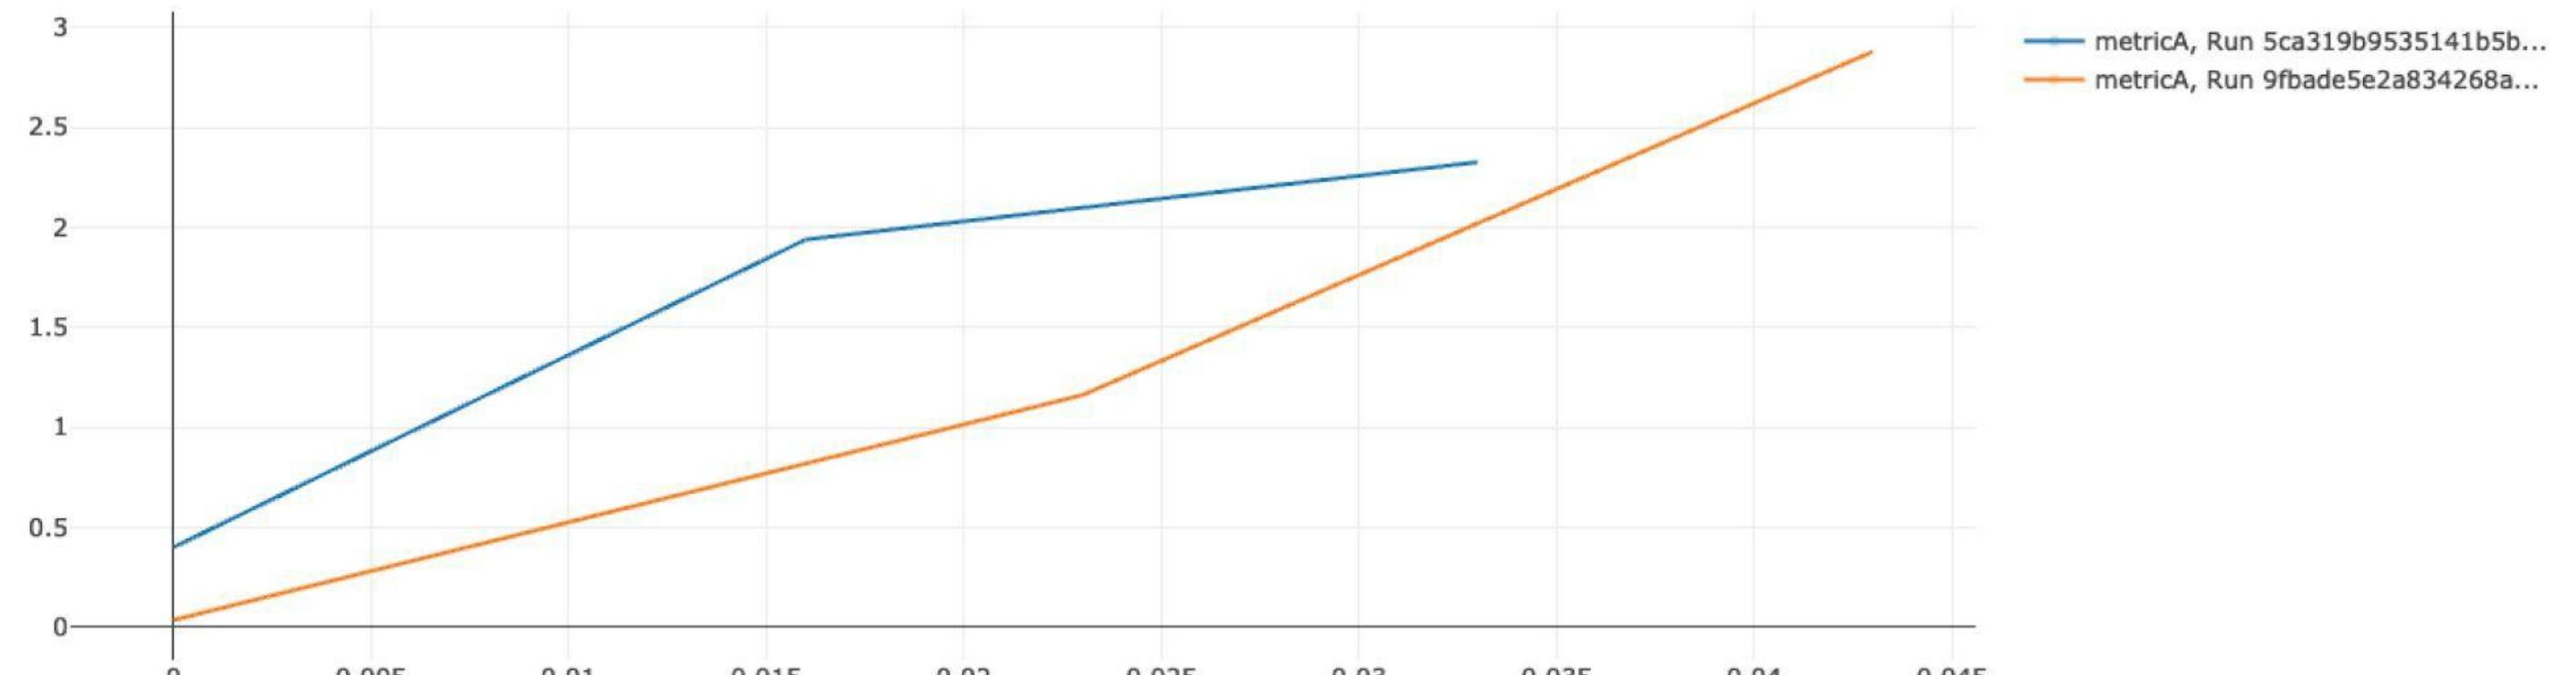
Time (Wall)

Time (Relative)

Y-axis:

metricA x

Y-axis Log Scale: Off



Model

Log->Load->deploy

```
mlflow.<model-type>.log_model(model, ...)
```

```
mlflow.<model-type>.load_model(modelpath)
```

```
mlflow.<model-type>.deploy()
```

```
mlflow.sklearn.log_model(lr, "model")
```

Built-In Model Flavors

<model-type>

- [Python Function \(python_function\)](#)
- [R Function \(crate\)](#)
- [H2O \(h2o\)](#)
- [Keras \(keras\)](#)
- [MLeap \(mleap\)](#)
- [PyTorch \(pytorch\)](#)
- [Scikit-learn \(sklearn\)](#)
- [Spark MLlib \(spark\)](#)
- [TensorFlow \(tensorflow\)](#)
- [ONNX \(onnx\)](#)
- [MXNet Gluon \(gluon\)](#)
- [XGBoost \(xgboost\)](#)
- [LightGBM \(lightgbm\)](#)
- [Spacy \(spaCy\)](#)
- [Fastai \(fastai\)](#)

▼ Artifacts

▼ model

MLmodel

conda.yaml

model.pkl



Full Path: /tmp/mlflow/artifactStore/4/557a222e5c8e454689a112cd...

Size: 349B



```
artifact_path: model
flavors:
  python_function:
    env: conda.yaml
    loader_module: mlflow.sklearn
    model_path: model.pkl
    python_version: 3.7.1
  sklearn:
    pickled_model: model.pkl
    serialization_format:云dpuclle
    sklearn_version: 0.23.2
run_id: 557a222e5c8e454689a112cd54dd7ff7
utc_time_created: '2020-09-04 00:21:11.388959'
```

▼ Artifacts

▼ model

MLmodel

conda.yaml

model.pkl



Full Path: /tmp/mlflow/artifactStore/4/557a222e5c8e454689a112cd...

Size: 150B

channels:

- defaults
- conda-forge

dependencies:

- python=3.7.1
- scikit-learn=0.23.2
- pip
- pip:
 - mlflow
 - cloudpickle==1.5.0

name: mlflow-env

Auto-Logging - TF

Manually logging

```
import mlflow  
  
mlflow.log_param("layers", layers)  
  
model = train_model()  
  
mlflow.log_metric("mse", model.mse())  
  
mlflow.log_artifact("plot", plot(model))  
  
mlflow.tensorflow.log_model(model)
```

With autologging

```
import mlflow  
  
mlflow.tensorflow.autolog()  
  
model = train_model()
```

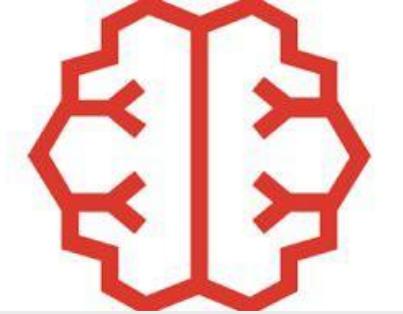
Capture TensorBoard metrics



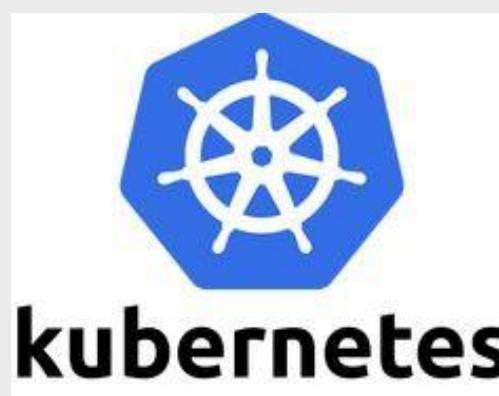
MLflow Tracking
MLflow Project
MLflow Models
MLflow Model registry
MLflow Deployment



Experiment tracking
Model deployment/serving
Model governance



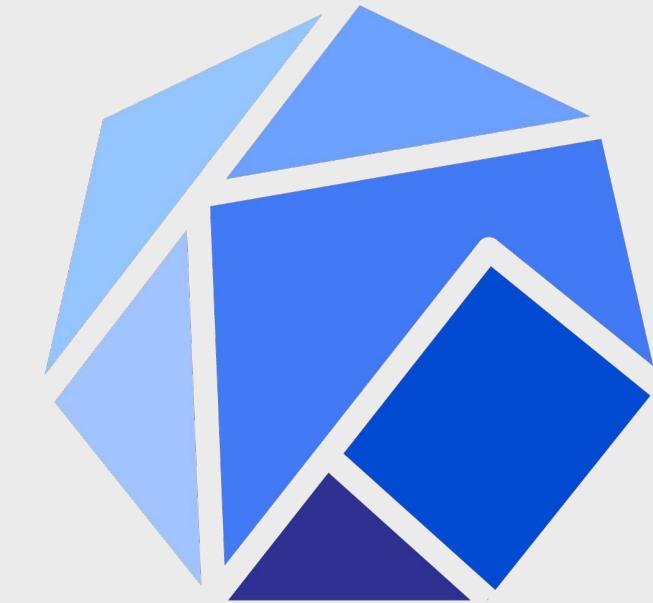
RedisAI



kubernetes

Kubeflow

- ML toolkit for Kubernetes
- Open-source and community-driven
- Support for multiple ML frameworks
- End-to-end workflows which can be shared, scaled and deployed



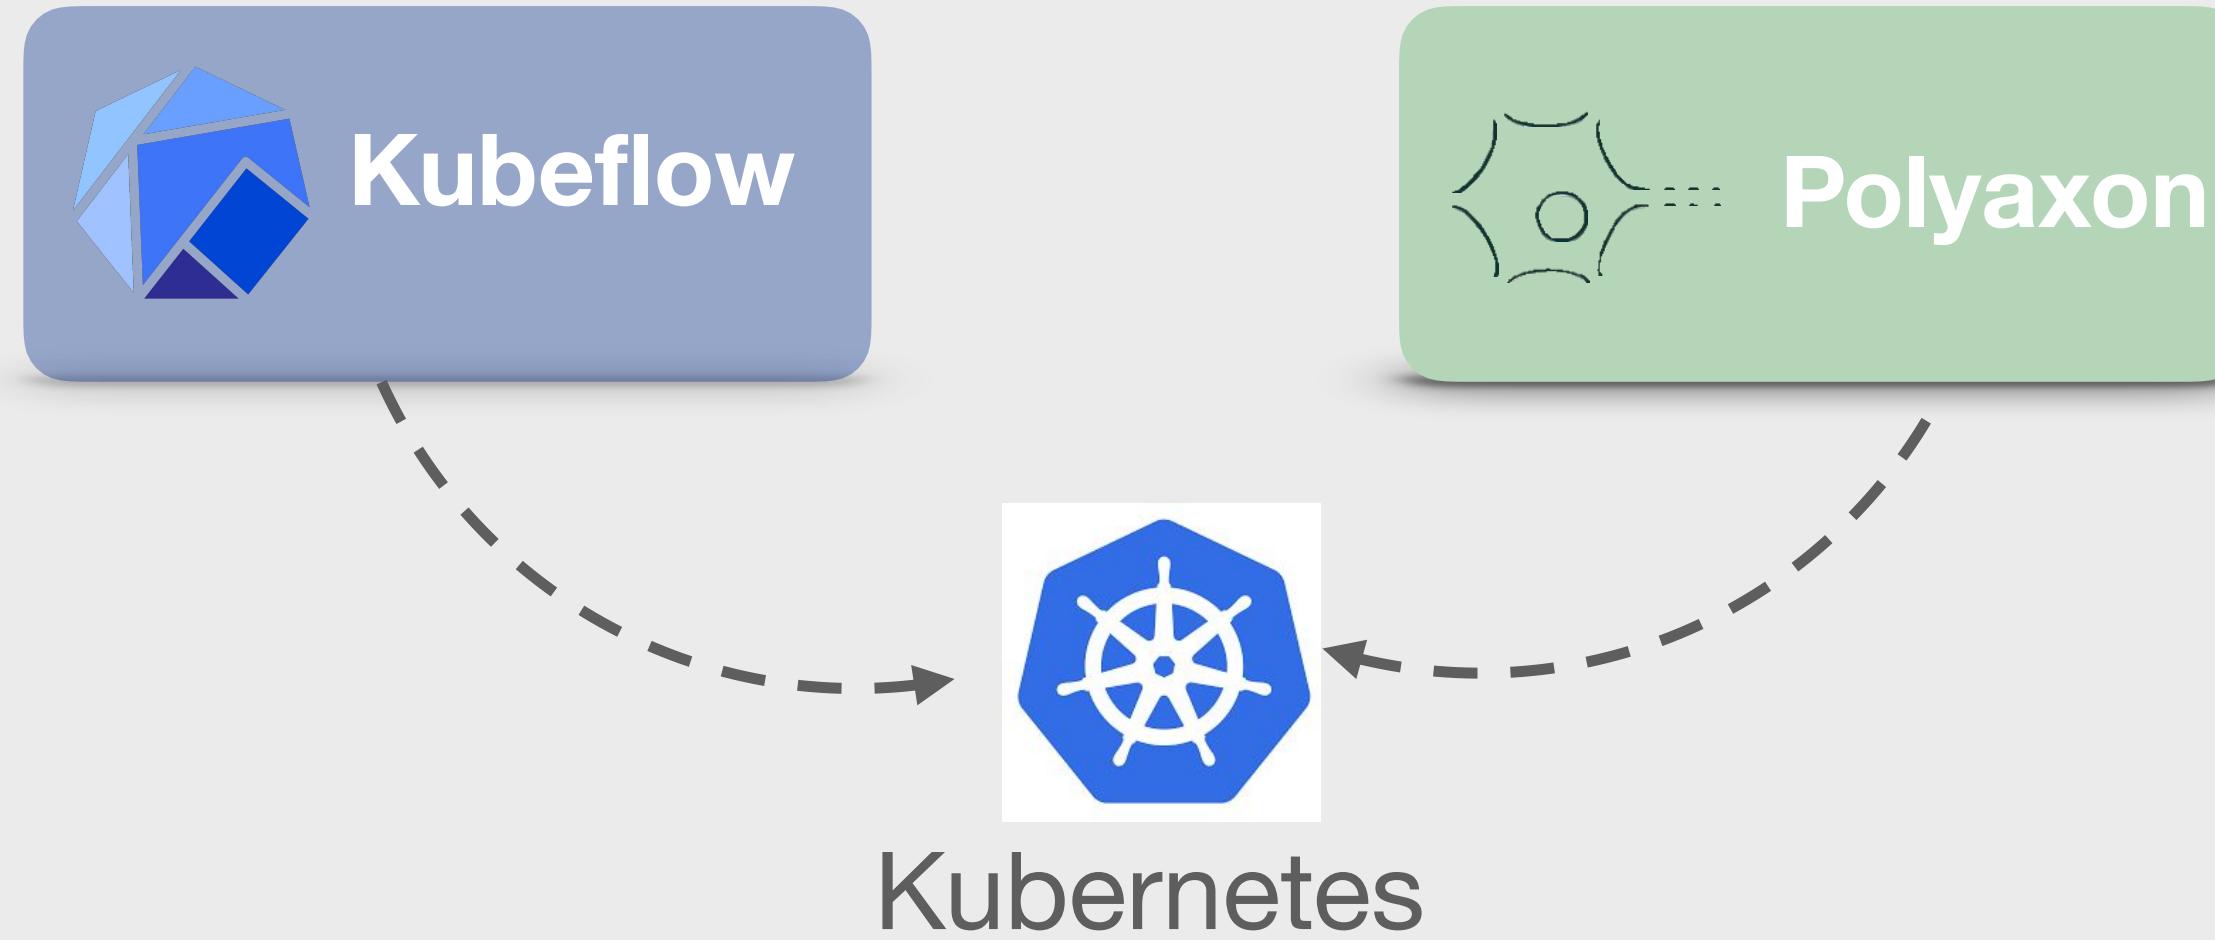
Kubeflow

Source:

<https://github.com/kubeflow/kubeflow/issues/187>

State of the Art Technologies

Support Machine Learning workloads in the Cloud



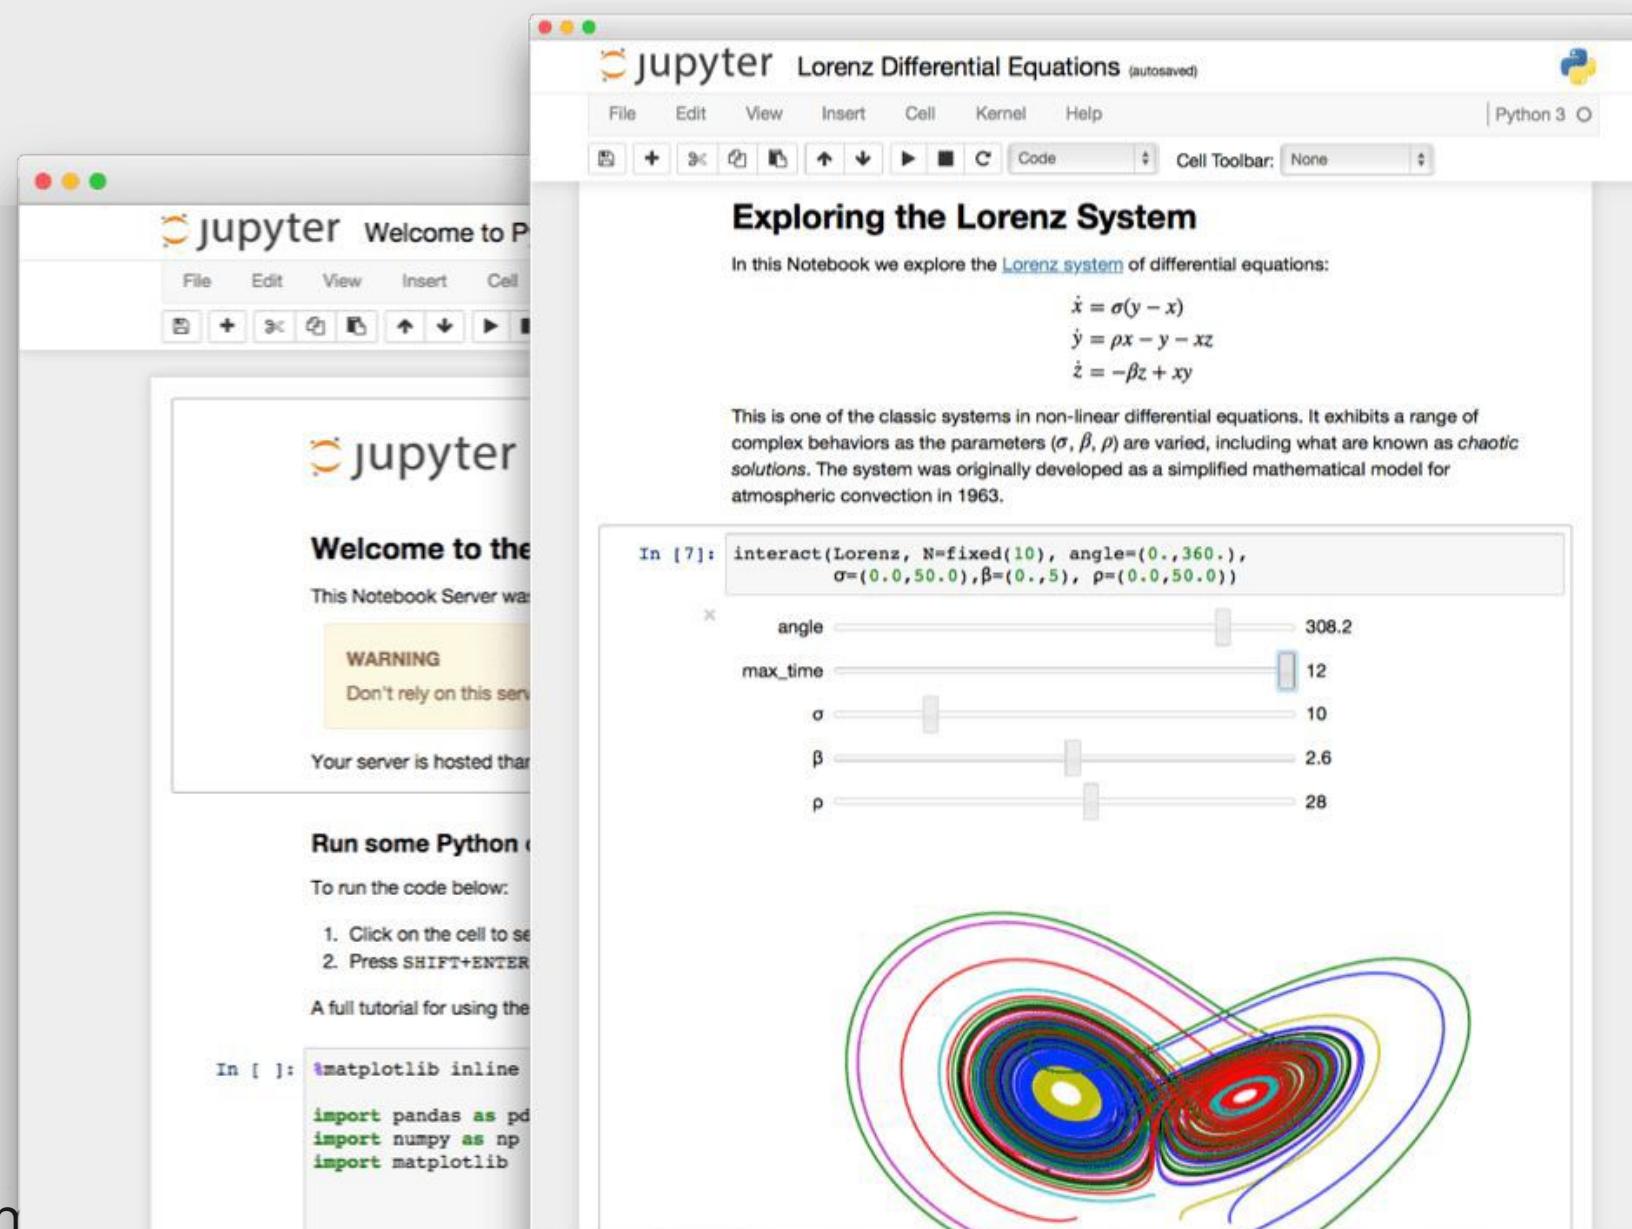
**Cloud
Agnostic**



Components

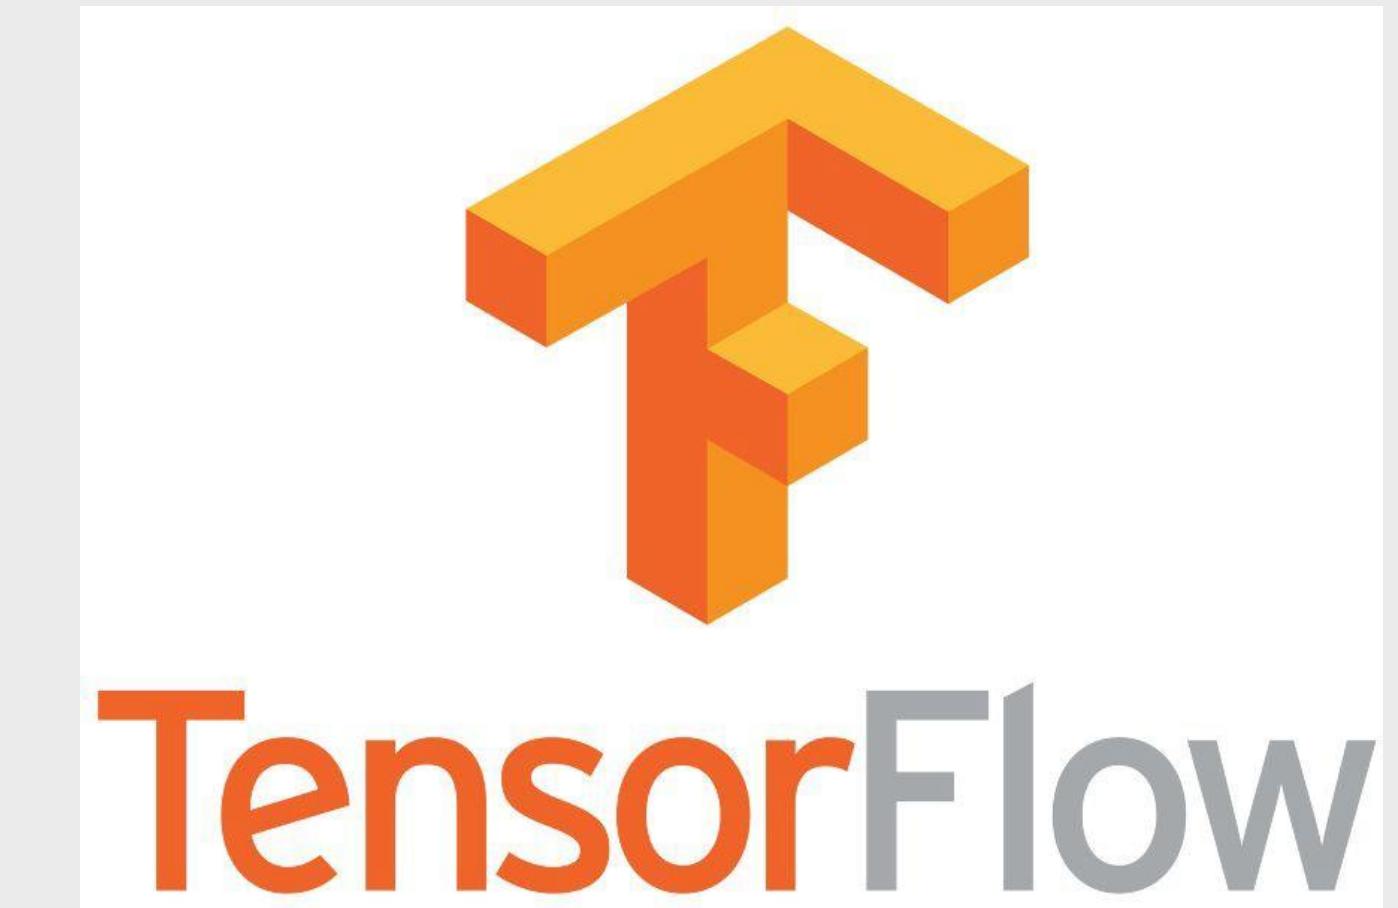
- Jupyterhub (collaboration and interactivity)
- K8s- native tensorflow controller (model building)
- K8s- native tensorflow serving deployment (model deployment)
- Ambassador (reverse proxy)
- Current and upcoming components for model tuning, model building and much more...
- Out-of-the-box setup for putting all of this together!

Jupyterhub



Tensorflow

- Open source numerical computing and ML
- Developed by Google, open-sourced in 2015
- Huge community and ecosystem
- Support for multiple ML models
- Tf-serving (model deployment), tensorboard (training visualization), etc.
- **Supports distributed training and deployment of models**



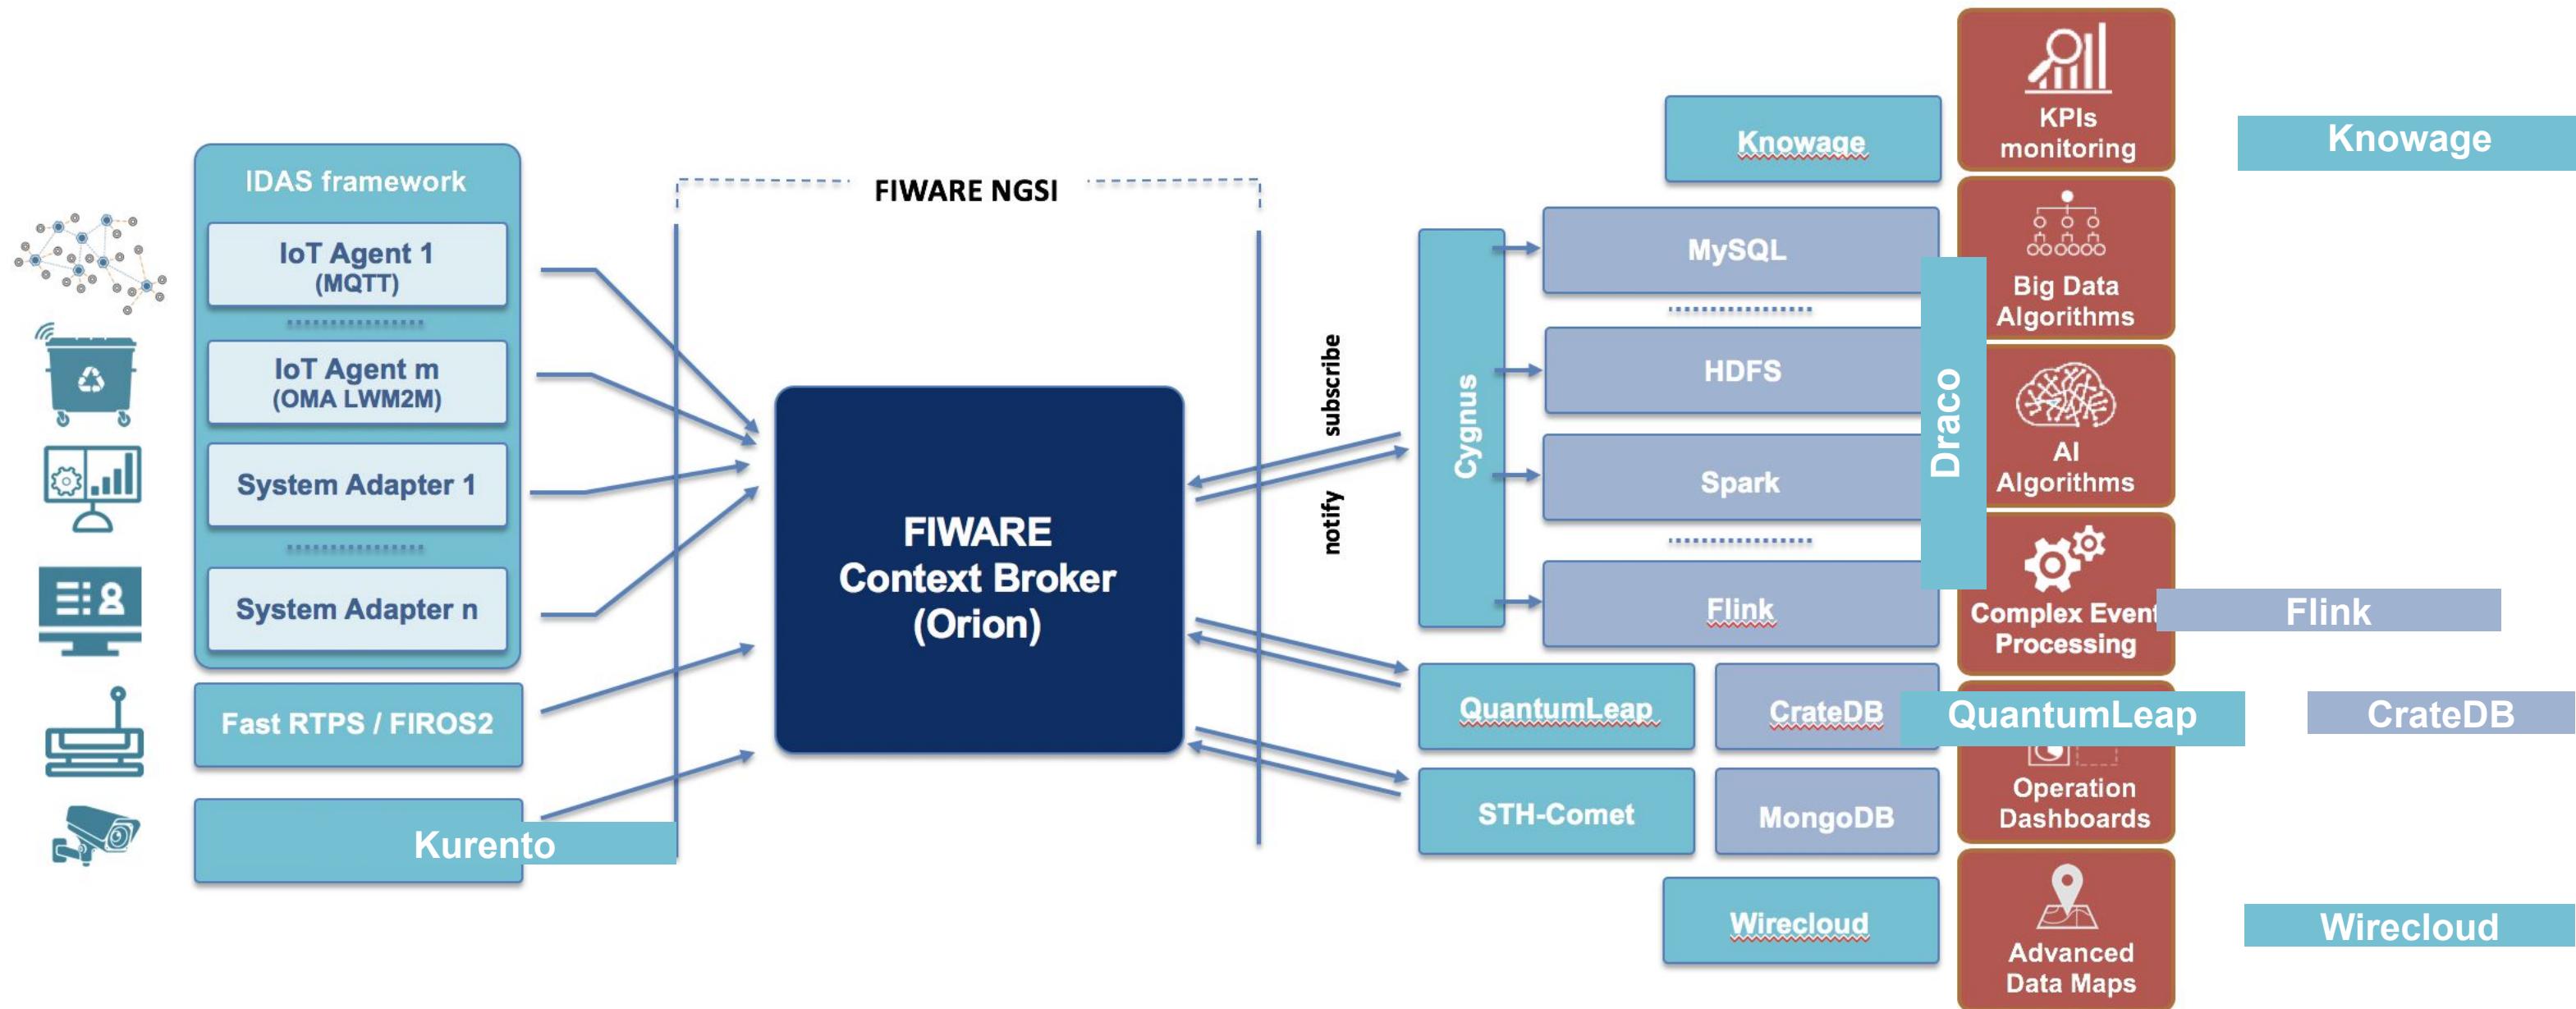
Why Kubeflow?

Based on current functionality you should consider using Kubeflow if:

- You want to train/serve TensorFlow models in different environments (e.g. local, on prem, and cloud)
- You want to use Jupyter notebooks to manage TensorFlow training jobs
- You want to launch training jobs that use resources – such as additional CPUs or GPUs – that aren't available on your personal computer
- You want to combine TensorFlow with other processes
 - For example, you may want to use [tensorflow/agents](#) to run simulations to generate data for training reinforcement learning models.

Refer <https://www.kubeflow.org/docs/started/getting-started/> for more info.

Simple Smart solutions: Reference Architecture



**FIWARE
Global
Summit**

**From Data
to Value**

OPEN SOURCE
OPEN STANDARDS
OPEN COMMUNITY

Draco: Persisting Context Data to MongoDB

José Andrés Muñoz Arcentales, Javier Conde Díaz, Joaquín
Salvachúa

Vienna, Austria

12-13 June, 2023

#FIWARESummit



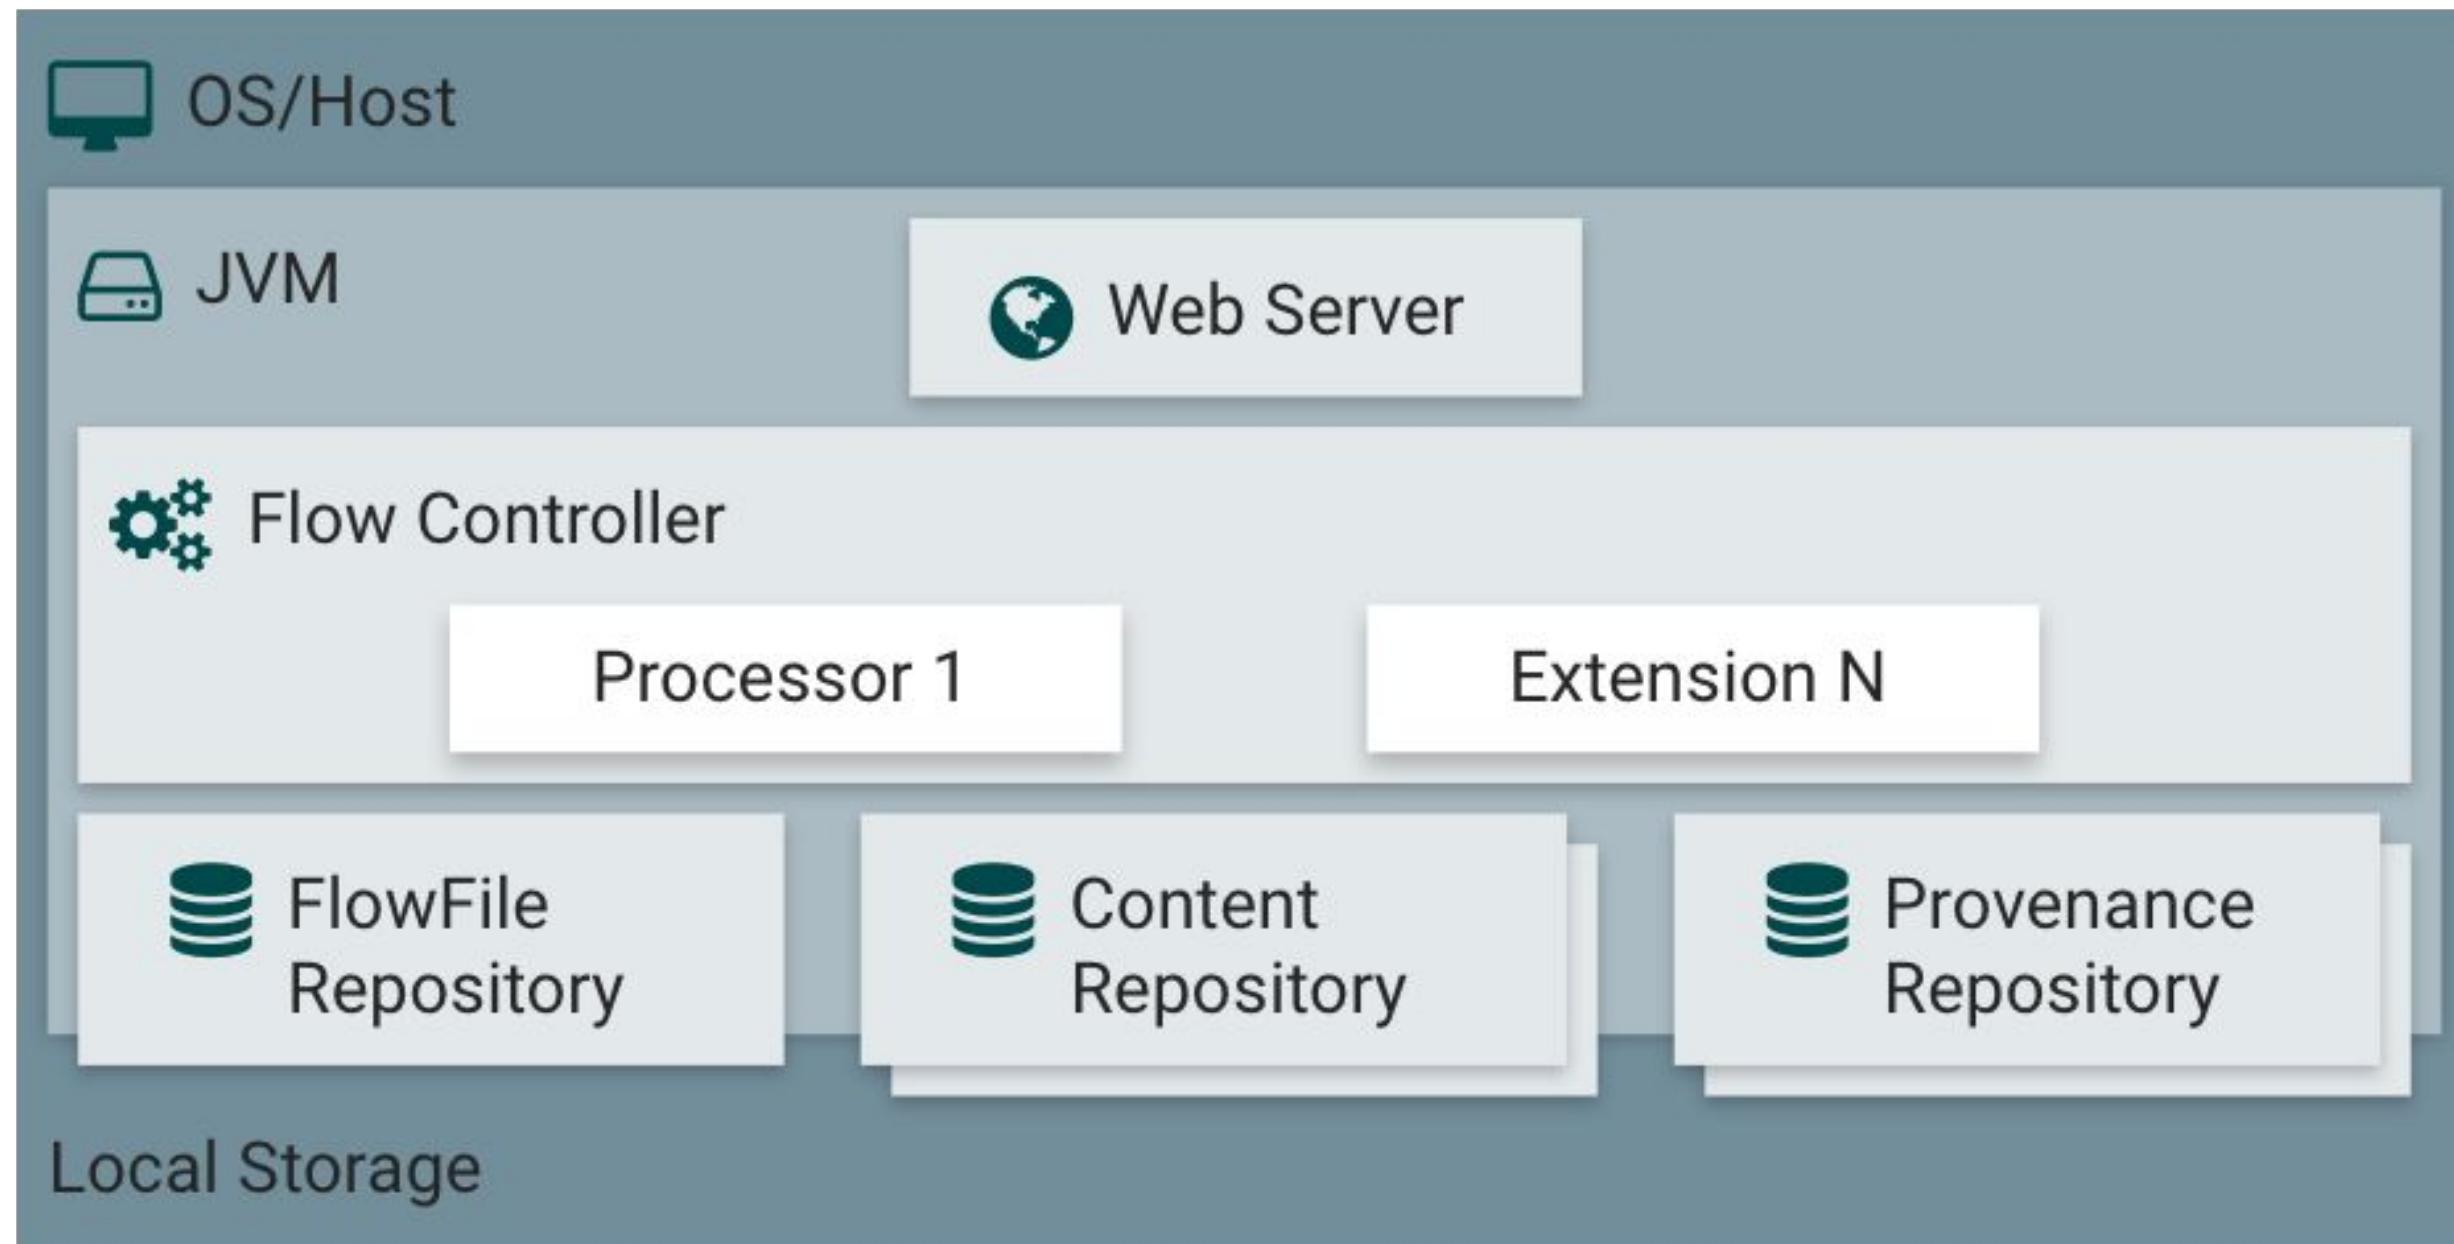


The Draco GE

- The Draco Generic Enabler takes care of the data ingestion and persistence. It is an easy to use, powerful, and reliable system for processing and distributing data. Internally, Draco is based on Apache NiFi.
- NiFi is a dataflow system based on the concepts of flow-based programming. It supports powerful and scalable directed graphs of data routing, transformation, and system mediation logic. It was built to automate the flow of data between systems.



Apache Nifi Architecture





Draco integration in the FIWARE ecosystem



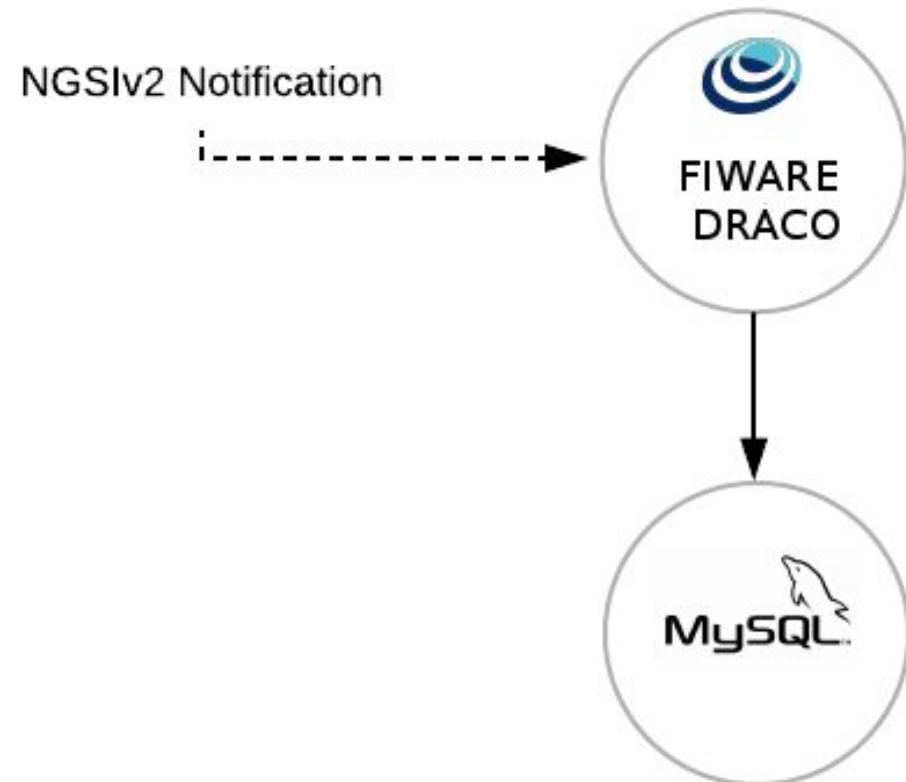
MiniNiFI (low profile version)



Features

- **Based on Apache NiFi.**
- **NGSI 2 Support both for ingestion and serialization to have full integration with the Orion Context Broker.**
- **Several persistent backends :**
 - **MySQL**, the well-know relational database manager.
 - **MongoDB**, the NoSQL document-oriented database.
 - **PostgreSQL**, the well-know relational database manager.
 - **HDFS**, Hadoop distributed file system.
 - **Cassandra**, Distributed database.
 - **CartoDB**, for geospatial Data
- **Templates for some common scenarios**
- **Rest API**

Basic Example



The screenshot shows the Apache NiFi user interface at the URL 127.0.0.1:9090/nifi/. The top navigation bar includes links for Home, Navigations, Process Groups, Schedules, Monitors, Metrics, and Help, along with a search bar. The main workspace displays a process group named "85e743f4-0165-1000-6a73-a1b37b01cd4e".

The process flow consists of three main components:

- NGSIV2-HTTP-INPUT**: A ListenHTTP processor with the following metrics:
 - In: 0 (0 bytes)
 - Read/Write: 0 bytes / 0 bytes
 - Out: 0 (0 bytes)
 - Tasks/Time: 0 / 00:00:00.000
- NGSIToMySQL**: An NGSIToMySQL processor with the following metrics:
 - In: 0 (0 bytes)
 - Read/Write: 0 bytes / 0 bytes
 - Out: 0 (0 bytes)
 - Tasks/Time: 0 / 00:00:00.000
- LogAttribute**: A LogAttribute processor with the following metrics:
 - In: 0 (0 bytes)
 - Read/Write: 0 bytes / 0 bytes
 - Out: 0 (0 bytes)
 - Tasks/Time: 0 / 00:00:00.000

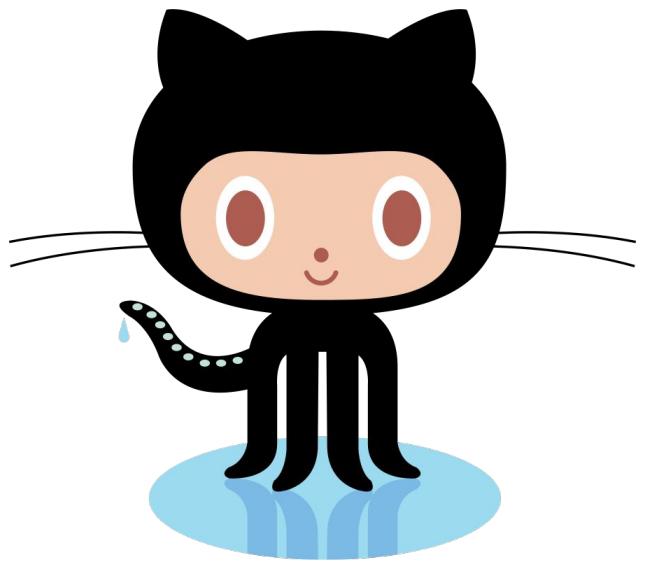
Flows connect the components in sequence: NGSIV2-HTTP-INPUT → NGSIToMySQL → LogAttribute. Each flow is labeled "Name success" and "Queued 0 (0 bytes)".

Demo



Get the code!

<https://github.com/ging/fiware-draco-data-ingestion-supermarket>



**FIWARE
Global
Summit**

**From Data
to Value**

OPEN SOURCE
OPEN STANDARDS
OPEN COMMUNITY

Cosmos: Loading Streaming Data using Flink and Spark

José Andrés Muñoz Arcentales, Javier Conde Díaz, Joaquín
Salvachúa

Vienna, Austria

12-13 June, 2023

#FIWARESummit

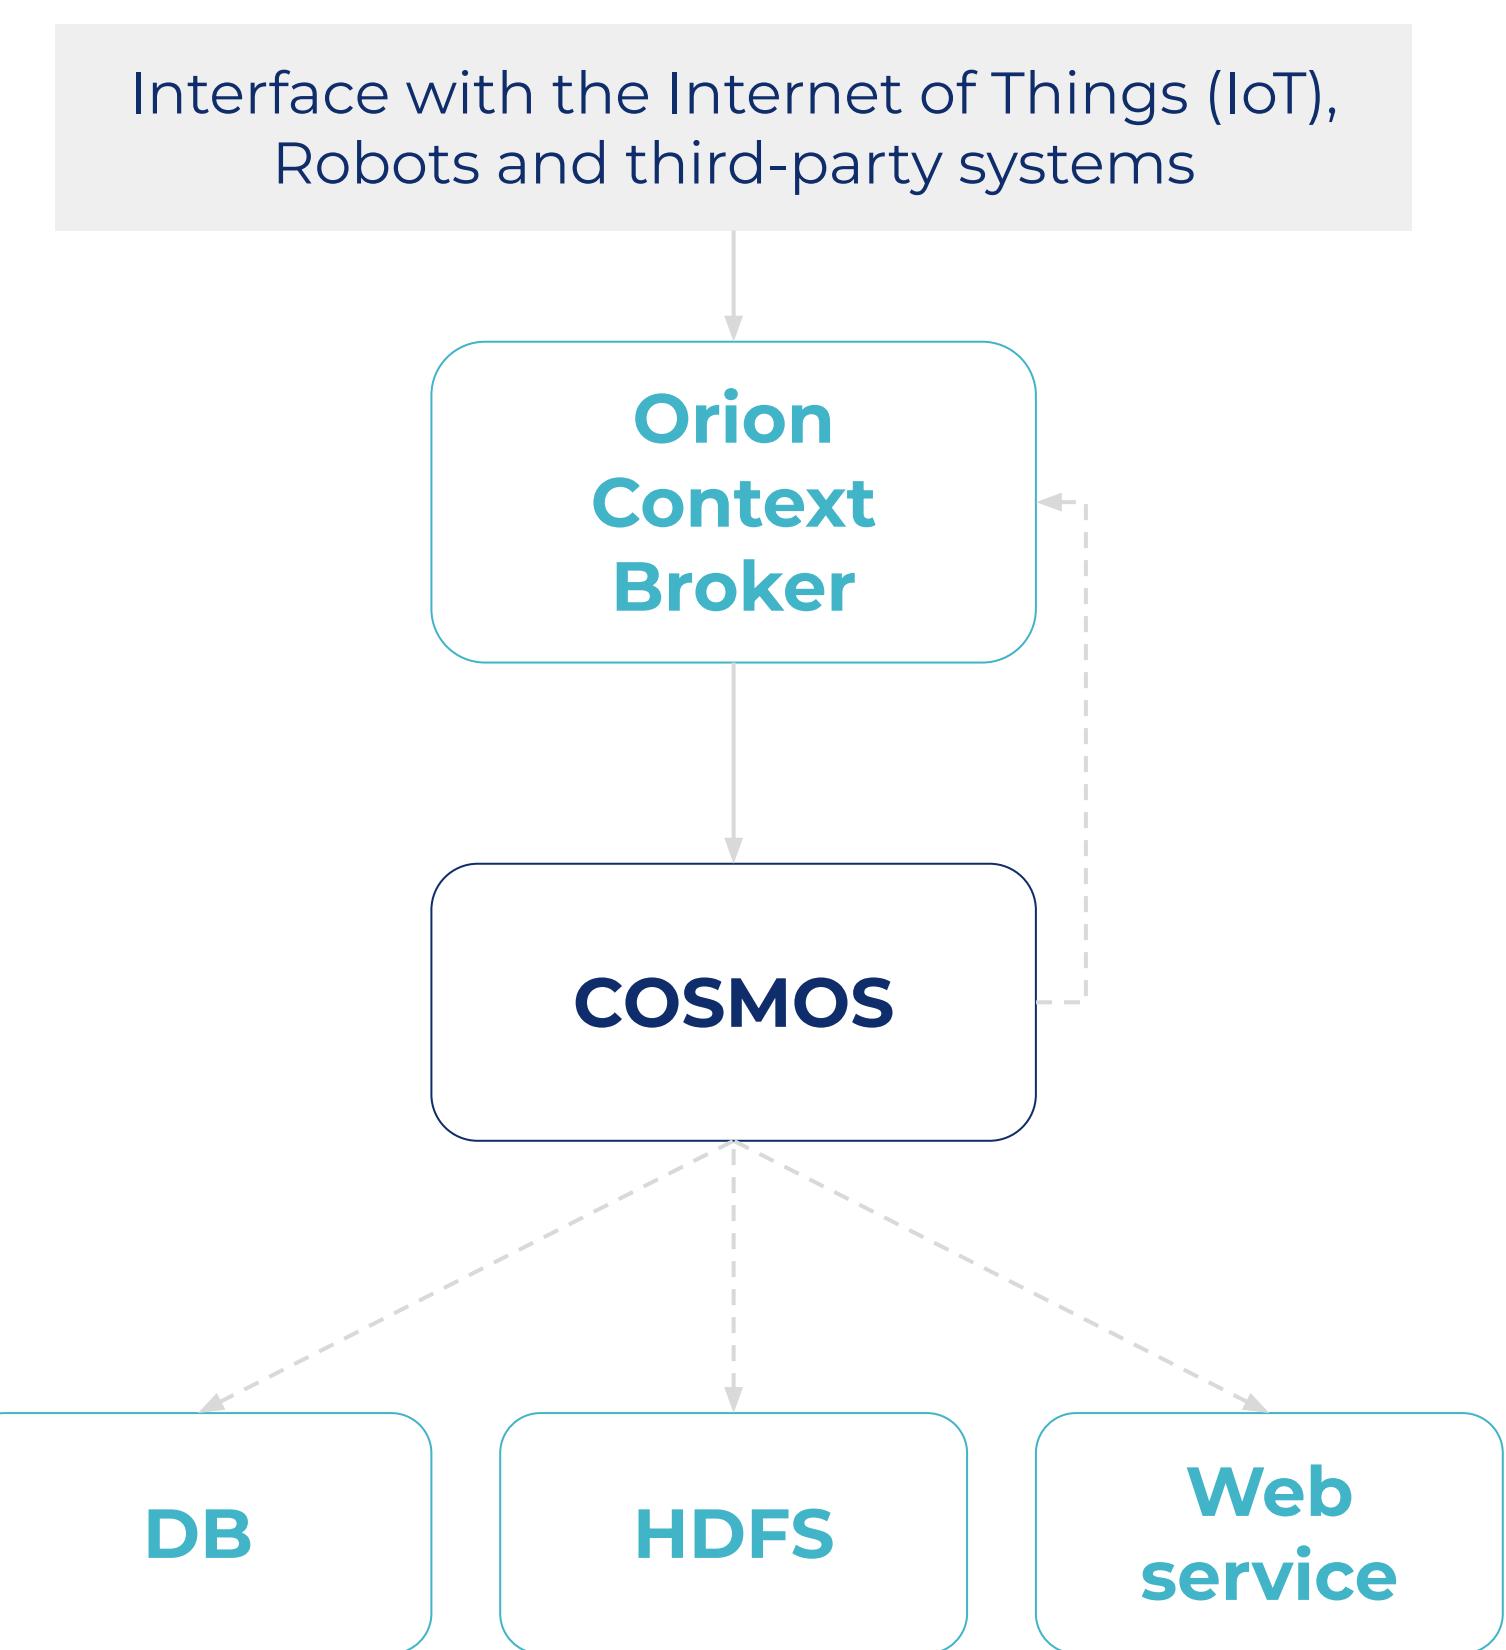


The Cosmos GE

The Cosmos Generic Enabler enables an easier BigData analysis over context integrated with some of the most popular BigData platforms.

Features

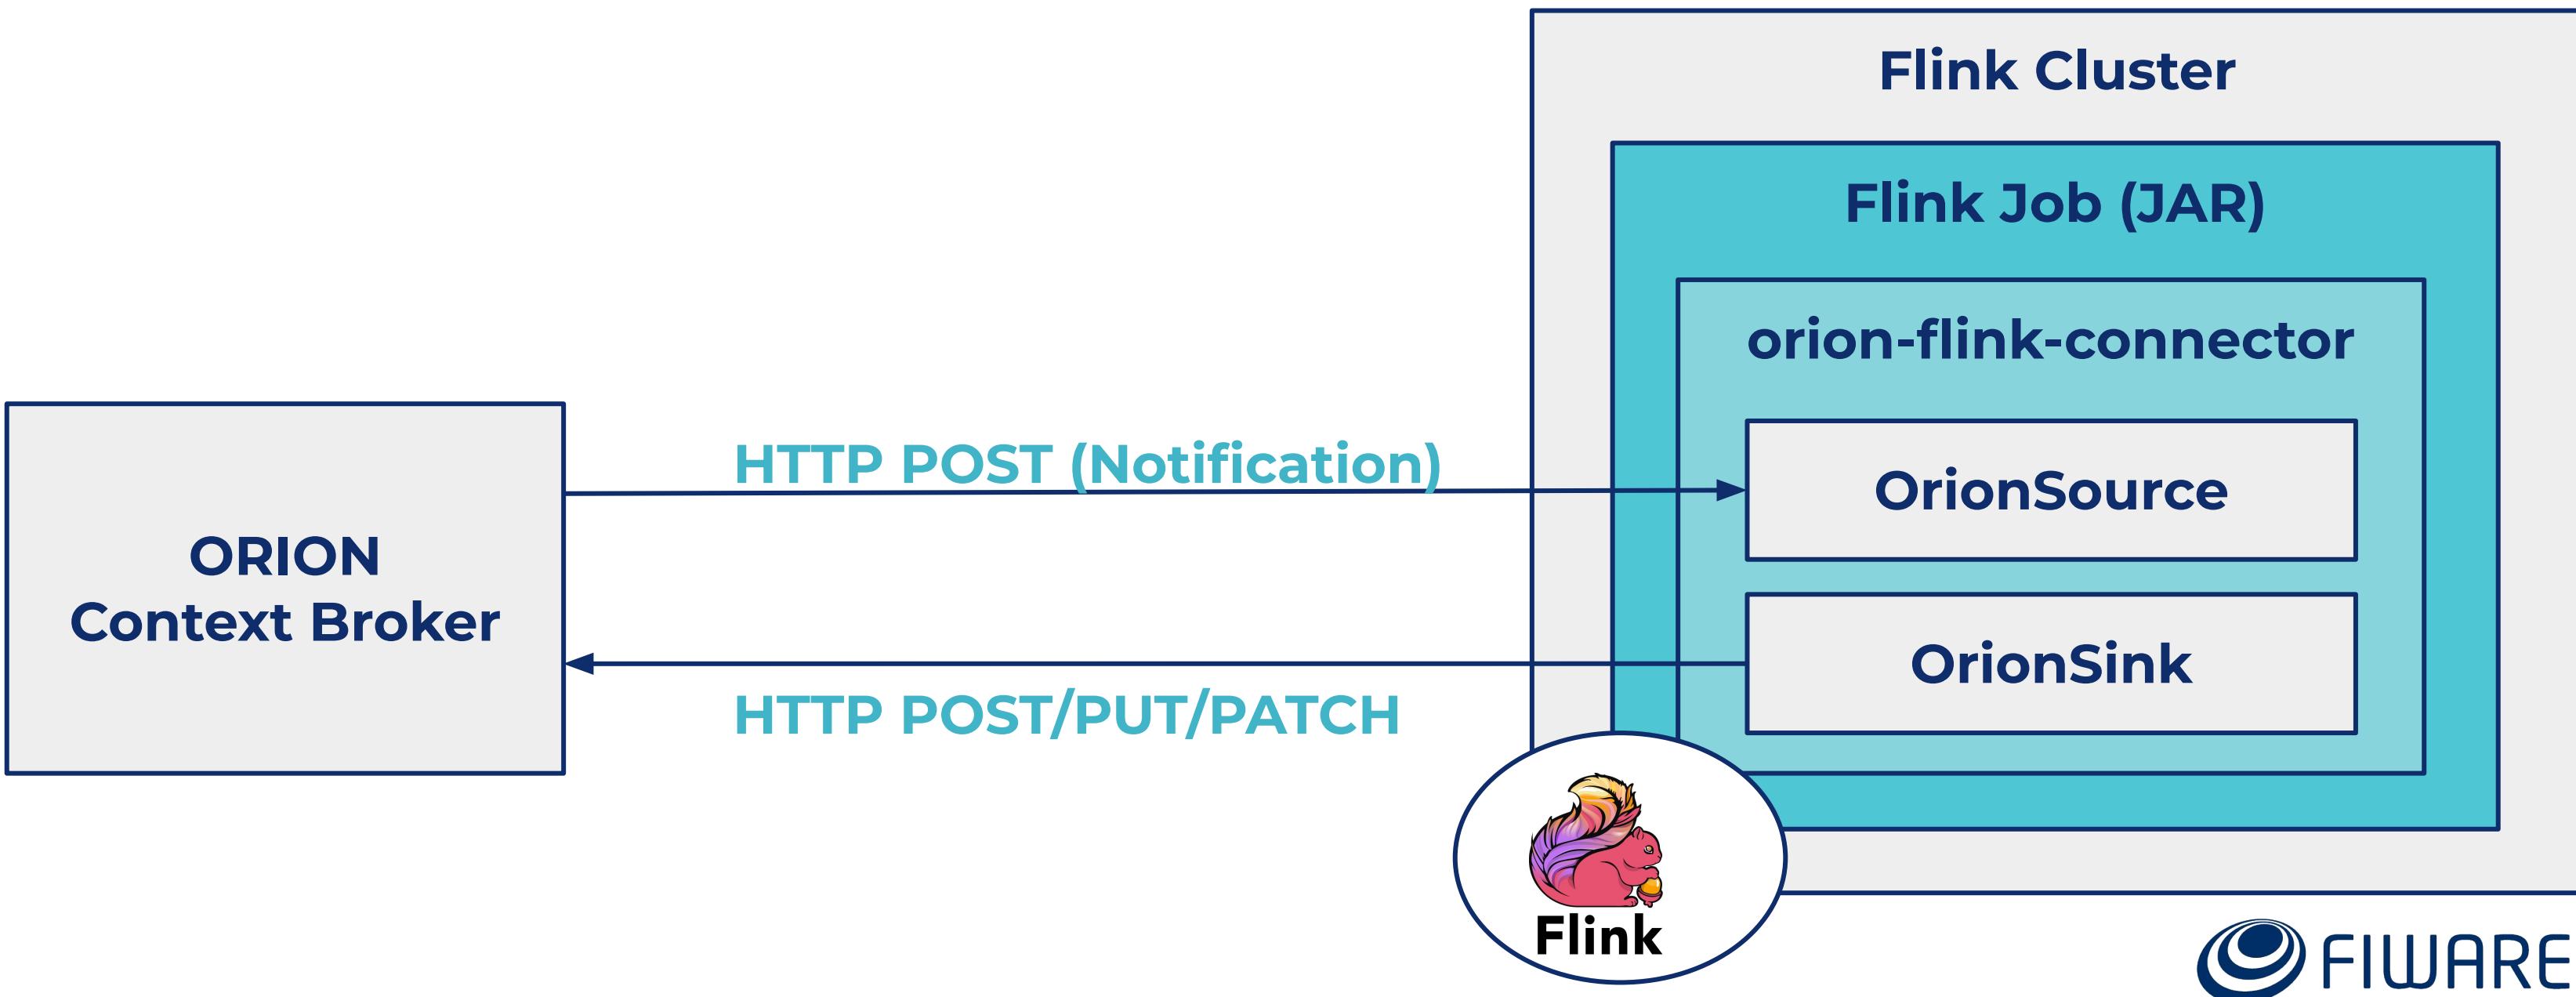
- ✓ Batch Processing
- ✓ Stream Processing (Real-time)
- ✓ Direct data ingestion
- ✓ Direct connection with Orion
- ✓ Multiple Sinks



<https://github.com/ging/fiware-cosmos>

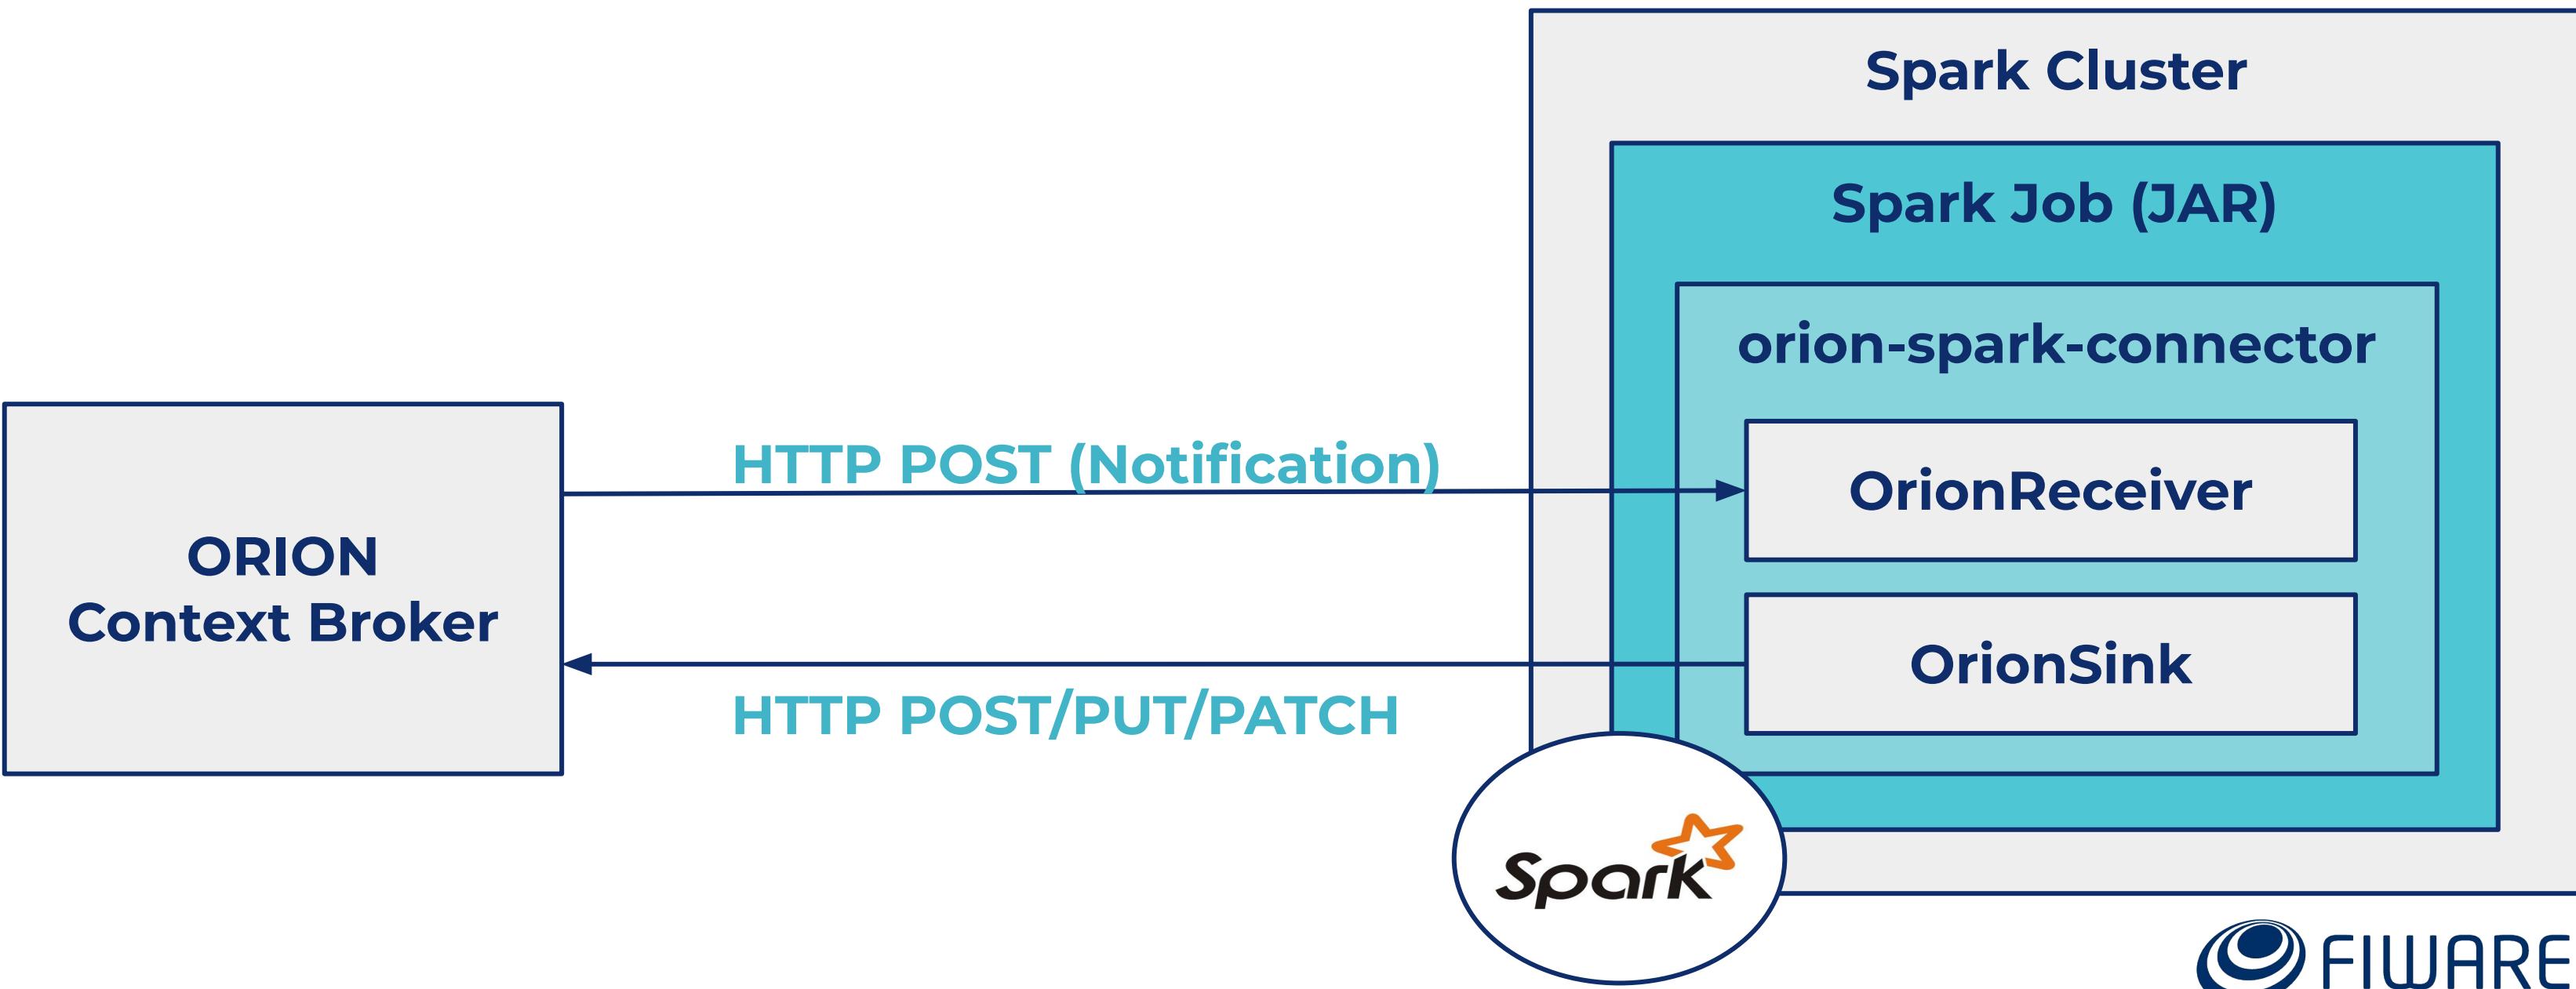
fiware-cosmos-orion-flink-connector

- ⌚ <https://github.com/ging/fiware-cosmos-orion-flink-connector>
- ⌚ <https://github.com/ging/fiware-cosmos-orion-flink-connector-examples>



fiware-cosmos-orion-spark-connector

- ⌚ <https://github.com/ging/fiware-cosmos-orion-spark-connector>
- ⌚ <https://github.com/ging/fiware-cosmos-orion-spark-connector-examples>



Demo: Average temperature for each entity

```
def main(args: Array[String]): Unit = {
  val env = StreamExecutionEnvironment.getExecutionEnvironment

  // Create Orion Source. Receive notifications on port 9001
  val eventStream = env.addSource(new OrionSource(9001))

  // Process event stream
  val processedDataStream = eventStream
    .flatMap(event => event.entities)
    .map(entity => {
      val temp = entity.attrs("temperature").value.asInstanceOf[Number].floatValue()
      (entity.id, temp)
    })
    .keyBy(0)
    .timeWindow(Time.seconds(10))
    .aggregate(new Average)

  // print the results with a single thread, rather than in parallel
  processedDataStream.print().setParallelism(1)
```



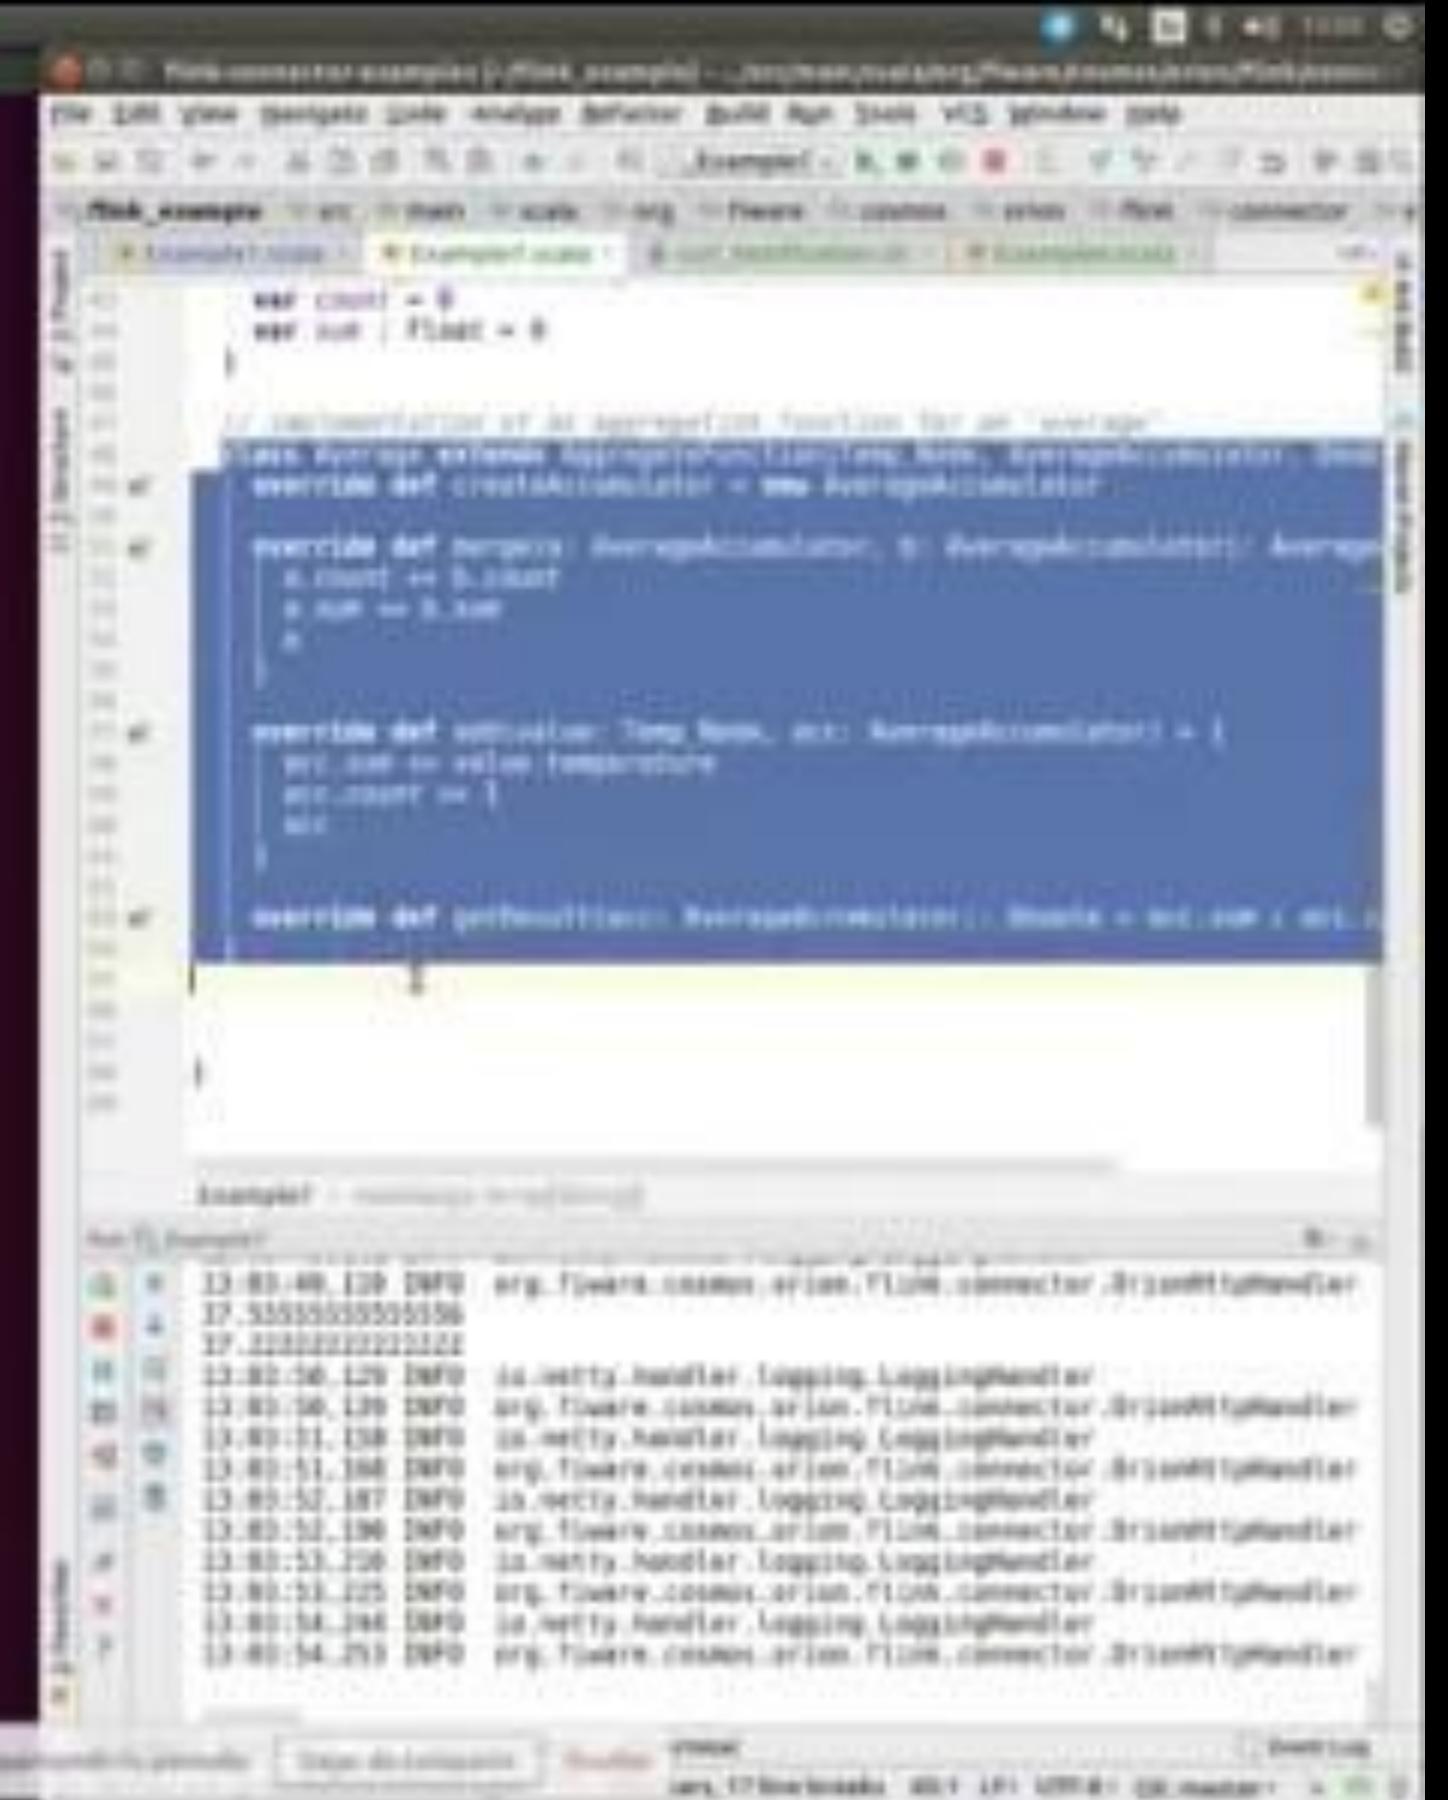
```
    -> Cloning connection: 8
curl: (3) Could not resolve host: www
-> Rebuilt URL: http://localhost:8080/
      trying 127.0.0.1...
-> HTTP_ACCEPT: text/html
-> Connected to localhost (127.0.0.1) port 8080 (#0)
-> POST / HTTP/1.1
-> Host: localhost:8080
-> Content-Type: application/x-www-form-urlencoded
-> Accept: application/html
-> User-Agent: curl/7.28.0
-> Fluence-Service: demo
-> Fluence-ServicePath: /test
-> Content-Length: 1148
-> Expect: 200-continue

-> HTTP/1.1 200 continue
=> we are completely uploaded and fine
-> HTTP/1.1 200 continue
-> HTTP/1.1 200 ok
-> Content-Type: text/plain
-> Content-Length: 0

-> Connection #0 to host localhost left intact
-> Rebuilt URL: http:///
-> Could not resolve host: /
-> Cloning connection: 8
curl: (3) Could not resolve host: /
-> Rebuilt URL: http://
-> Could not resolve host: http://
-> Cloning connection: 8
curl: (3) Could not resolve host: http://
-> Rebuilt URL: http://localhost:8080/
      trying 127.0.0.1...
-> HTTP_ACCEPT: text/html
-> Host: localhost:8080
-> Content-Type: application/x-www-form-urlencoded
-> Accept: application/html
-> User-Agent: curl/7.28.0
-> Fluence-Service: demo
-> Fluence-ServicePath: /test
-> Content-Length: 1148
-> Expect: 200-continue

-> HTTP/1.1 200 continue
=> we are completely uploaded and fine
-> HTTP/1.1 200 ok
-> Content-Type: text/plain
-> Content-Length: 0

-> Connection #0 to host localhost left intact
```



Current status

	 Flink	 Spark
Orion Connector Orion Source/Receiver + Orion Sink	✓	✓
RTD Documentation	✓	✓
Unit Tests	✓	✓
Examples	✓	✓
Step-by-step tutorial	✓	
Support NGSI LD		

**FIWARE
Global
Summit**

**From Data
to Value**

OPEN SOURCE
OPEN STANDARDS
OPEN COMMUNITY

Inference system with ML

José Andrés Muñoz Arcentales, Javier Conde Díaz, Joaquín Salvachúa

Vienna, Austria

12-13 June, 2023

#FIWARESummit



Predicting supermarket purchases

An use case



Purchase data

How we model each purchase:

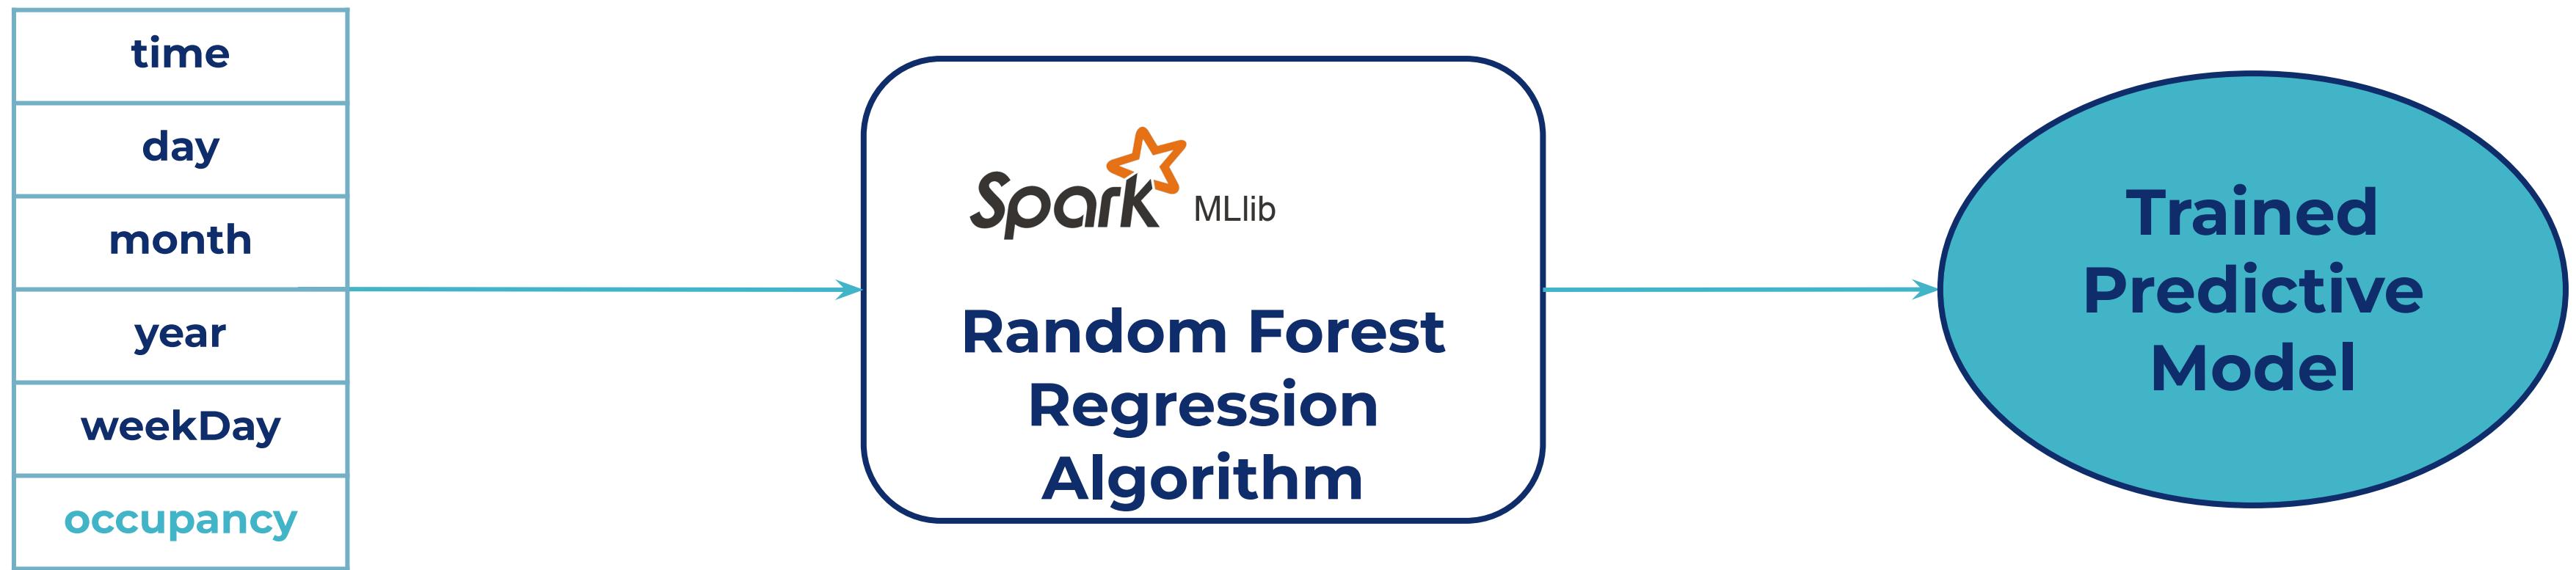
- **date**
- **client_id**
- **supermarket_id**
- **product_list**
 - **description**
 - **n_items**
 - **price**



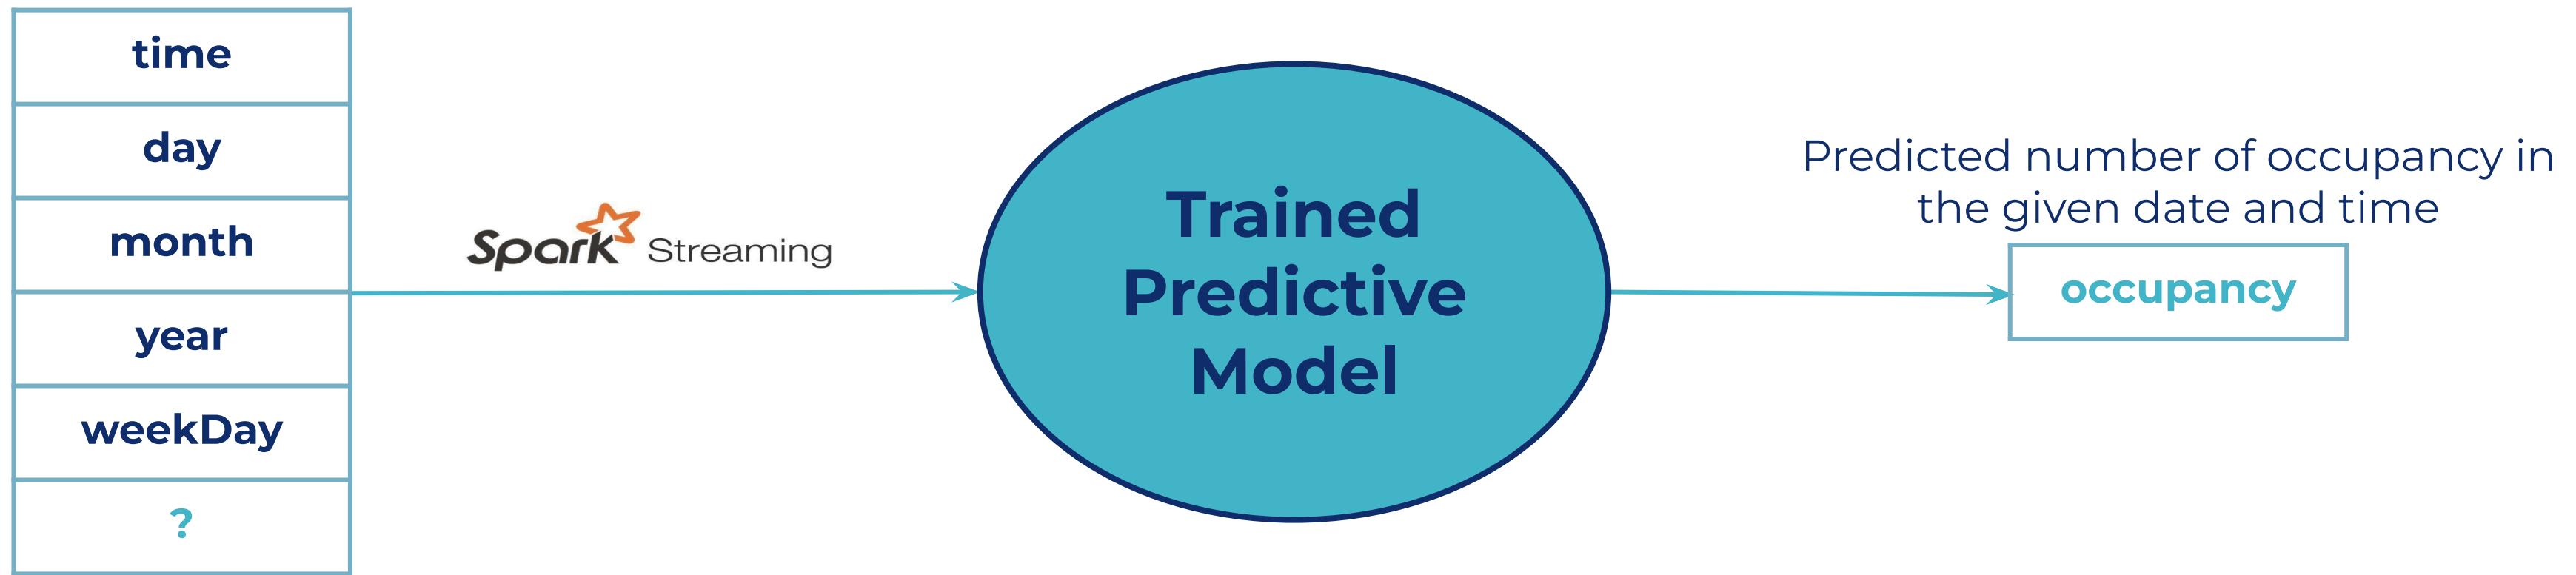
Data aggregation (for each store)

time	day	month	year	weekDay	ocupaccy
0	14	1	2016	3	5
1	14	1	2016	3	3
2	14	1	2016	3	4
3	14	1	2016	3	3
4	14	1	2016	3	2
5	14	1	2016	3	8
6	14	1	2016	3	12
7	14	1	2016	3	12
8	14	1	2016	3	23
9	14	1	2016	3	45
10	14	1	2016	3	55
11	14	1	2016	3	37
12	14	1	2016	3	42
13	14	1	2016	3	41
14	14	1	2016	3	38
15	14	1	2016	3	29
16	14	1	2016	3	33

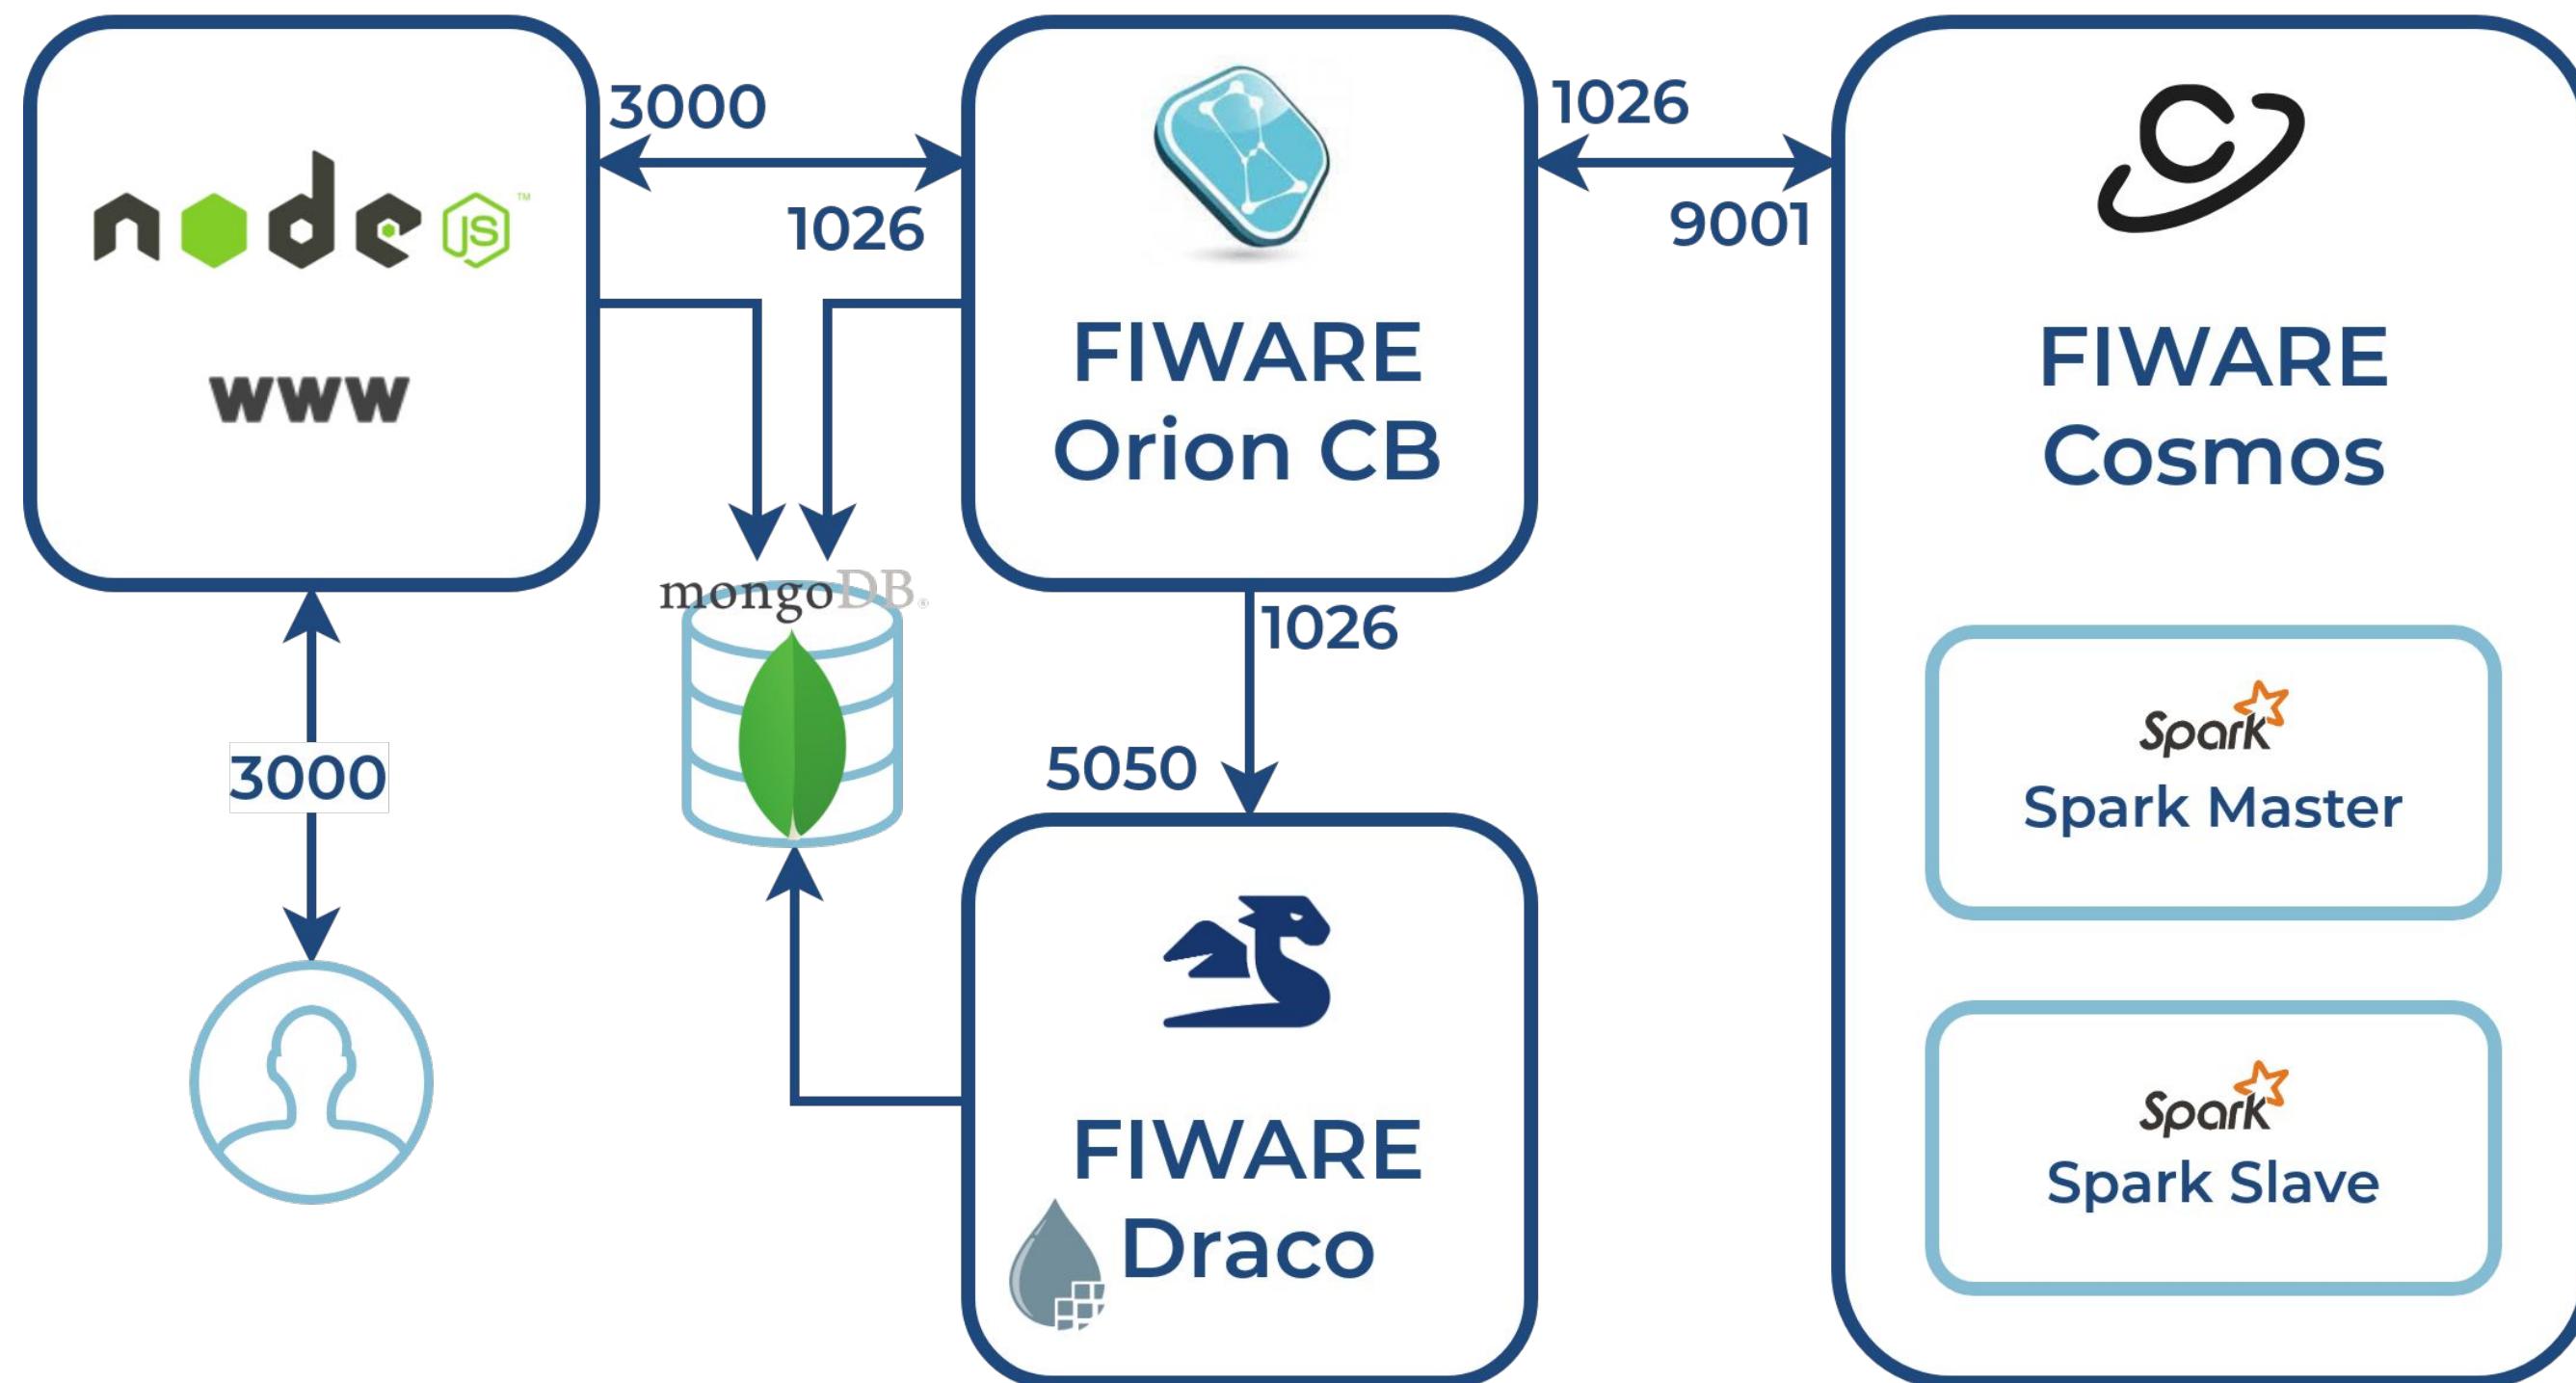
Training our model



Using our model to predict purchases

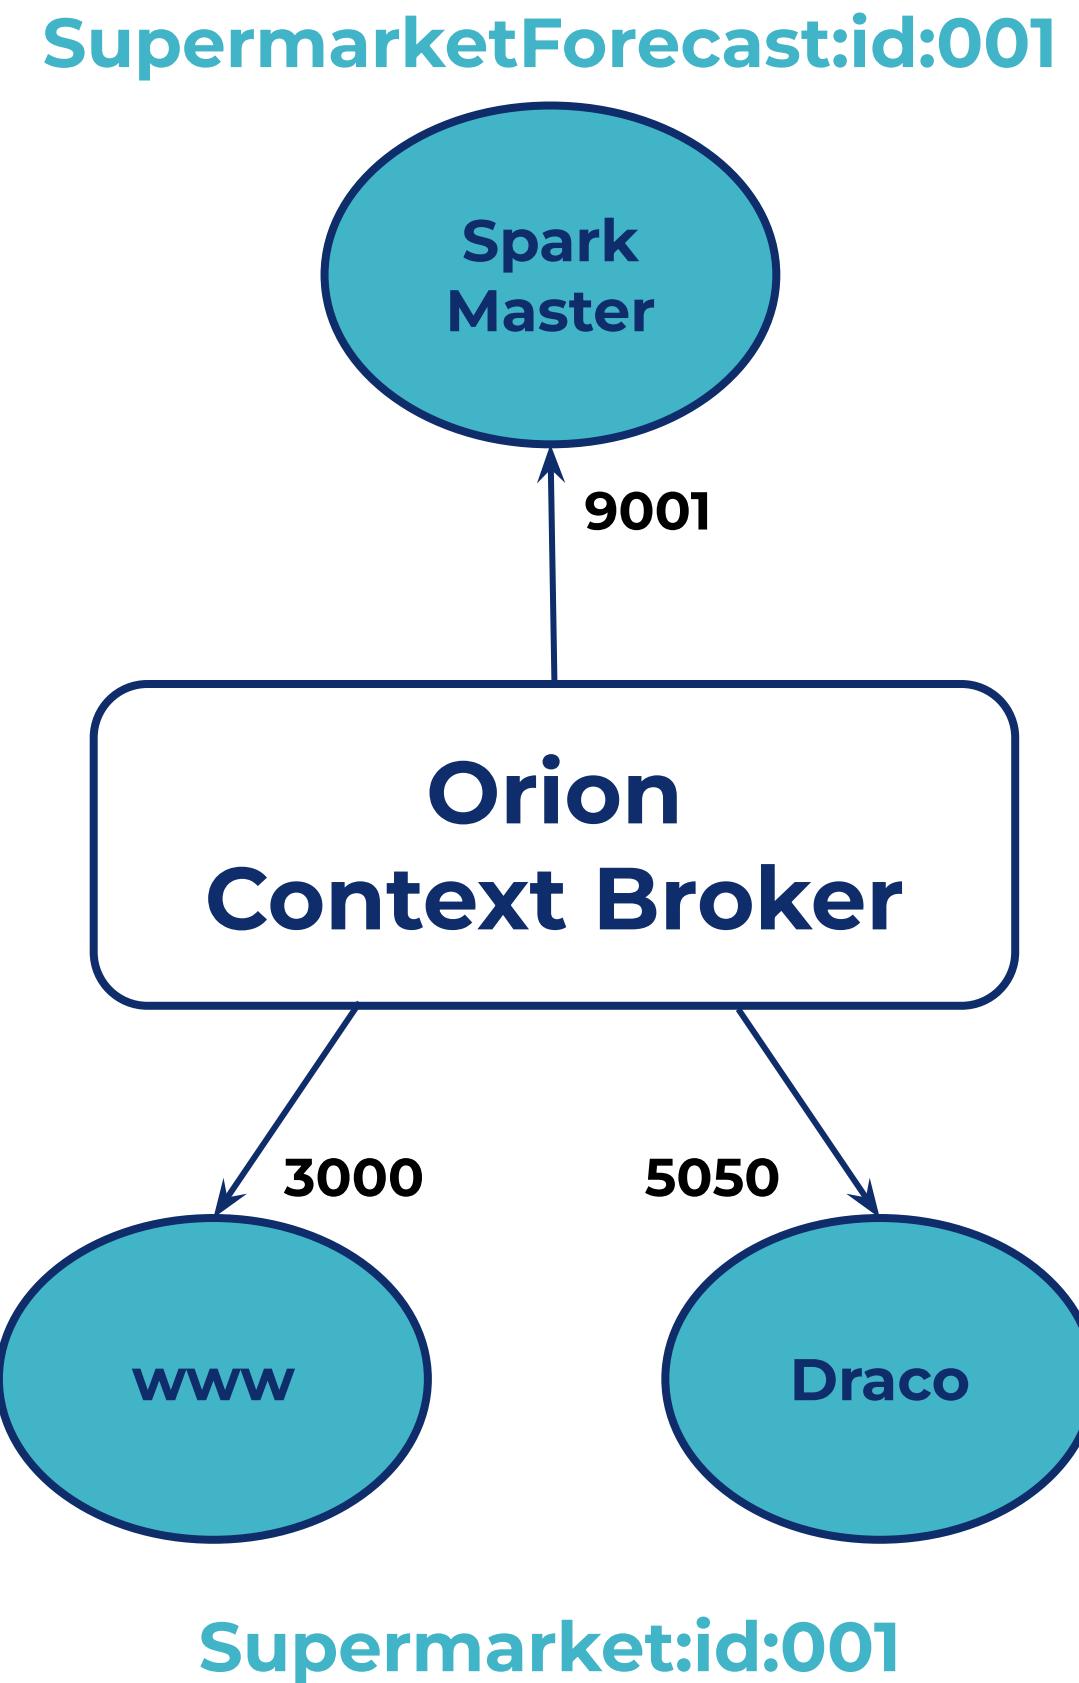


Architecture



Orion entities and subscriptions

```
{  
  "id": "urn:ngsi-ld:Supermarket:id:001",  
  "type": "Supermarket",  
  "name": "F Summit Supermarket",  
  "capacity": 200,  
  "occupancy": 50,  
  "dateObserved": "2023-06-11T21:09:57Z",  
  "@context": [  
    "https://smartdatamodels.org/context.jsonld"  
  ]  
}  
  
{  
  "id": "urn:ngsi-ld:SupermarketForecast:id:001",  
  "type": "Supermarket",  
  "name": "F Summit Supermarket Forecast",  
  "capacity": 200,  
  "occupancy": 50,  
  "dateObserved": "2023-06-11T23:09:57Z",  
  "@context": [  
    "https://smartdatamodels.org/context.jsonld"  
  ]  
}
```

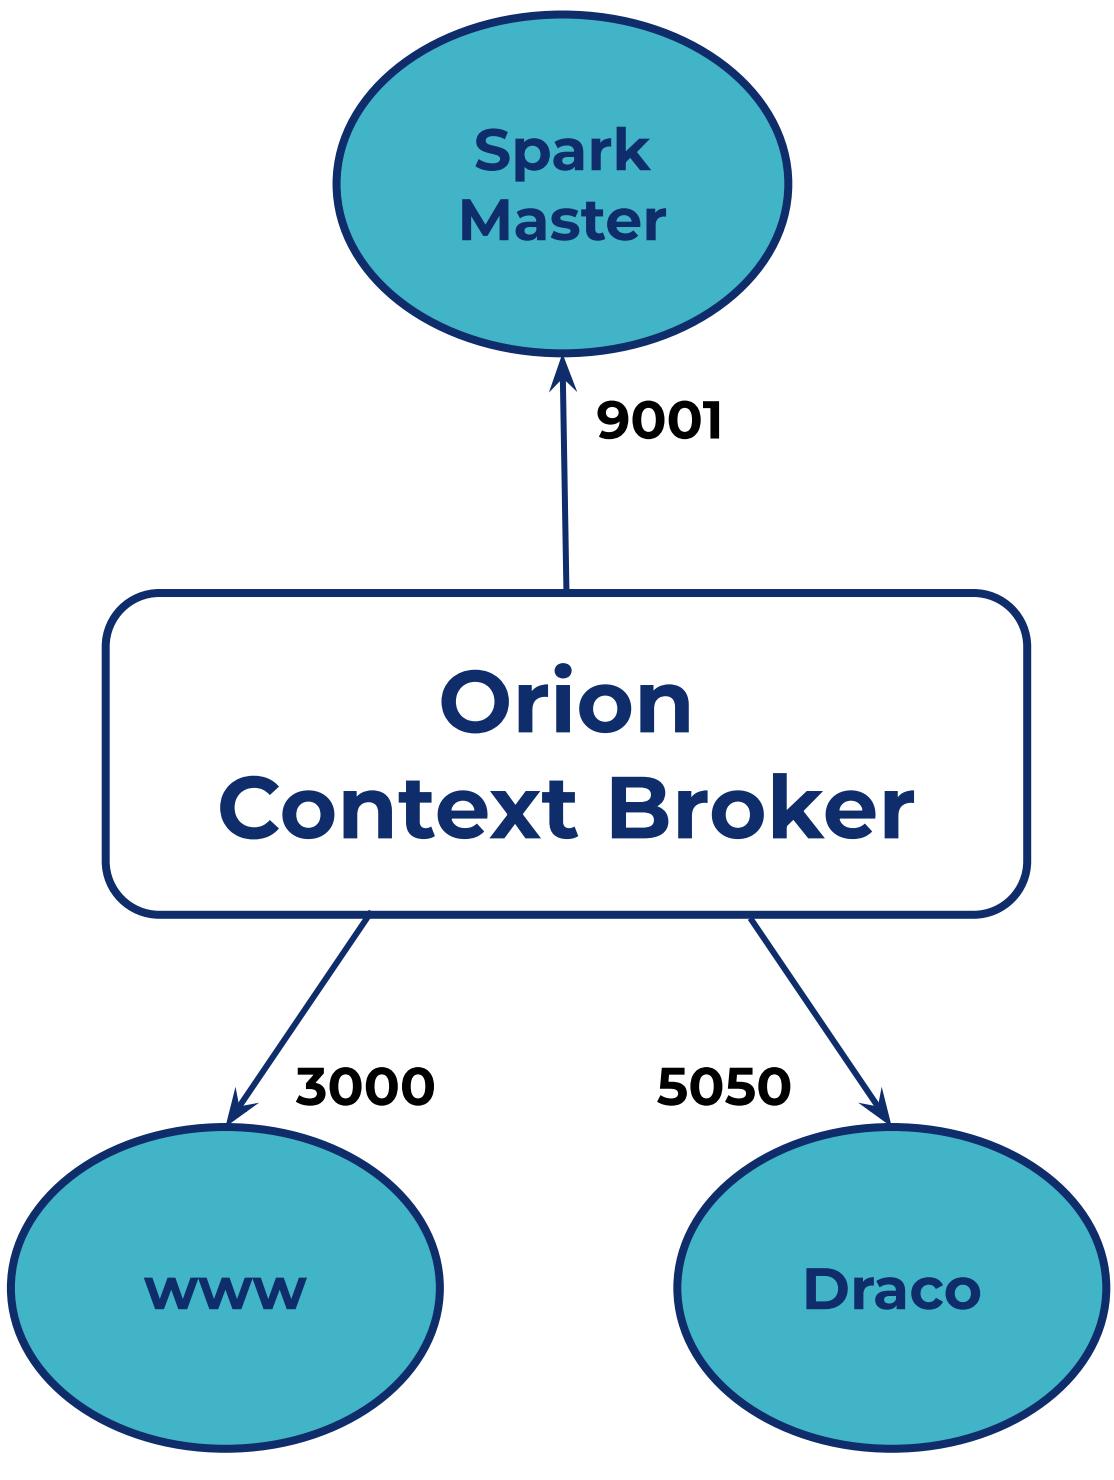


Orion entities and subscriptions

```
{  
  "id": "urn:ngsi-ld:Supermarket:001",  
  "type": "Supermarket",  
  "name": {  
    "value": "F. Summit Supermarket"  
    "type": "Property"  
  },  
  "capacity": {  
    "value": 50,  
    "type": "Property"  
  },  
  "occupancy": {  
    "value": 20,  
    "type": "Property"  
  },  
  "year":{  
    "value": 2023,  
    "type": "Property"  
  },  
  "month":{  
    "value": 11,  
    "type": "Property"  
  },  
  "day":{  
    "value": 22,  
    "type": "Property"  
  },  
  "time": {  
    "value": 22,  
    "type": "Property"  
  },  
  "weekDay": {  
    "value": 5,  
    "type": "Property"  
  },  
  "dateObserved": {  
    "value": "2023-11-22T22:01:12Z",  
    "type": "Property"  
  }  
  "@context": [  
  
  "https://uri.etsi.org/ngsi-ld/v1/ngsi-ld-core-context.jsonld"  "https://uri.etsi.org/ngsi-ld/v1/ngsi-ld-core-context.jsonld"  
]  
}
```

```
{  
  "id": "urn:ngsi-ld:SupermarketForecast:001",  
  "type": "Supermarket",  
  "name": {  
    "value": "F. Summit Supermarket Forecast"  
    "type": "Property"  
  },  
  "capacity": {  
    "value": 70,  
    "type": "Property"  
  },  
  "occupancy": {  
    "value": 50,  
    "type": "Property"  
  },  
  "year":{  
    "value": 2023,  
    "type": "Property"  
  },  
  "month":{  
    "value": 11,  
    "type": "Property"  
  },  
  "day":{  
    "value": 22,  
    "type": "Property"  
  },  
  "time": {  
    "value": 23,  
    "type": "Property"  
  },  
  "weekDay": {  
    "value": 5,  
    "type": "Property"  
  },  
  "dateObserved": {  
    "value": "2023-11-22T23:01:12Z",  
    "type": "Property"  
  }  
  "@context": [  
  
  "https://uri.etsi.org/ngsi-ld/v1/ngsi-ld-core-context.jsonld"  "https://uri.etsi.org/ngsi-ld/v1/ngsi-ld-core-context.jsonld"  
]  
}
```

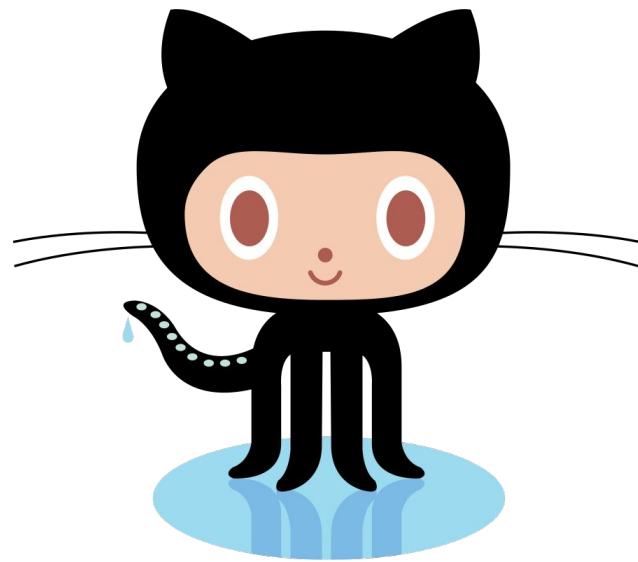
SupermarketForecast:id:001



Supermarket:id:001

Get the code!

<https://github.com/ging/fiware-mlops-supermarket>



Open your laptop

```
git clone https://github.com/qing/fiware-mlops-supermarket
```

```
docker-compose up
```

Open your browser: <http://localhost:3000>

IT WORKS !!!!

Demo



```
amunoz@amunoz: ~/Desktop/fiware-global-summit-berlin-2019-ml
```

```
amunoz@amunoz: ~/Desktop/fiware-global-summit-berlin-2019-ml 140x54
```

```
\n[INFO] Including org.apache.httpcomponents:httpclient:jar:4.5.9 in the shaded jar.\n[INFO] Including org.apache.httpcomponents:httpcore:ja  
r:4.4.11 in the shaded jar.\n[INFO] Including commons-logging:commons-logging:jar:1.2 in the shaded jar.\n[INFO] Including commons-codec:com  
mons-codec:jar:1.11 in the shaded jar.\n[WARNING] httpcore-4.4.11.jar, orion.spark.connector.prediction-1.0  
.jar define 249 overlapping classes: \n[WARNING] - org.apache.http.protocol.HttpRequestHandler\n[WARNING] - org.apache.http.impl.Chun  
kedOutputStream\n[WARNING] hStrategy\n[WARNING]  
- org.apache.http.H  
onVerifier\n[WARNING]  
on-1.0.1.jar, orion.sp  
e.RFC2109Spec\n[WARNING]  
n[WARNING] - org.apac  
.apache.http.impl.clien  
.protocol.RequestAuthCa  
eConnectionEvictor\n[WA  
417 overlapping classe  
handler.codec.http.mult  
g.apache.hadoop.metrics  
ntNamenodeProtocolProto  
oper.server.quorum.Lear  
codahale.metrics.Gauge  
overlapping classe: \n  
\n[WARNING] - org.apa  
ommons.codec.language.D  
- org.apache.commons.co  
g.apache.commons.codec.  
..\n[WARNING] orion.sp  
s: \n[WARNING] - org.  
ache.commons.codec.lang  
- org.apache.common  
he.commons.codec.langua  
e.bm.Languages$1\n[WAR  
ctor.prediction-1.0.1.j  
\n[WARNING] - org.apa  
NG] - org.apache.http  
- org.apache.http.conn  
pl.cookie.RFC2109Spec$1  
dle\n[WARNING] orion.sp  
ses: \n[WARNING] - or  
- org.apache.common  
ctoryImpl$2\n[WARNING]  
] - org.apache.common  
] - org.apache.common  
prediction-1.0.1.jar de  
ootstrap.WorkerPoolExe  
 class files are\n[WAR  
uber jar.\n[WARNING] Usually this is not harmful and you can skip these warnings,\n[WARNING] otherwise try to manually exclude artifacts ba  
sed on\n[WARNING] mvn dependency:tree -Ddetail=true and the above output.\n[WARNING] See http://maven.apache.org/plugins/maven-shade-plugin/  
\n[INFO] Replacing original artifact with shaded artifact.\n[INFO] Replacing /prediction-job/target/orion.spark.connector.prediction-1.0.1.j  
ar with /prediction-job/target/orion.spark.connector.prediction-1.0.1-shaded.jar\n[INFO] Dependency-reduced POM written at: /prediction-job/  
dependency-reduced-pom.xml\n[INFO] -----\n[INFO] BUILD SUCCESS\n[INFO] --  
-----\n[INFO] Total time: 10.832 s\n[INFO] Finished at: 2019-10-21T11:03:2  
7Z\n[INFO] -----
```

```
maven
```

```
| ./prediction-job
```

```
Open ▾
```



```
docker-compose.yml  
~/Desktop/fiware-global-summit-berlin-2019-ml
```

```
Save
```



```
14 orion:  
15   container_name: orion  
16   image: fiware/orion  
17   links:  
18     - mongo  
19   ports:  
20     - "1026:1026"  
21   command: -dbhost mongo -dbuser root -dbpwd example #-logLevel DEBUG  
22   networks:  
23     - fiware  
24 web:  
25   container_name: web  
26   build:  
27     context: ./web  
28   ports:  
29     - "3000:3000"  
30   depends_on:  
31     - orion  
32   networks:  
33     - fiware  
34   command: bash -c "sh /entities/createPredictionEntities.sh && sh /entities/  
subscribeReqPredictionTicket.sh && sh /entities/subscribeResPredictionTicket.sh && npm start"  
35   environment:  
36     - URL_CB=http://orion:1026/v2/entities/ReqTicketPrediction1/attrs  
37     - MONGO_URI=mongodb://root@example:mongo:27017/sth_test?authSource=admin  
38   volumes:  
39     - ./entities:/entities  
40 spark-master:  
41   image: bde2020/spark-master:2.4.4-hadoop2.7  
42   container_name: spark-master  
43   ports:  
44     - "8080:8080"  
45     - "7077:7077"  
46     - "9001:9001"  
47 # depends_on:  
48 #   - maven  
49   environment:
```

```
YAML ▾
```

```
Tab Width: 8 ▾
```

```
Ln 24, Col 7 ▾
```

```
INS
```

```
-----  
Uber jar.\n[WARNING] Usually this is not harmful and you can skip these warnings,\n[WARNING] otherwise try to manually exclude artifacts ba  
sed on\n[WARNING] mvn dependency:tree -Ddetail=true and the above output.\n[WARNING] See http://maven.apache.org/plugins/maven-shade-plugin/  
\n[INFO] Replacing original artifact with shaded artifact.\n[INFO] Replacing /prediction-job/target/orion.spark.connector.prediction-1.0.1.j  
ar with /prediction-job/target/orion.spark.connector.prediction-1.0.1-shaded.jar\n[INFO] Dependency-reduced POM written at: /prediction-job/  
dependency-reduced-pom.xml\n[INFO] -----  
-----\n[INFO] Total time: 10.832 s\n[INFO] Finished at: 2019-10-21T11:03:2  
7Z\n[INFO] -----
```

**FIWARE
Global
Summit**

**From Data
to Value**

OPEN SOURCE
OPEN STANDARDS
OPEN COMMUNITY

MLOps Life cycle real-time application

José Andrés Muñoz Arcentales, Javier Conde Díaz, Joaquín Salvachúa

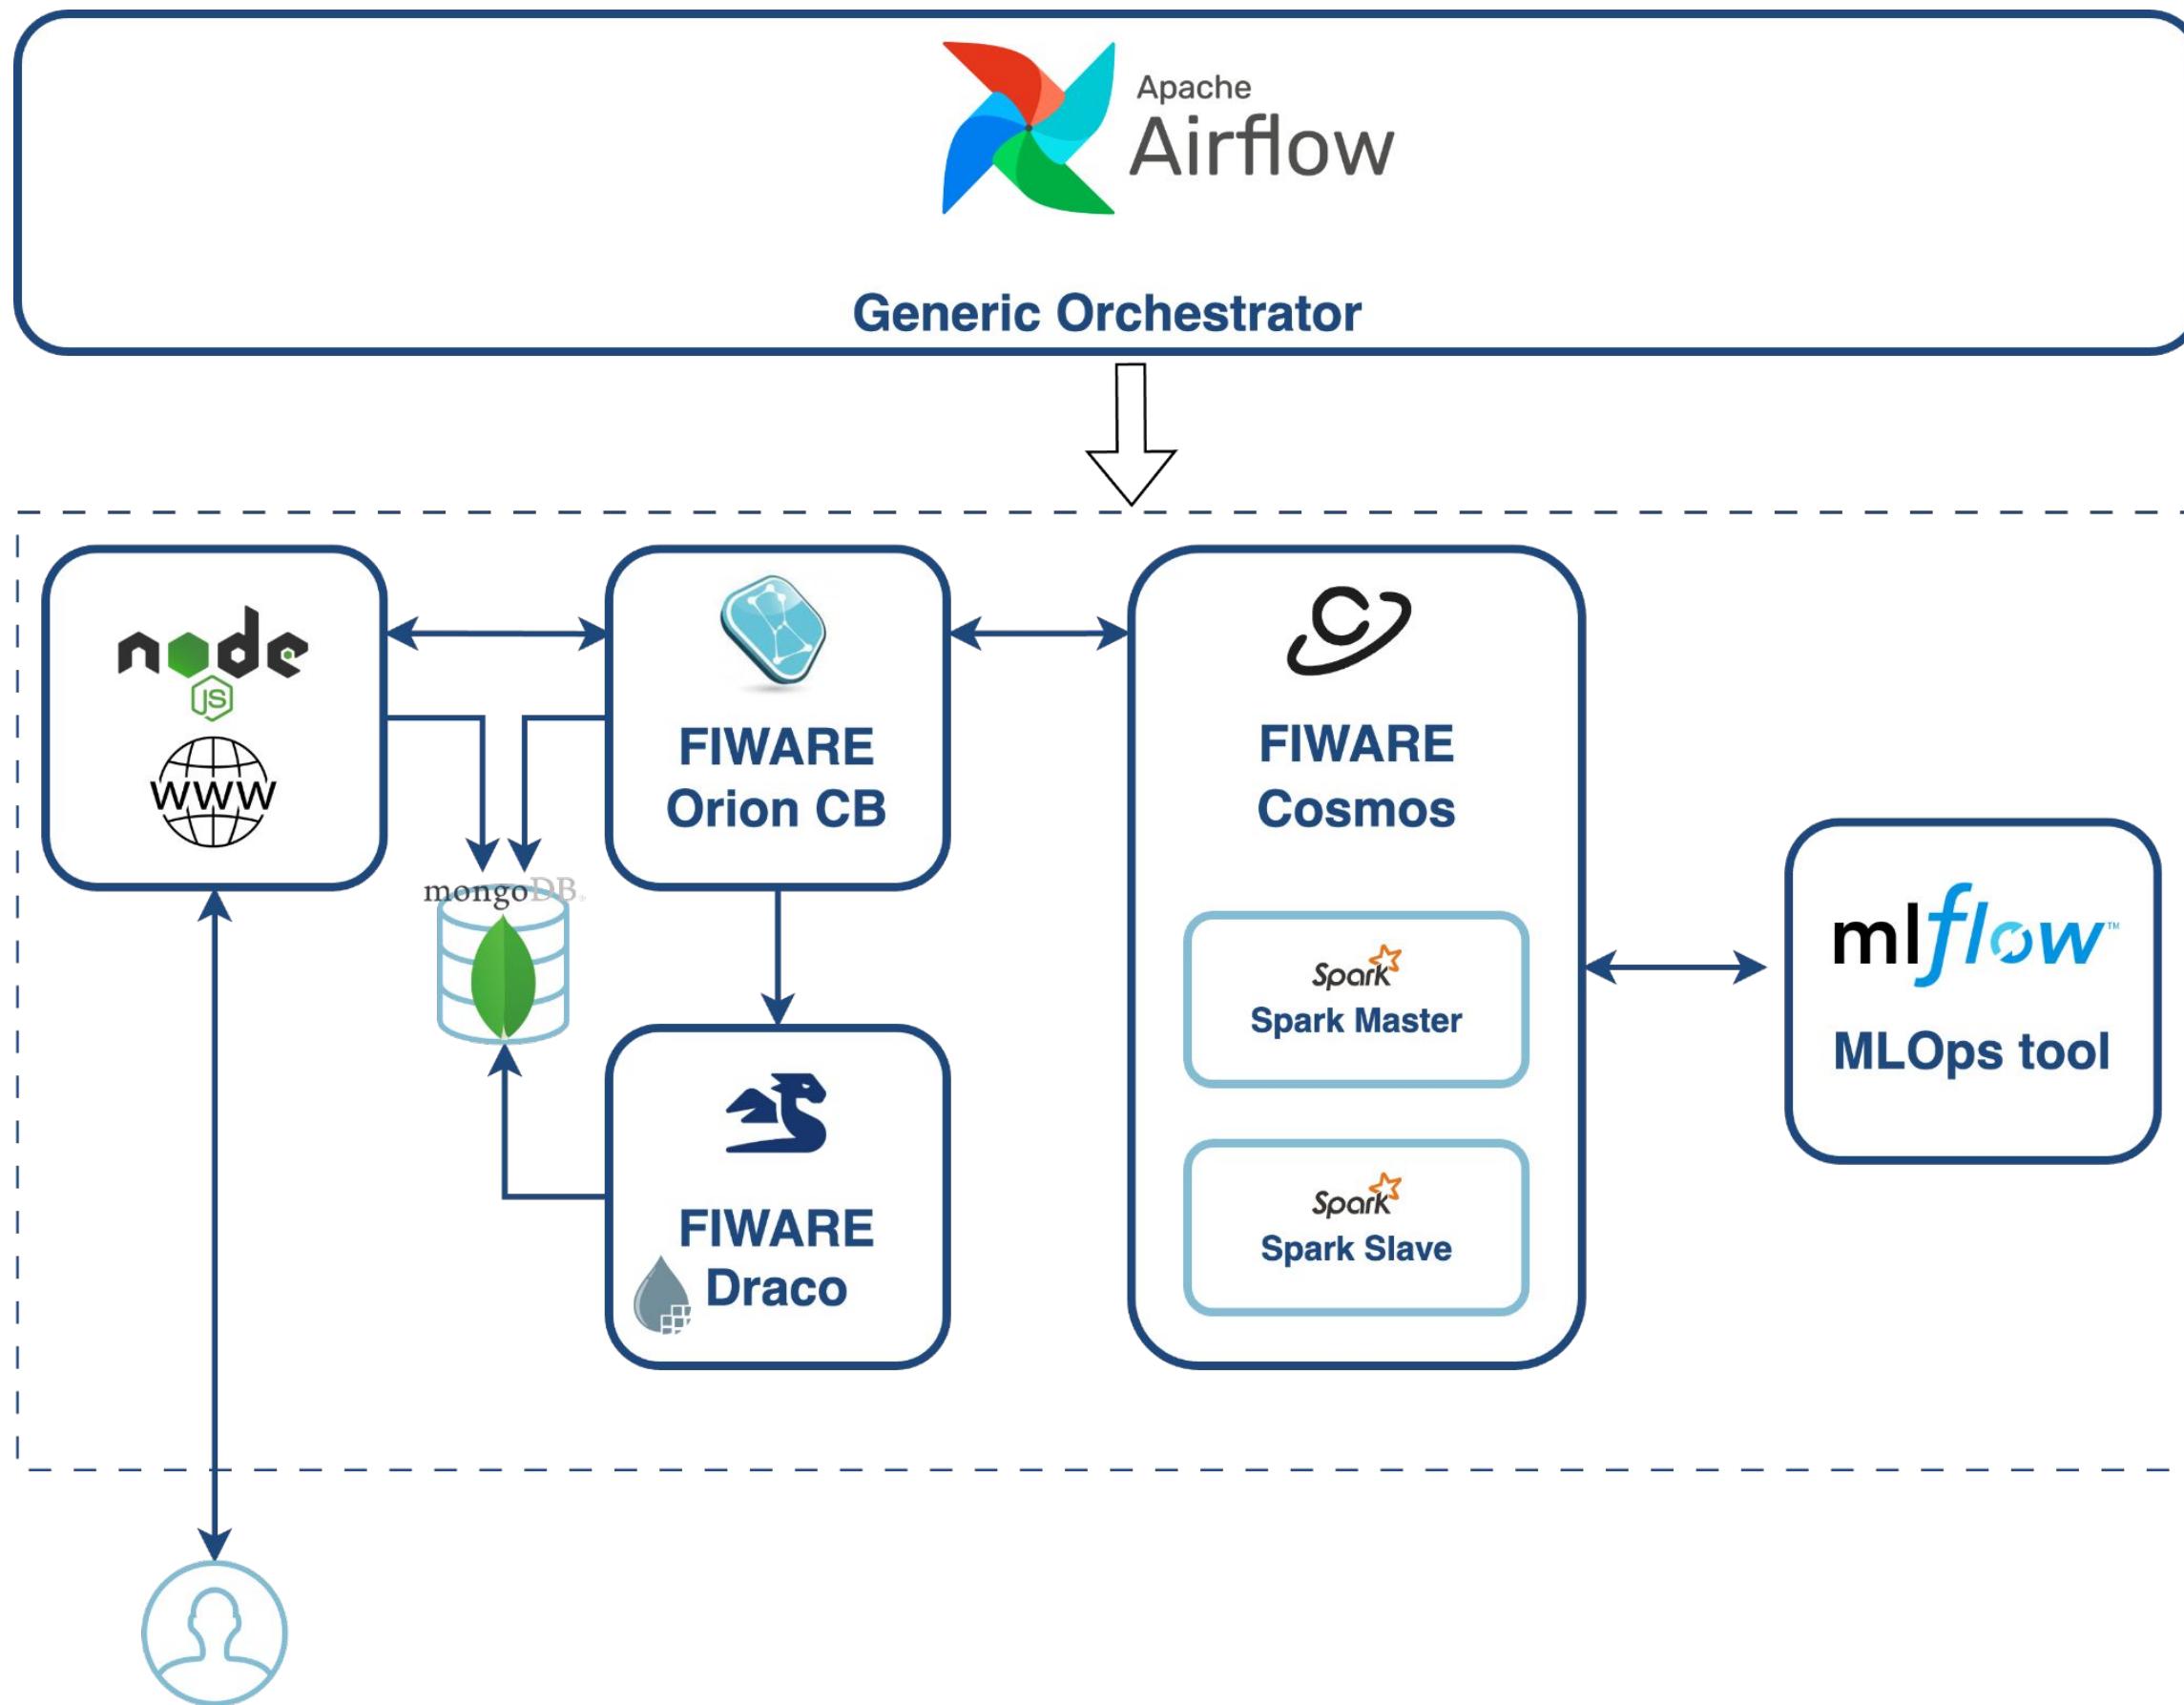
Vienna, Austria

12-13 June, 2023

#FIWARESummit



Architecture





Find Us On



Stay up to date

JOIN OUR NEWSLETTER

Be certified and featured



Hosting Partner



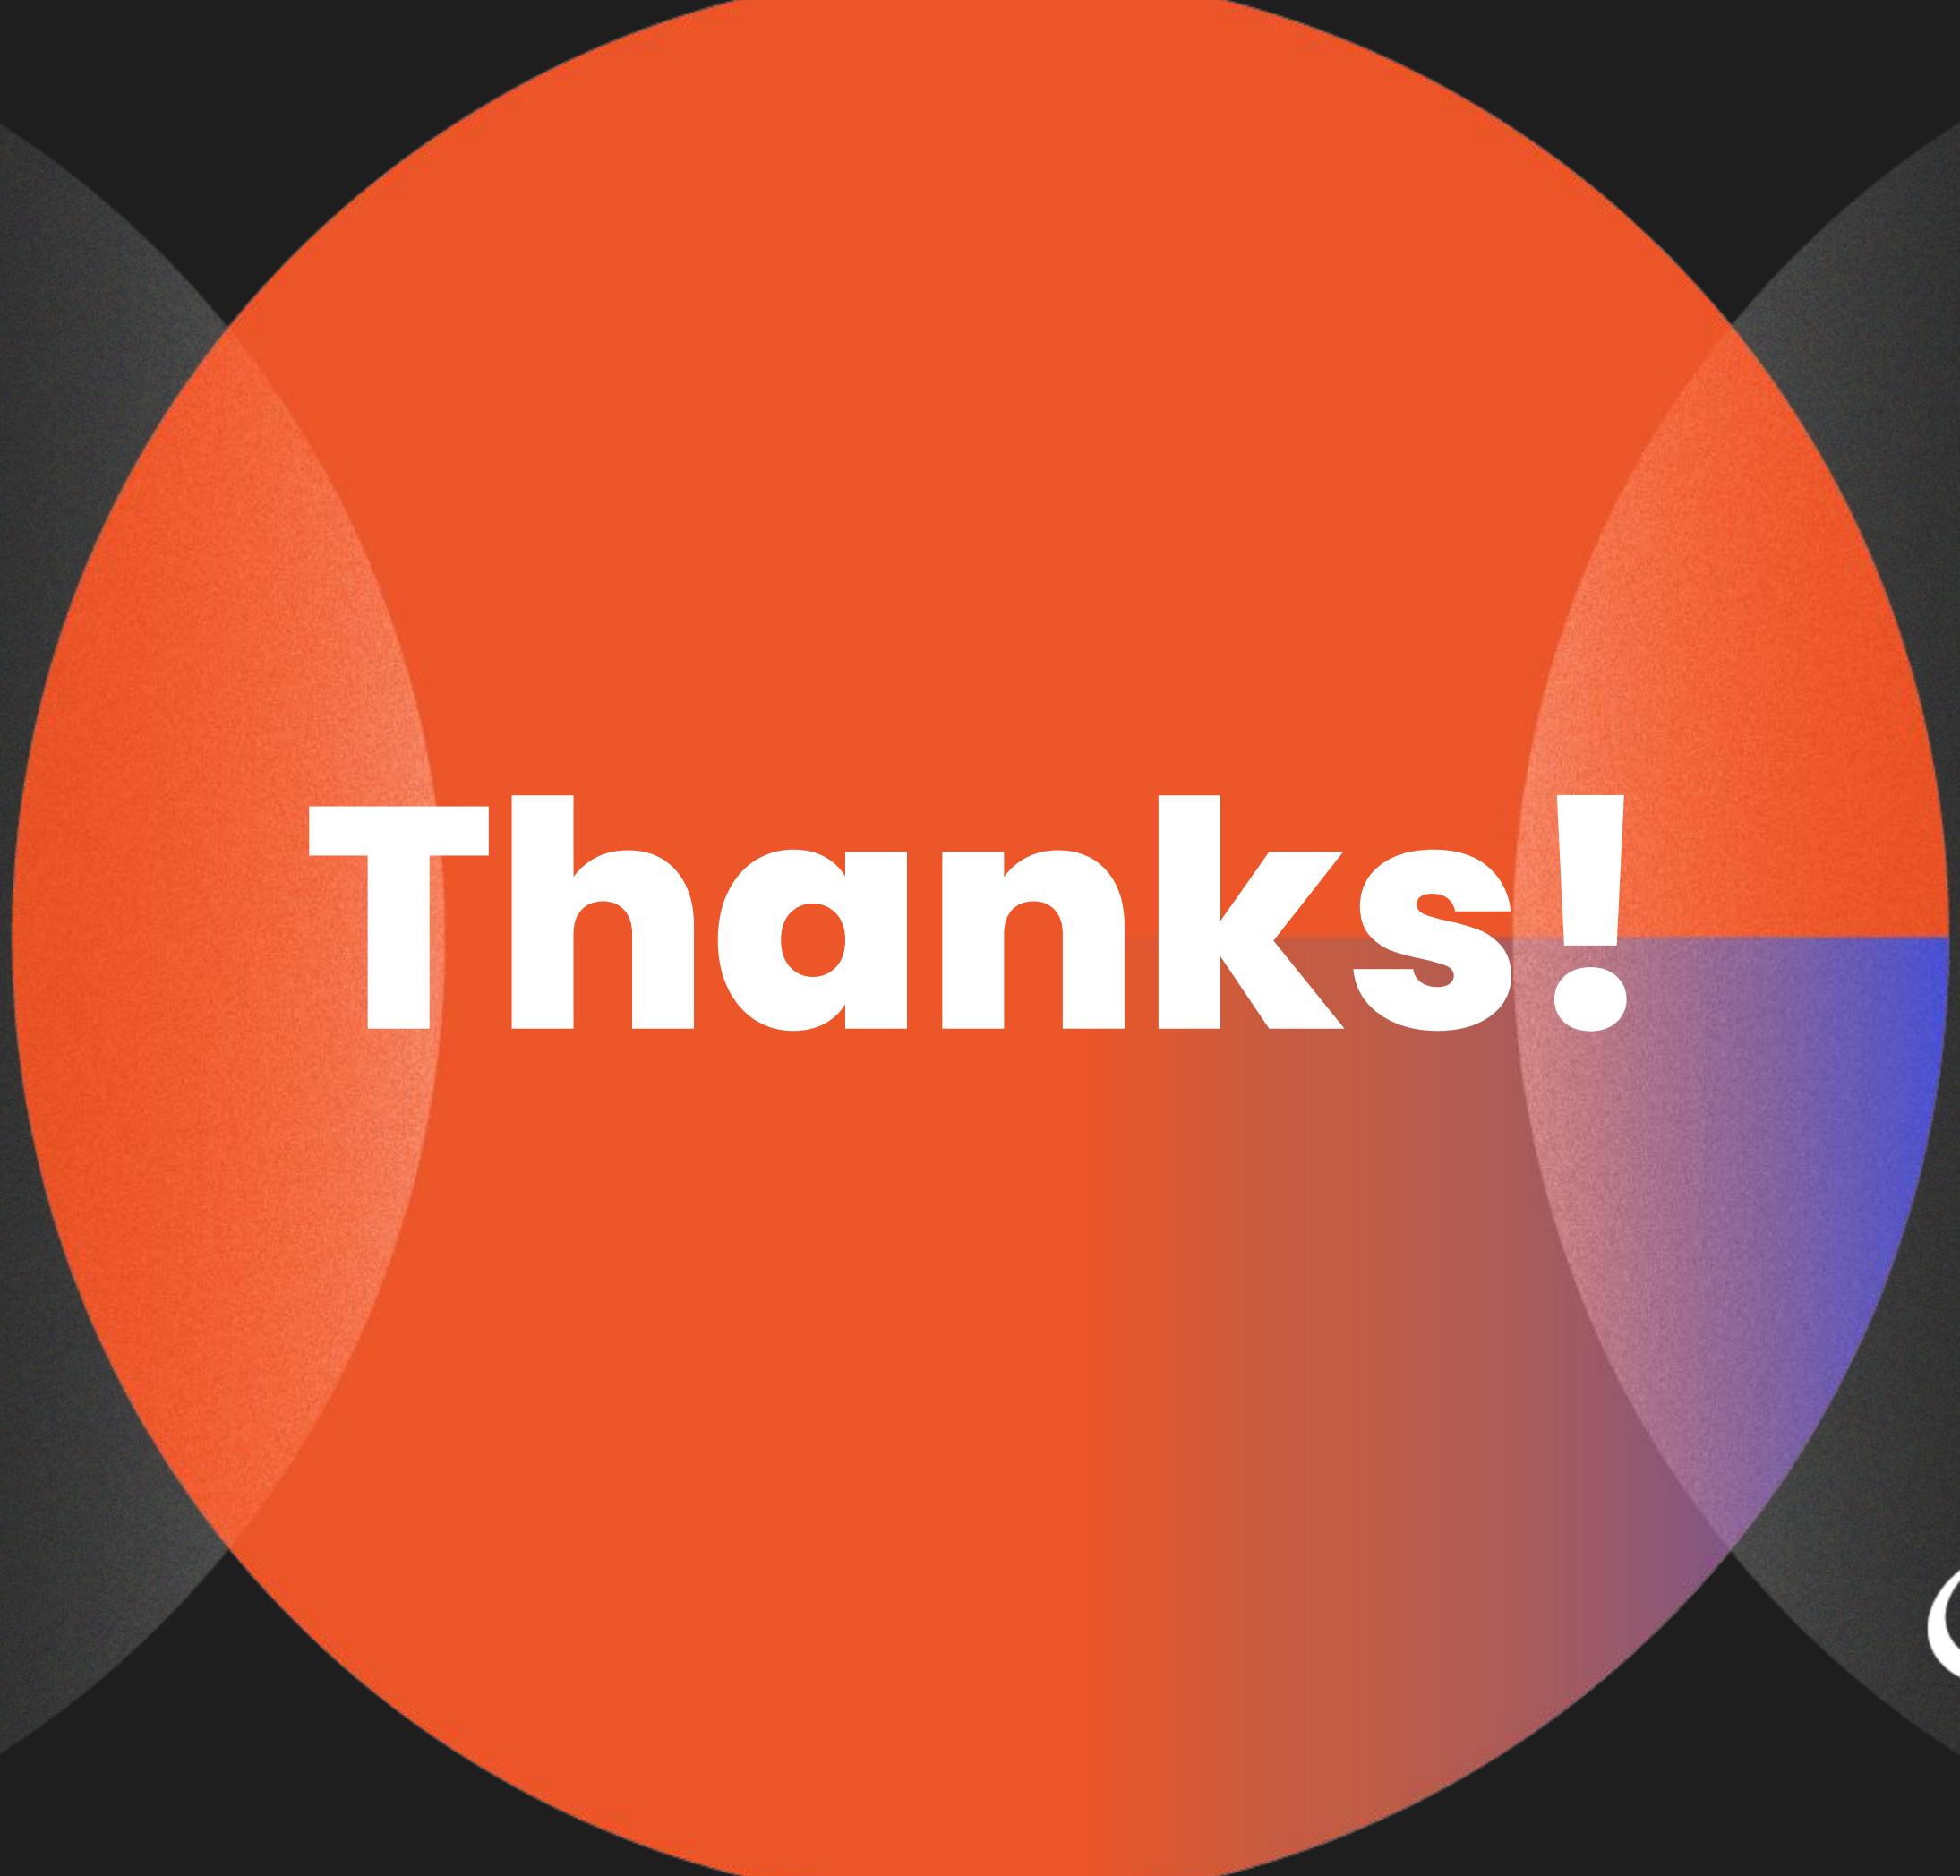
Keystone Sponsors



Media Partners



FIWARE
Global
Summit



Thanks!

Vienna, Austria
12–13 June, 2023
#FIWARESummit

